

数据组织与存储



徐辰
华东师范大学
数据科学与工程学院
cxu@dase.ecnu.edu.cn

Outline

- How to store (data) record on disk
- How the DBMS locate records
- How to store record with variable size
- How to modify record

Overview

Structured data has a **structure** in the storage format

Name Addr ...



Files



Blocks



Records



Data Items

Each row is a **record**

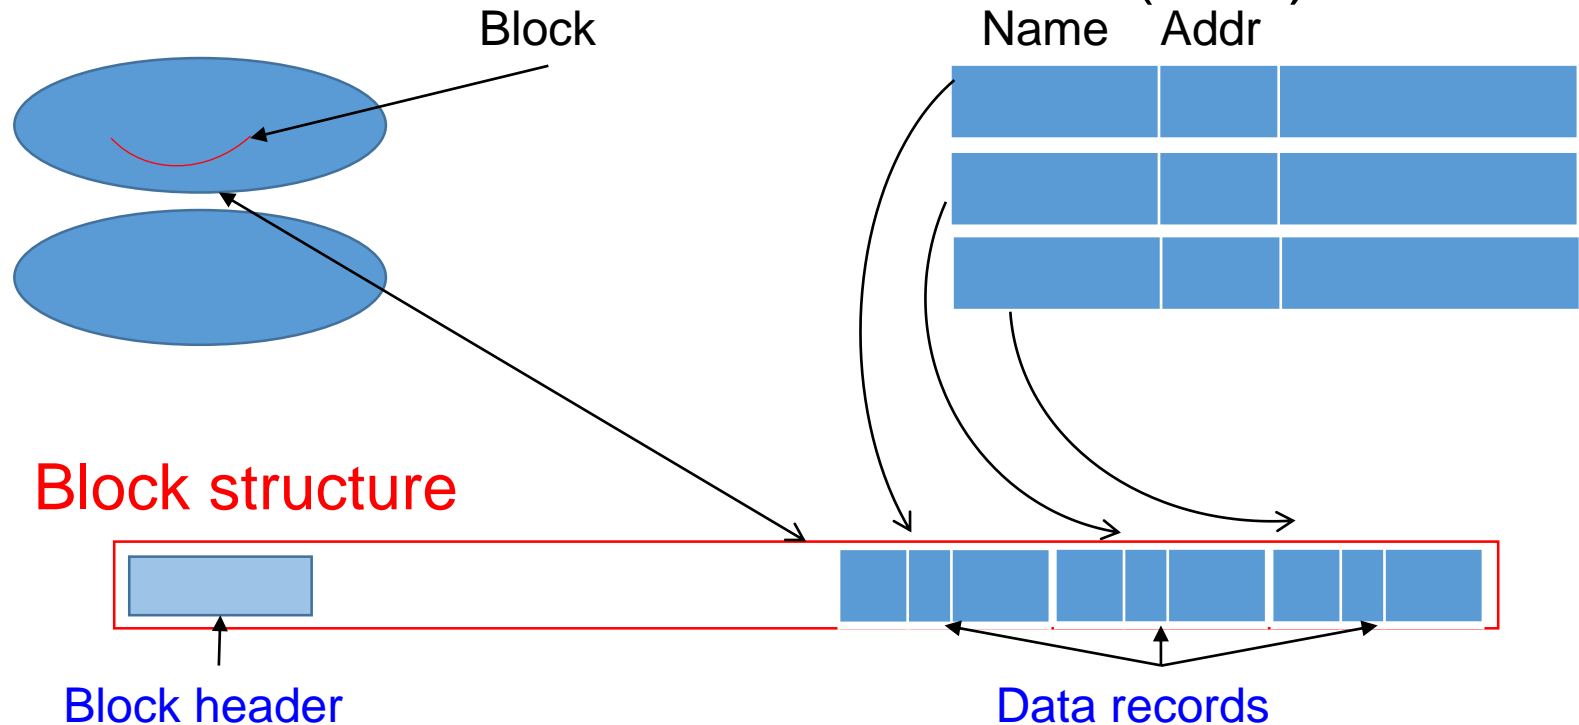
Records are **stored** in a **file**

A file consists of one/more **disk blocks**

How do we store records in disk blocks?

Block Structure

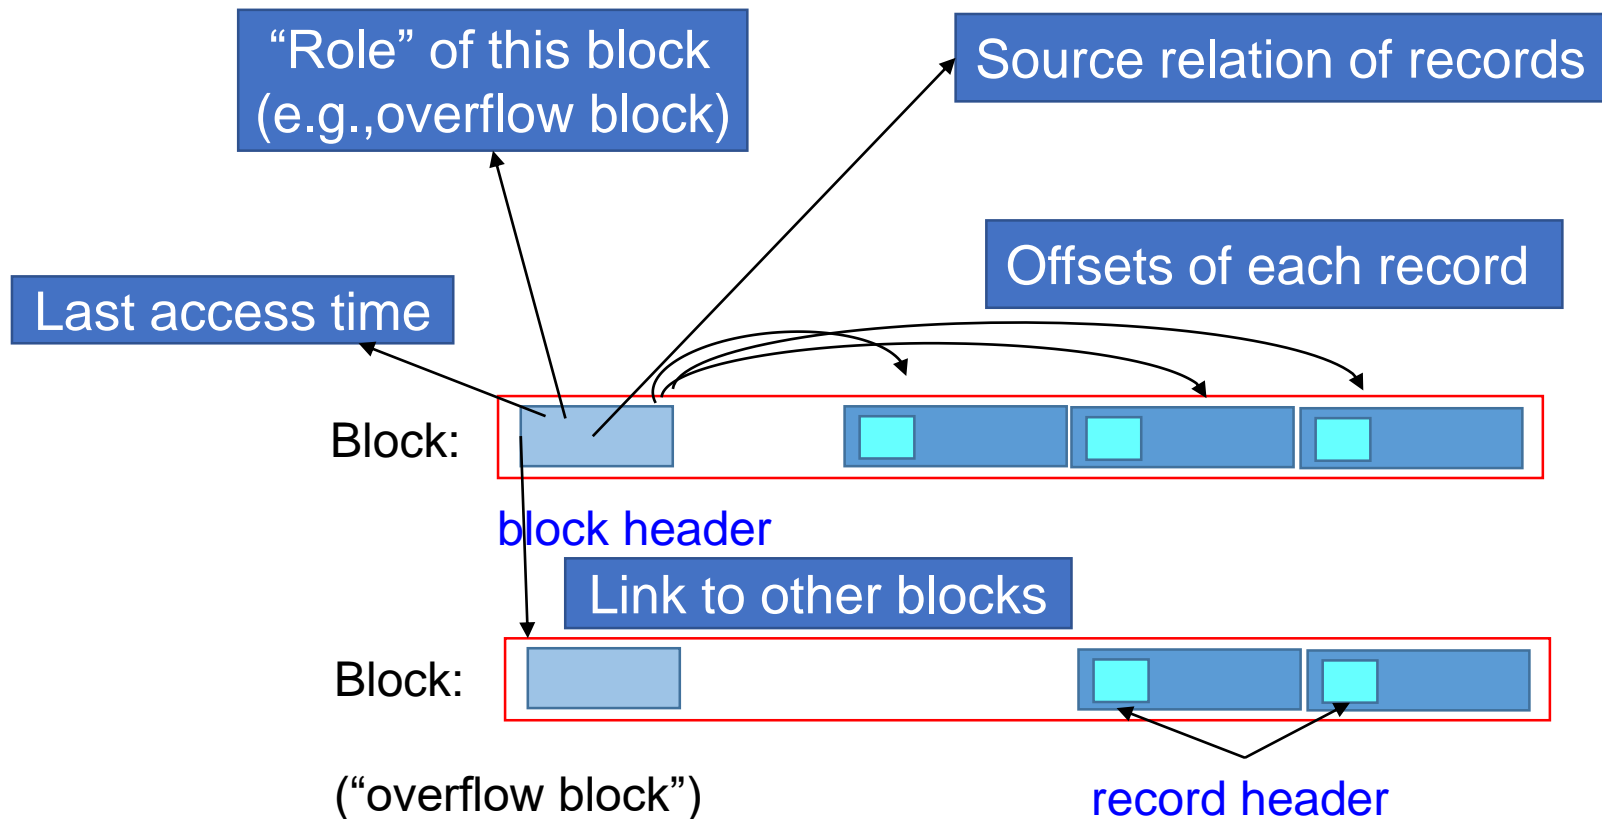
There is a **block structure (organization)** in a disk block when records are stored in a (disk) block



Block structure = block header + records


Block Header

Contain **info** on the records stored in this block



Record Format

- A record consist of
 - **Record header**
 - Pointer to record scheme
 - Length of record
 - Timestamp that the record was last updated
 - **Record fields**
- A record will contain “**padding**” bytes to satisfy **alignment requirement**



Outline

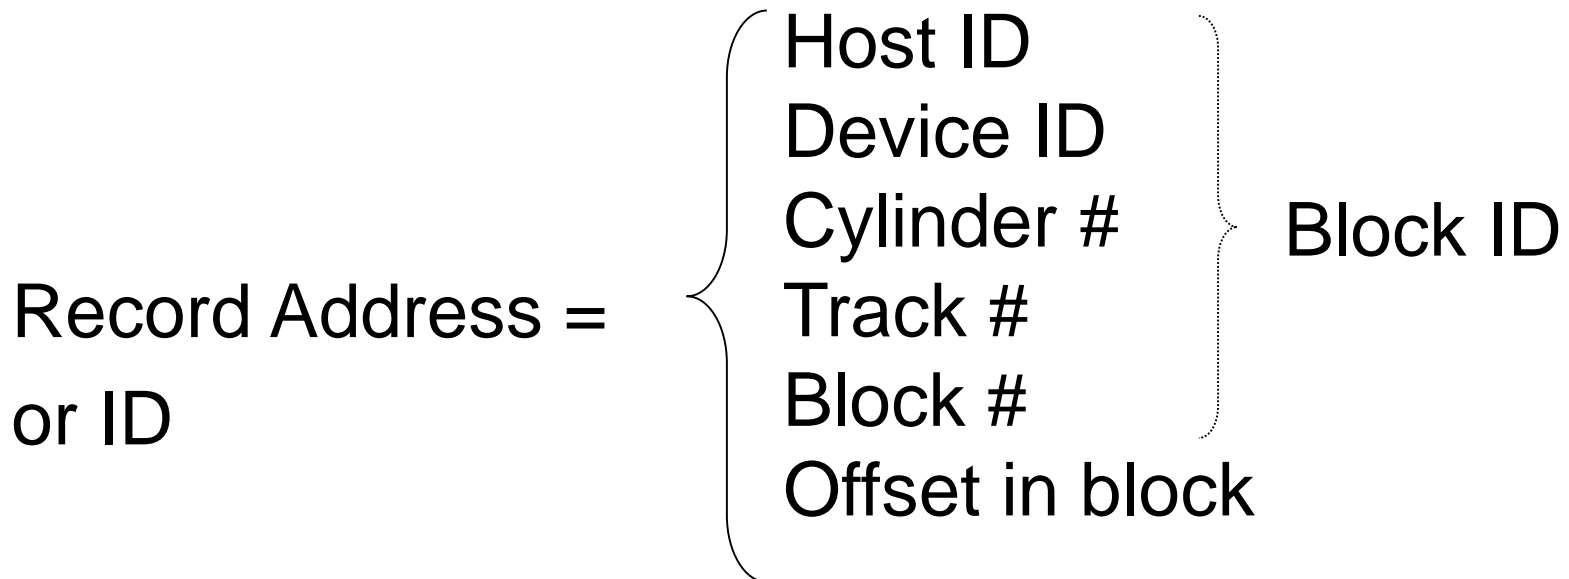
- How to store (data) record on disk
- How the DBMS locate records
- How to store record with variable size
- How to modify record

Blocks/records Address

- Two types of address
 - **Database address**: data stored **on disk**
 - Physical address
 - Logical address
 - **(Virtual) memory address**: data stored **in memory**


Physical Address

- **Direct** addressing format
- Byte strings determine the place of block or record in the secondary storage.



Problem of Physical Address

2 records X1 and X2 are referencing the record Y using physical database address.

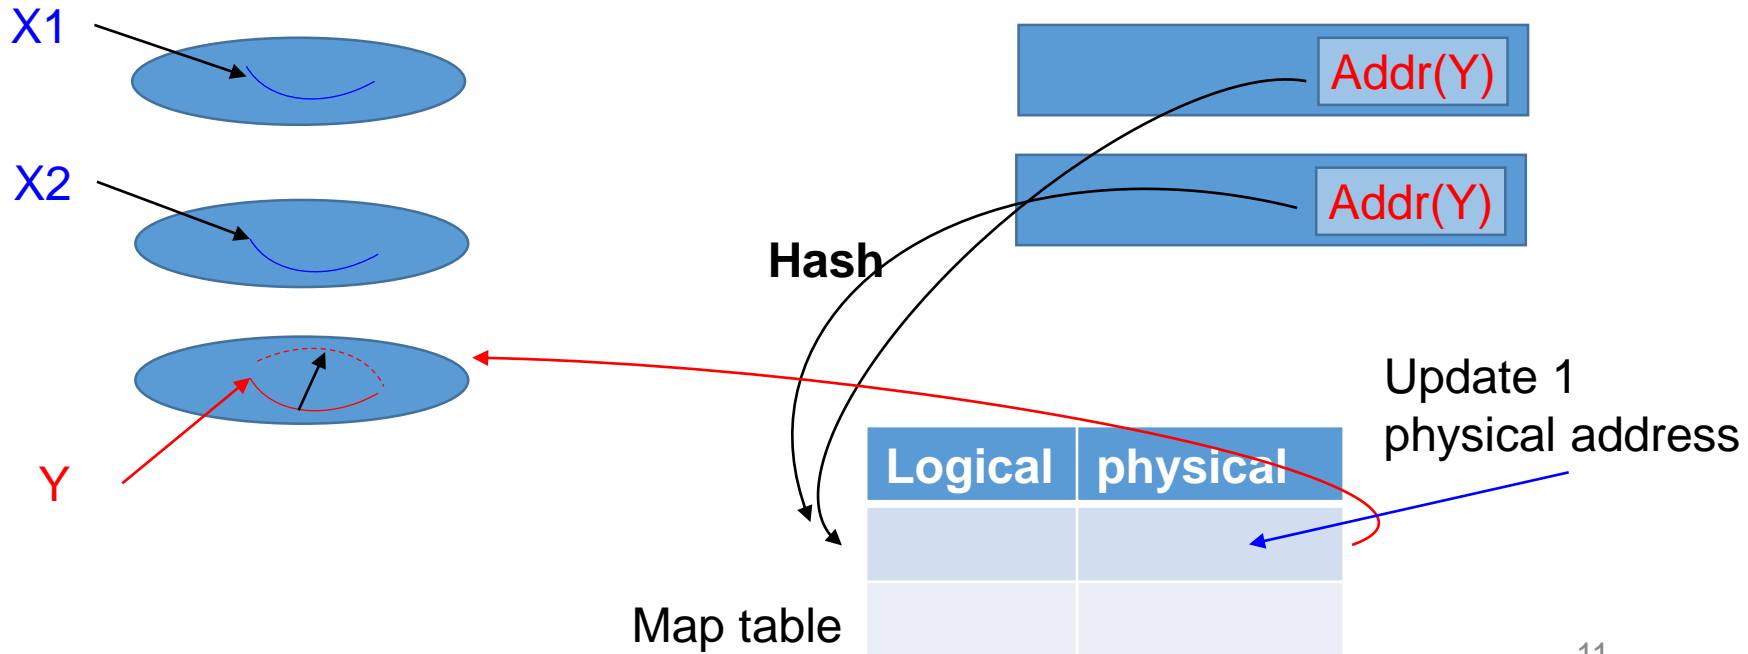


Problem: Move Y to different block on the disk, we must **update** multiple physical addresses.

Logical Address

- **Indirect** addressing format: **map table** (on disk)
- Record ID is arbitrary bit string


Only update physical address in the map table



Structured Address

- **Combination**

- Physical address of block + offset of record
- 2 way represent offset
 - **Direct**: (actual) offset in the block (e.g., offset = 345)
 - **Indirect**: index in the offset table (e.g., offset = 1)

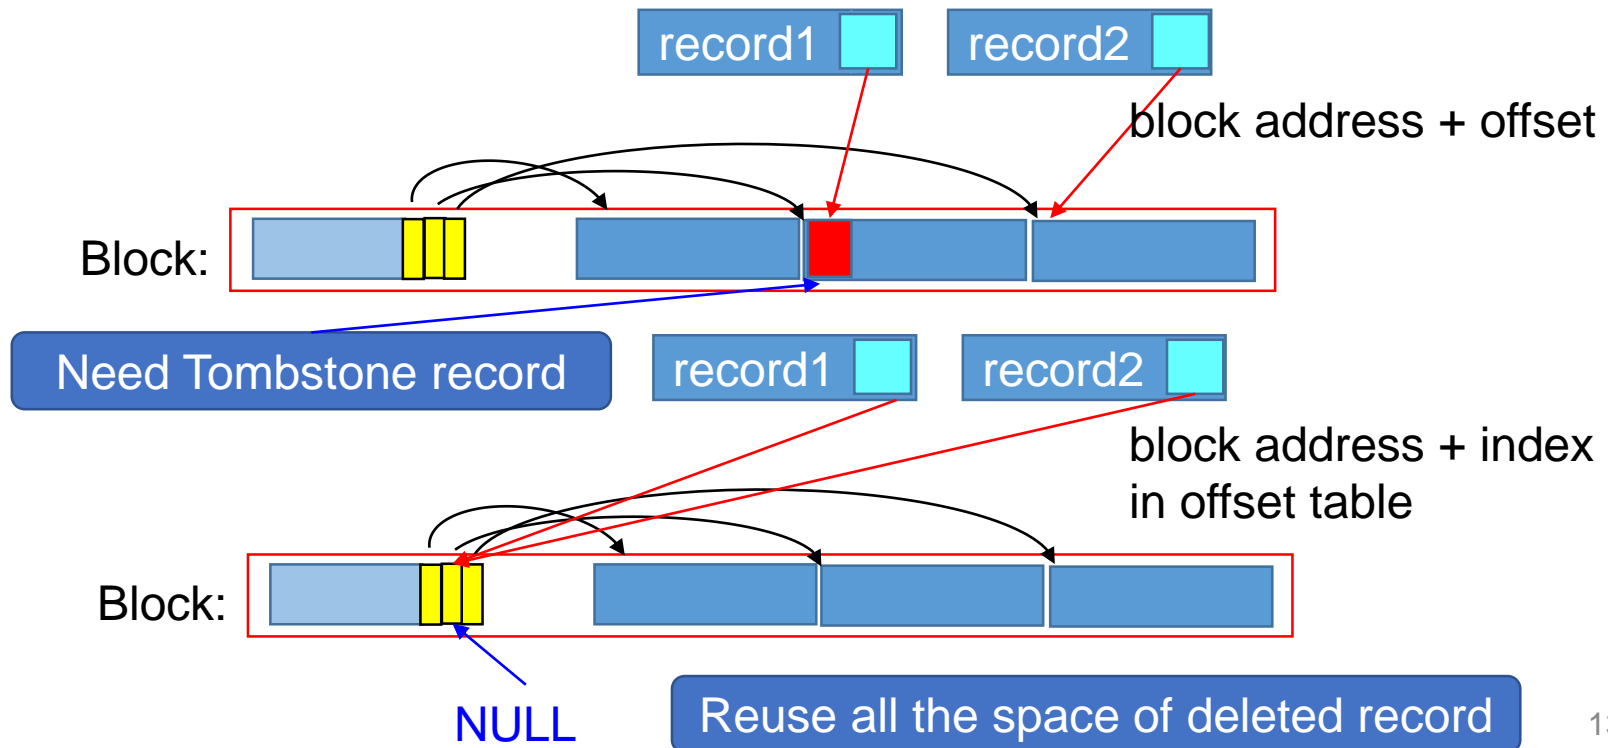


- **Advantage**
 - No need a global of map table

Direct Offset vs Indirect Offset

Consider the **deletion** of a record

- Direct offset: need a **tombstone record**
- Indirect offset: set offset to **NULL**



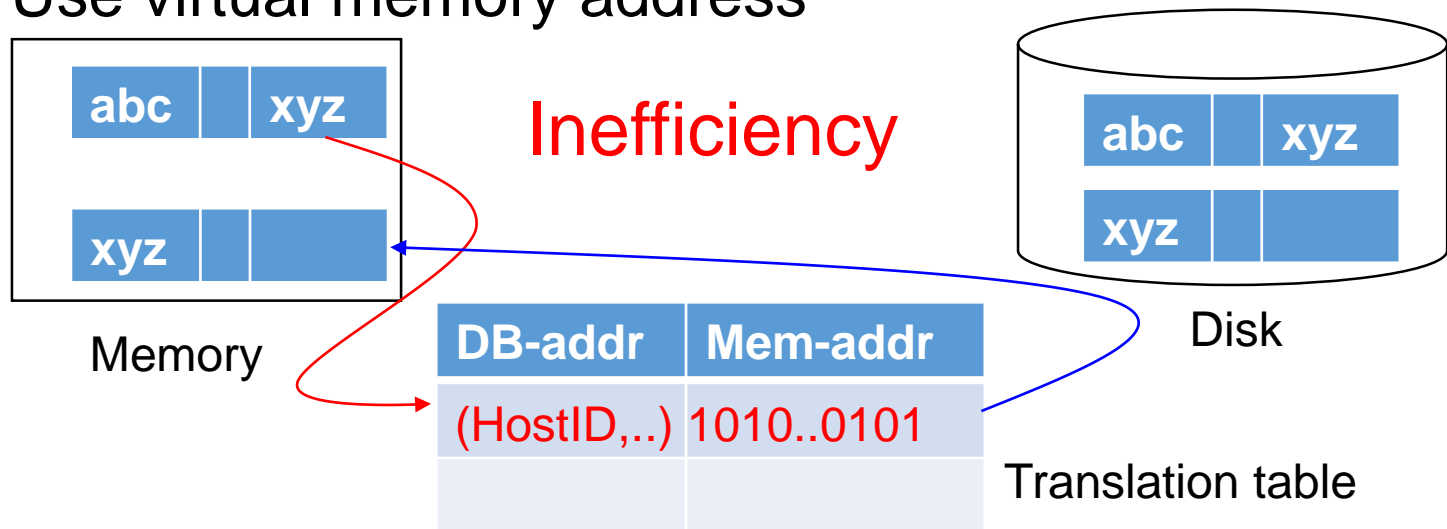
Blocks/records in Memory

- Block/record is read into main memory
 - **Must** use virtual memory address
- Problem
 - **Where** is the block right now?
 - **How to find** the **memory address** of a block?
- Solution for both problems
 - Use a **translation table**
 - **Map : Database address - > Virtual memory address**

How to Use Translation Table

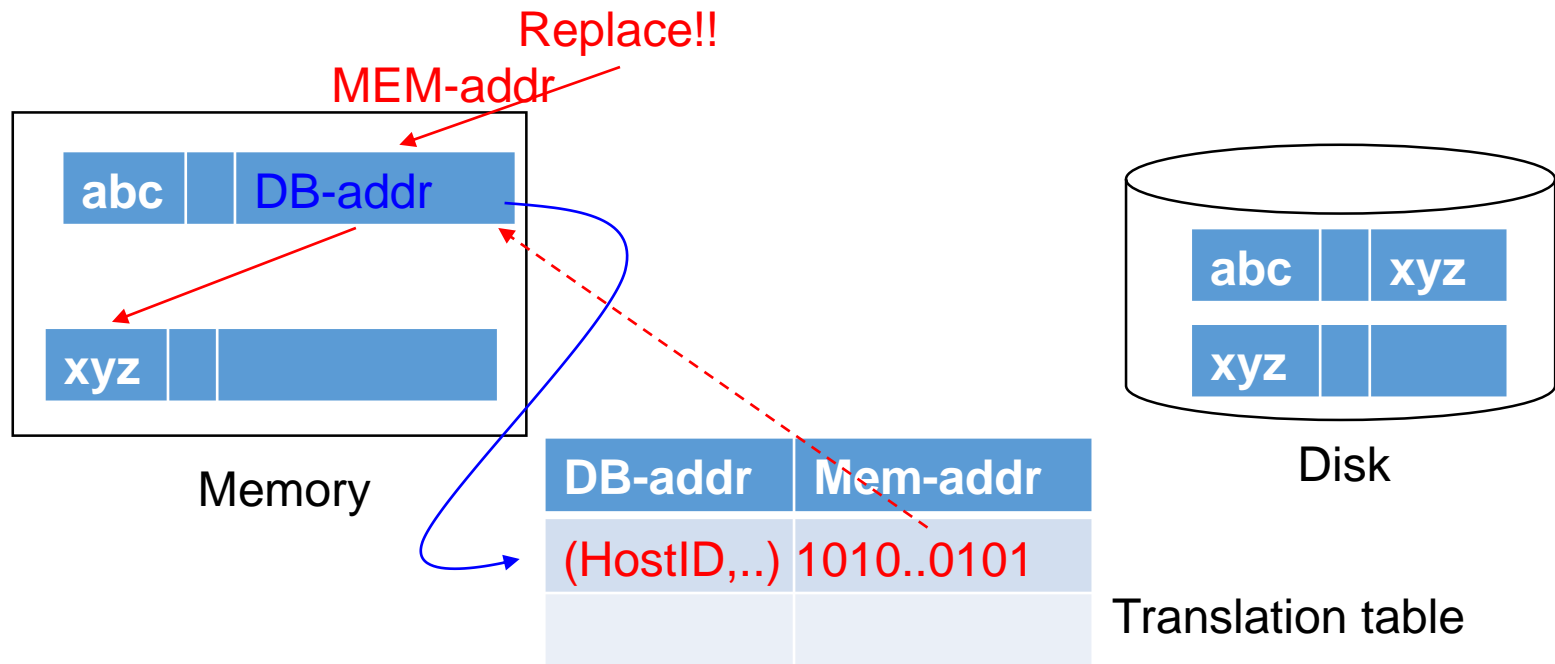
Hash database address x

- x is **not found** in translation table
 - Read the database object into memory (update the translation table)
- x is **found** in translation table
 - Use virtual memory address



Pointer Swizzling

Replace database address in a record by **virtual memory address** when the referenced data block/record resides in memory.

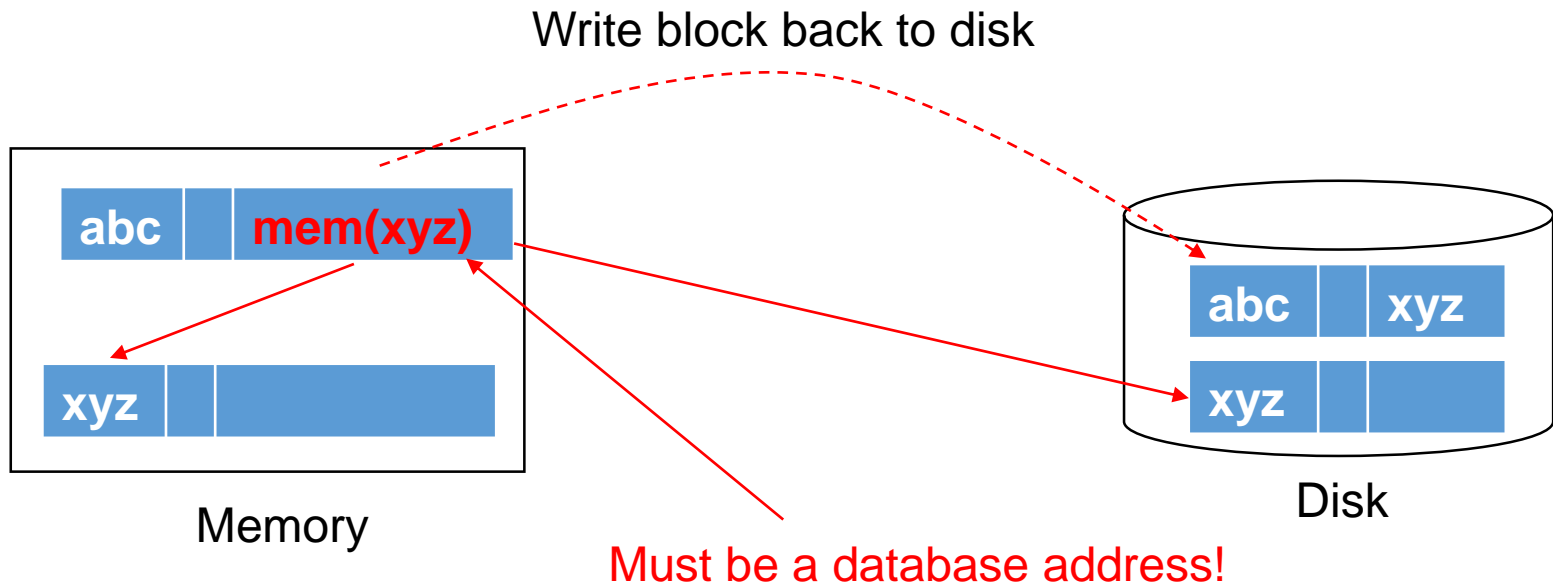


Swizzling Policies

- Automatic swizzling
 - Replace all database address in the records stored in the block when a block is read into memory.
- On-demand swizzling
 - Swizzle the database address when we access a record that has not been swizzled.
- No swizzling
 - Never to swizzle pointers
- Programmer control

Problems caused by Pointer Swizzling


- Problem 1
 - **Unswizzle** all virtual addresses before we write a block back to disk.



Problems caused by Pointer Swizzling

- Problem 2

- Release memory used by a disk block can result in **invalid record references** in other records.



Outline

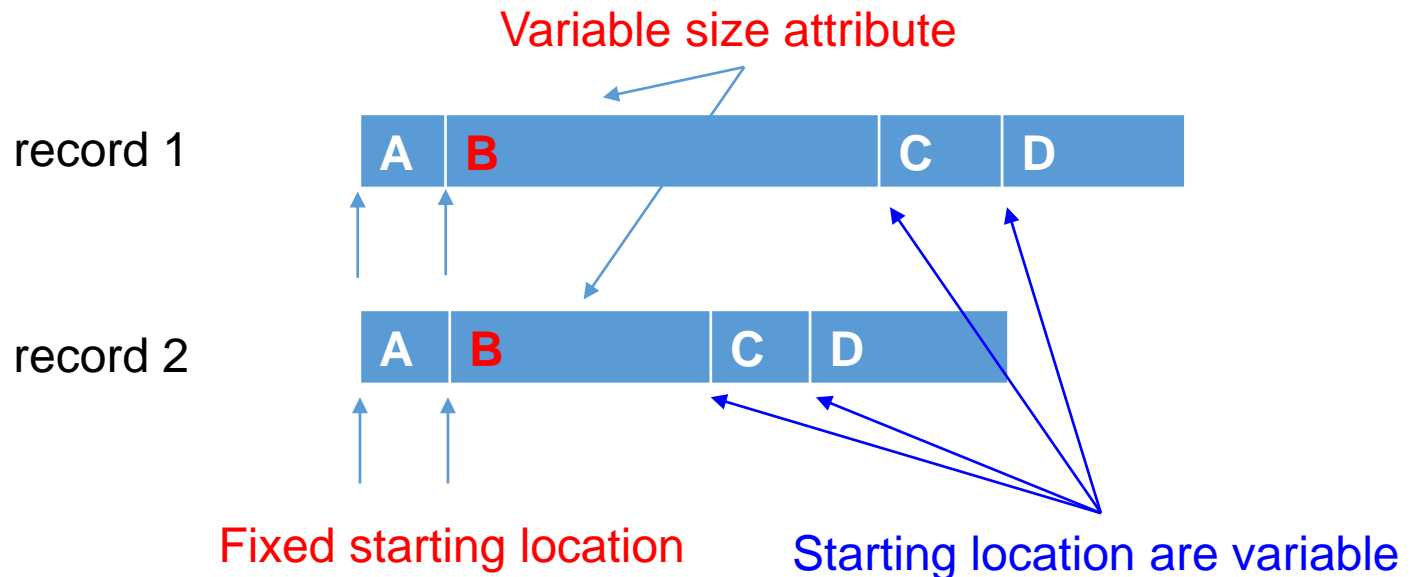
- How to store (data) record on disk
- How the DBMS locate records
- How to store record with variable size
- How to modify record

Need for Variable-Length Record

- Variable size record
 - String typed attribute
 - Repeating field
- Variable format
 - XML element
- Large size field

Problem of Storing Variable Size Record

Must store the **starting location** of **every field** in the record




How to store records with variable length field that uses minimal number of starting location information ?

Solution to Store Variable Size

Solution 1:

Store all fixed length fields first, then store the variable size attributes

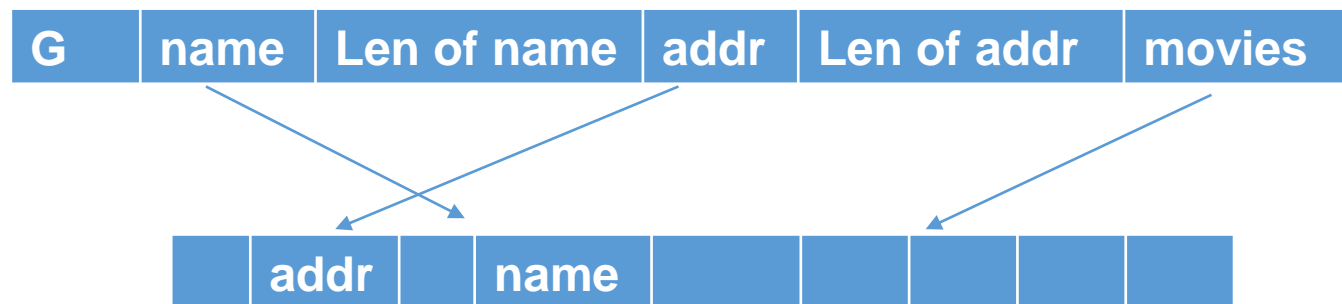


Solution to Store Variable Size

Solution 2:


Use **2 separate block**

- A fixed length record with data + pointer to variable length fields
- The variable length fields are stored separately



Advantages and Disadvantages of Using 2 Separate Block

- Advantage
 - Search more efficiently
 - Minimize the overhead in block header
 - **Easy to move record** within or among blocks
- Disadvantage
 - Increase the disk IO (need **2 disk access**)



Storing Variable Format Record

Problem:

- Record structure is unknown or variable

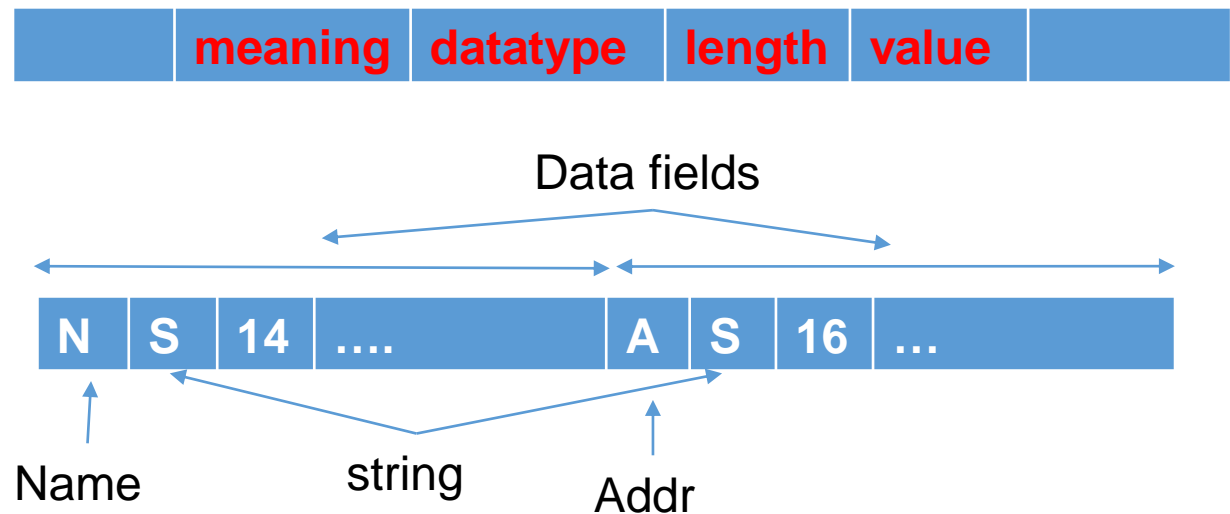
Solution:

- First define an encoding scheme

Example

N = Name
A = Addr

S = String
I = Integer
...




Spanned and Unspanned Record

- Problem
 - Packing whole record into blocks **wastes space**
 - Example
 - If record are just slightly larger than half a block, it waste space can approach 50%
- Solution
 - **Spanned record**

- Example

Unspanned

Spanned




Outline

- How to store (data) record on disk
- How the DBMS locate records
- How to store record with variable size
- How to modify record

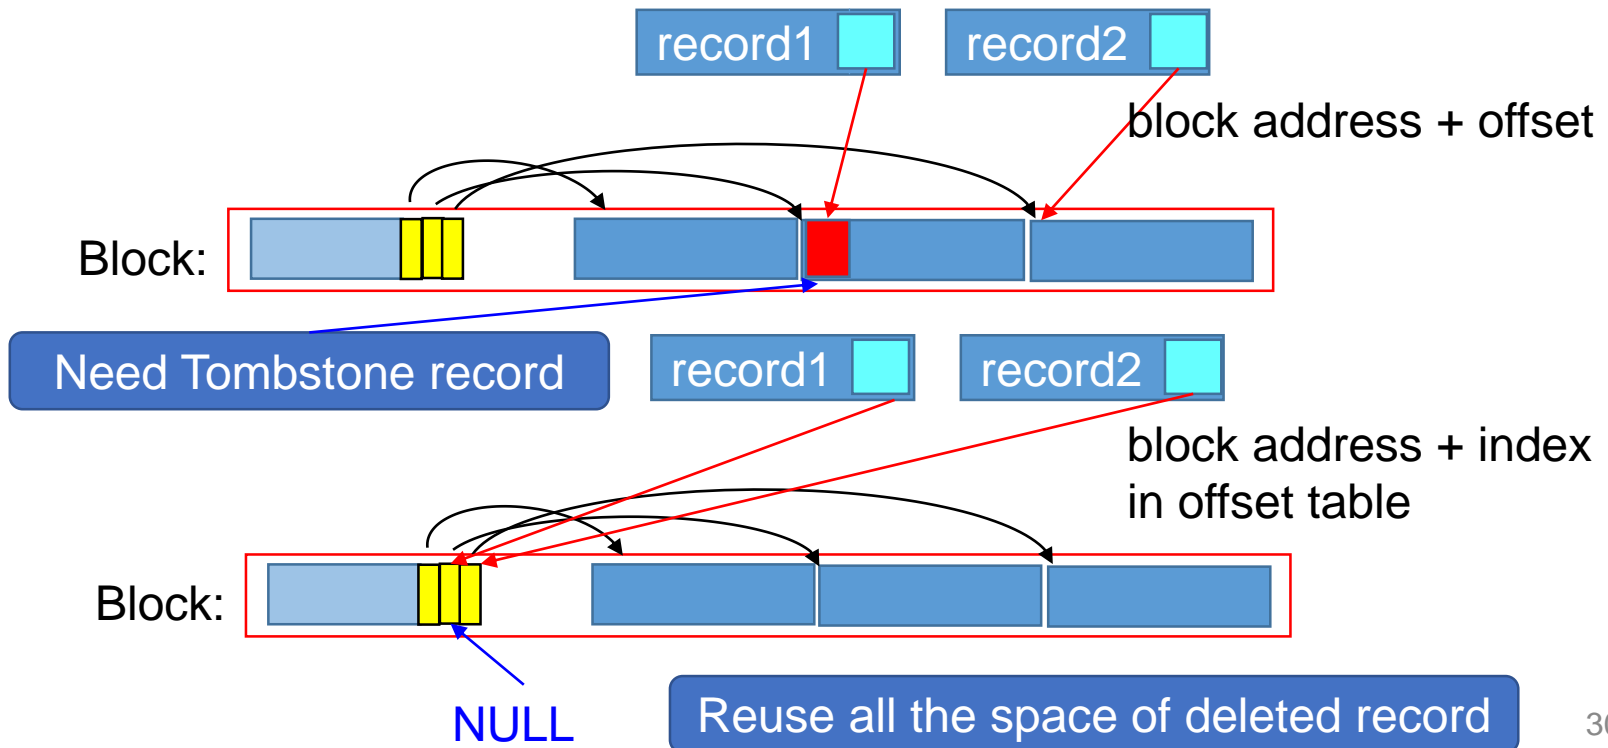
Insertion

- Easy case: records not ordered
 - Insert new record at end of file or in a deleted slot
- Hard case: records are ordered
 - Has room within the block
 - Slide records and insert
 - No room within the block
 - Find room outside the block
 - Nearby block
 - Overflow block



Deletion

- Need a **tombstone** record
- Set offset to **NULL**



Thank you!
Q & A