Microkernel

Jinyu Gu Sept. 27, 2020





Microkernel-based Harmony OS



INTRODUCTION

The Role of Operating Systems

- Manage hardware resources
 - E.g., CPU, Memory, I/O
- Provide ease-to-use interfaces to access resources
 - E.g., network sockets

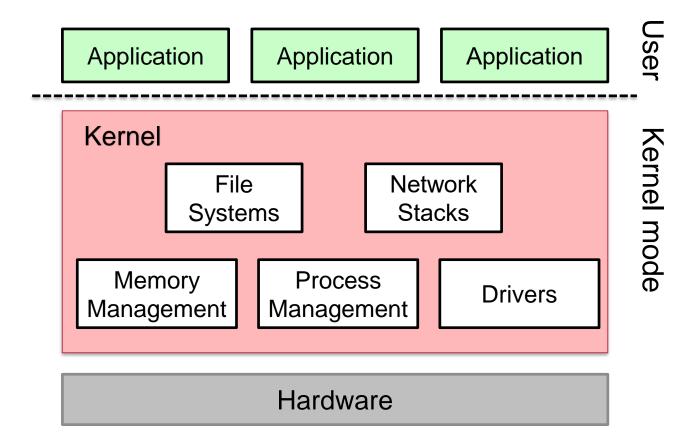
- Perform privileged/HW-specific operations
 - E.g., ring 0/3
- Provide separation and collaboration
 - E.g., users/processes isolation and communication

OS Architectures

Monolithic vs. Microkernel

- Monolithic Kernel
 - Linux
 - Windows
 - BSD

Monolithic Kernel Design



Disadvantages of Monolithic OSes

- System components run in privileged mode
 - No protection between system components
 - Security issues: e.g., faulty driver
 - No isolation between system components
 - Resilience issues: e.g., crash together
 - Big and inflexible
 - Difficult to replace system components
 - Difficult to understand and maintain

Monolithic vs. Microkernel

Monolithic Kernel

- Linux
- Windows
- MacOS

Microkernel

- QNX (car)
- Fuchsia/Zircon (Google)
- Fiasco (L4)
- Harmony (mobile, TV, car)
- RT-Thread Smart (IoT)

The Microkernel Vision

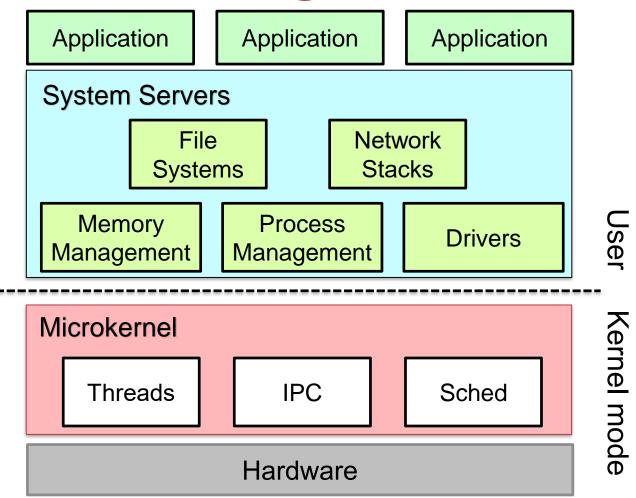
Minimal OS kernel

- Small code size
- Suitable for verification

System services in user-level servers

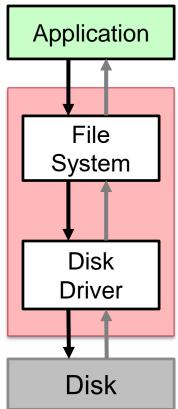
- Flexible and extensible
- Isolation between individual components
 - More resilient (no crash whole)
 - More secure (inter-component isolation)

Microkernel Design



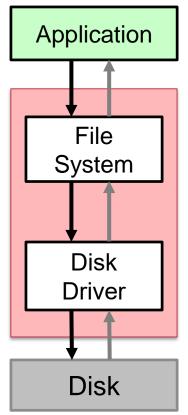
Interaction Becomes Different

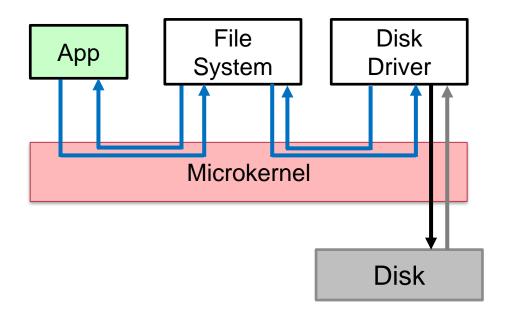
Example: file operations



Interaction Becomes Different

Example: file operations





From the View of Developers

Customizability

- Configure suitable servers (embedded, desktop ...)
- Remove unneeded servers

Enforce reasonable system design

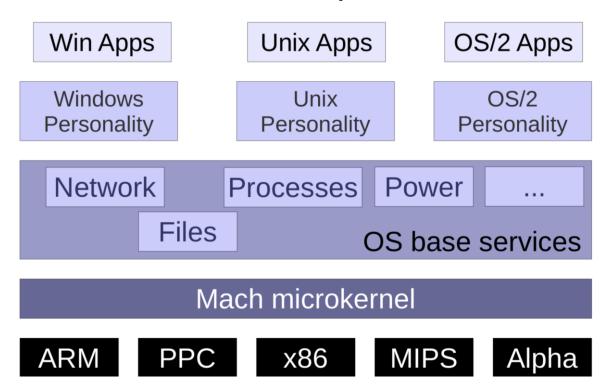
- Well-defined interfaces between components
- No access to components besides the interfaces
- Improved maintainability

Mach: The First Microkernel

- Mach developed at CMU, 1985 1994
 - Rick Rashid (former head of MS Research)
 - Avie Tevanian (former Apple CTO)
 - Brian Bershad (professor @ U. of Washington)
- Foundation for several real systems
 - IBM Workplace OS
 - Next OS --> Mac OS X

IBM Workplace OS

- Main goals: 1. multiple OS personalities
 - 2. run on multiple HW architectures



Lessons Learned

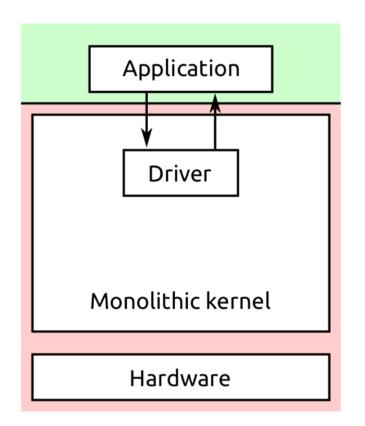
- Microkernel worked, but system atop it did not
- Underestimated difficulties in creating OS personalities
- subsystem protection/isolation
- small code size
- customizability:
 - Tailored memory management/scheduling/... algorithms
 - Adaptable to embedded/real-time/secure/... systems

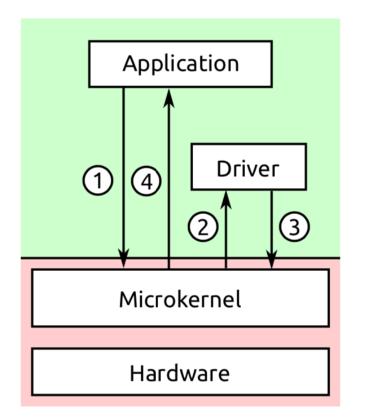
Review: Monolithic vs. Microkernel

- Flexibility and Customizable
 - Monolithic kernels are typically modular
- Maintainability and complexity
 - Monolithic kernels have layered architecture
- Robustness
 - Isolation & Trusted code size
- Performance
 - Application performance degraded

Monolithic vs. Microkernel: Syscalls

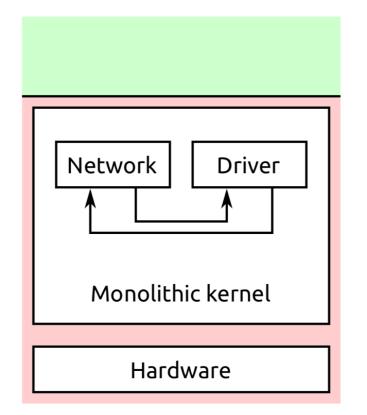
- Monolithic kernel: 2 kernel entries/exits
- Microkernel: 4 kernel entries/exits + 2 context switches

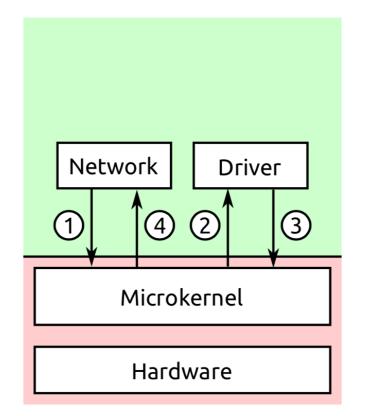




Calls Between Services

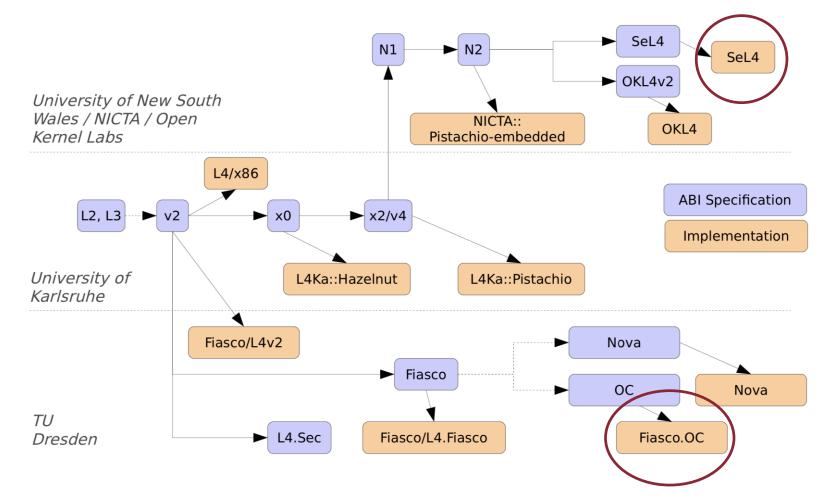
- Monolithic kernel: 2 function calls/returns
- Microkernel: 4 kernel entries/exits + 2 context switches





WHAT DO MICROKERNELS HAVE

L4 Family Microkernel



L4 Concepts

- Jochen Liedtke: "A microkernel does no real work"
 - Kernel only provides inevitable mechanisms
 - Kernel does not provide policies
 - Example: virtual memory management
 - Write page tables (mechanism)
 - How to write (policy)

What things are inevitable?

Case Study: L4/Fiasco.OC

" Everything is an object "

TaskAddress space

ThreadScheduling

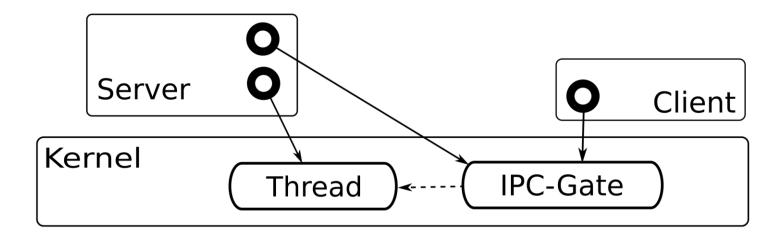
IRQ Asynchronous notification

IPC Gate Communication

One system call: invoke_object()

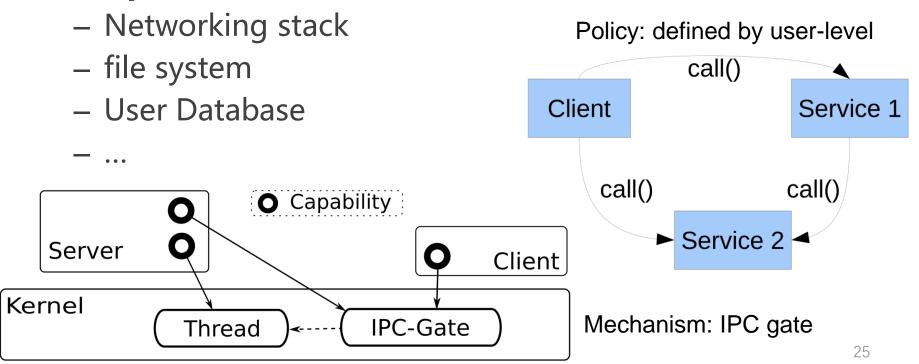
L4/Fiasco.OC: IPC Gate

- Generic communication object: IPC gate
 - Send message from sender to receiver
 - Used to provide services in user-level applications
- Example:



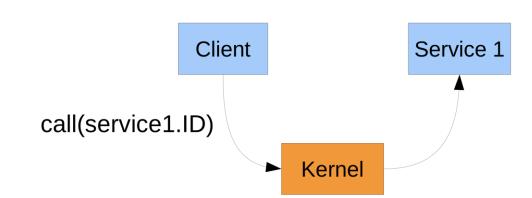
L4/Fiasco.OC: IPC Gate

 Everything above kernel is built using IPC gates that can provide different services



L4/Fiasco.OC: How to Call Objects?

- Recap: invoke_object()
- To call an object, we need an identifier:
 - Telephone number
 - Postal address
 - IP address
 - **–** ...
- Simple idea, right?
- ID is wrong? Kernel returns ENOTEXIST
- This scheme is insecure!
 - Client could simply guess IDs (brute-force)



L4/Fiasco.OC: Local Names for Objects

Using global object IDs

- Insecure (forgery)
- Inconvenient (programmers needs to know about partitioning in advance)

Solution in Fiasco.OC

- Task-local capability space as an indirection
- Object capability <u>required</u> to invoke the object

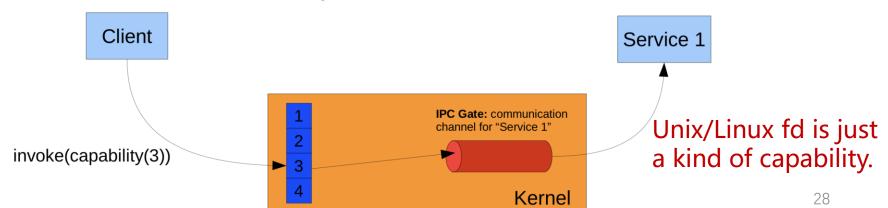
Per-task namespace

- Get object capabilities at start time (set by the farther task)
- Get object capabilities at runtime (grant by other tasks)

L4/Fiasco.OC: Object Capabilities

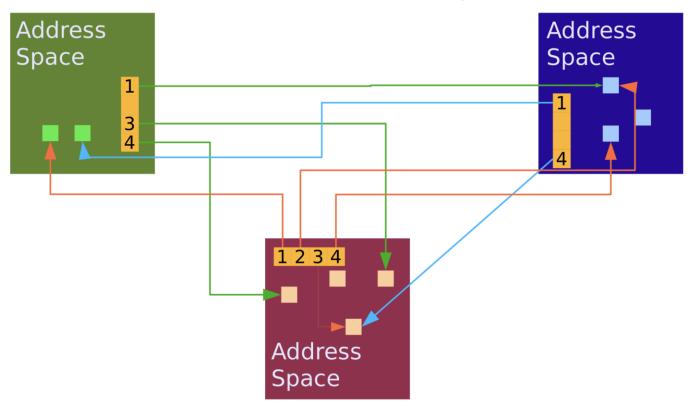
Capability:

- Reference to an object
- Protected by the kernel
 - Kernel knows all capability-object mappings
 - Managed as a per-process capability table
 - User processes only use indexes into the table



L4/Fiasco.OC: Communication

Indirection allows for security and flexibility



GO DEEPER INTO DETAILS

Abstraction: Thread

- An independent flow of control inside an address space
- Communicates with other threads using IPC
- Characterized by a set of registers and the thread state
- Dispatched by the kernel according to a defined schedule

Implementation in L4/Fiasco.OC

Execution Context

- Register state
- Address Space
- FPU state
- Continuation
- Msg buffer

Scheduling Context

- Priority
- Budget
- Remaining Budget
- Prev/next pointer

Thread Variants

Global Thread

- Needs a scheduling context, i.e., CPU time, to execute
- Causes exception on startup to let creator set register state

Local Thread (Passive Thread)

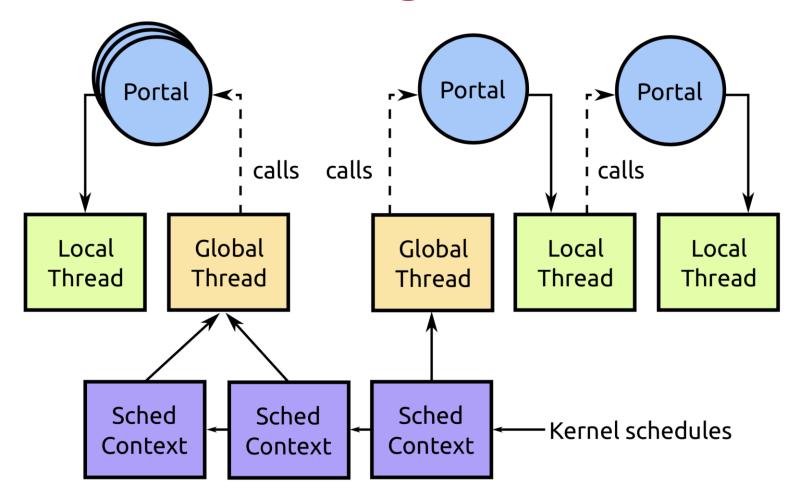
- Has no scheduling context
- Only used to handle portal calls
- Waits in the kernel until someone called an associated portal

Portals

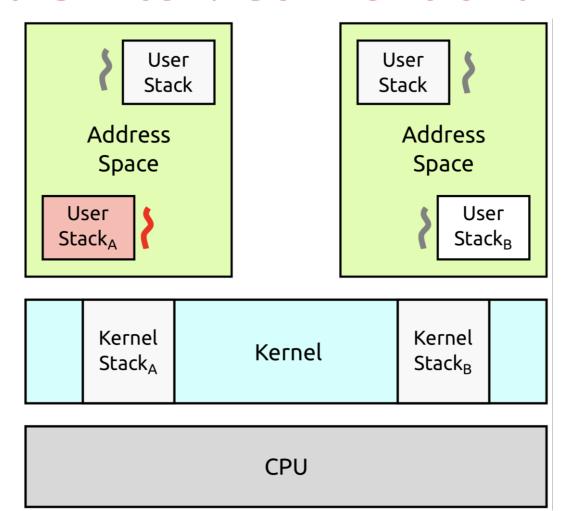
- A portal is an IPC gate mentioned before
 - Called via system call

- Associated with local threads
 - CPU time is donated from caller
 - Message is transferred from sender msg buffer to receiver msg buffer

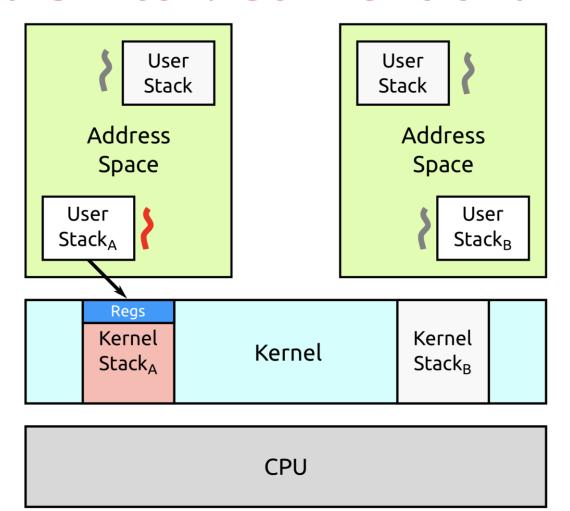
Thread Scheduling Overview



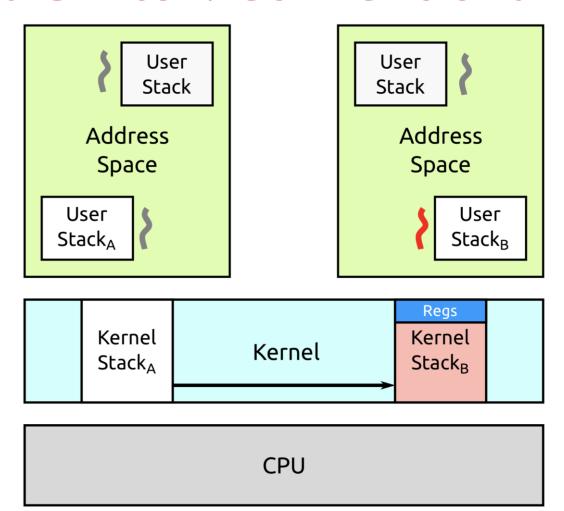
Thread Switch: Conventional



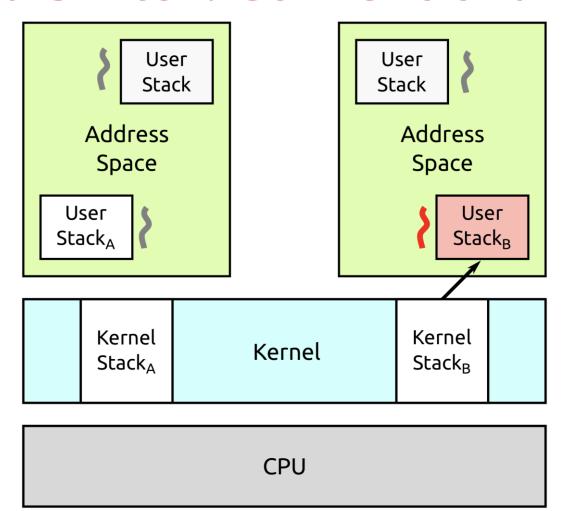
Thread Switch: Conventional

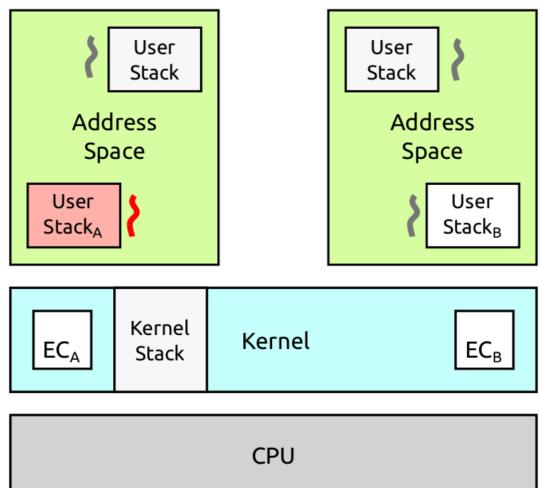


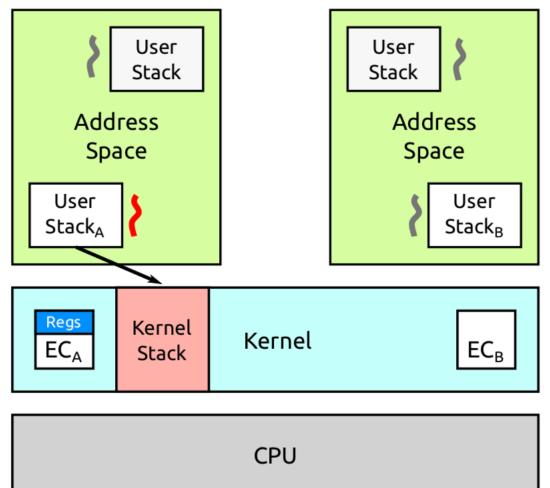
Thread Switch: Conventional

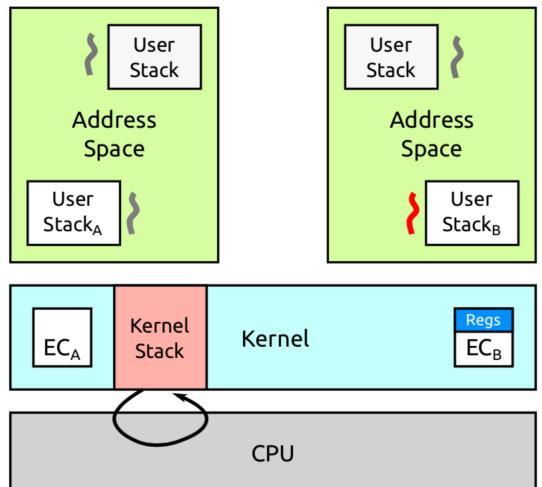


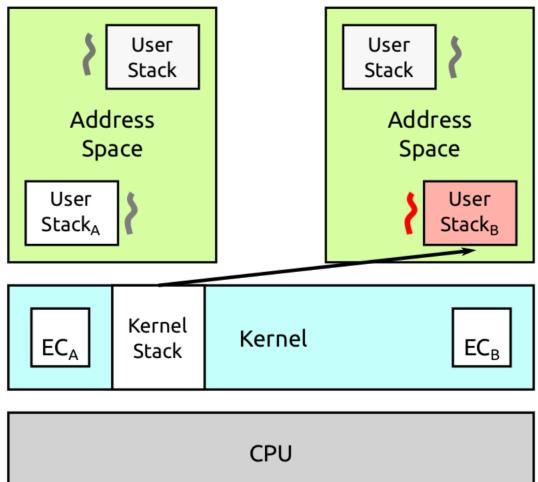
Thread Switch: Conventional











Improving IPC Performance

IPC is critical to the performance of the whole microkernel OS

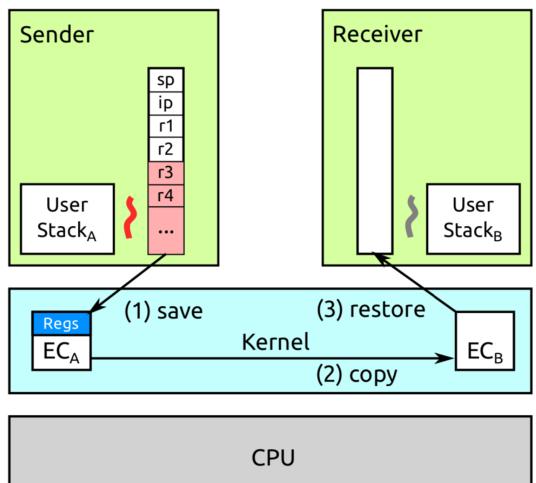
- First generation microkernel systems
 - Exhibited poor performance when compared to monolithic UNIX implementations
 - Particularly Mach, the best-known example
- Typical results:
 - Move OS services back into the kernel for performance
 - Mac OS X (Darwin): complete BSD kernel linked to Mach

IPC Optimization Techniques

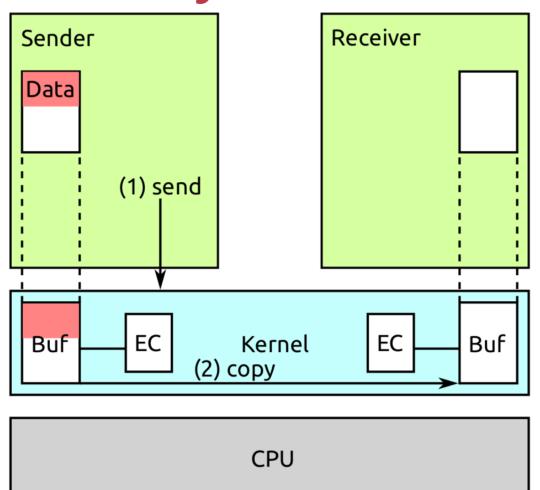
- Minimize number of system calls
 - Combined operations: Call, ReplyWait
 - Complex messages
 - Combines multiple messages into one operation
 - As many arguments as possible in registers
- Copy messages only once
- Reducing cache and TLB misses
- Avoid jumps and checks on fast path
- Direct process switch to receiver

•

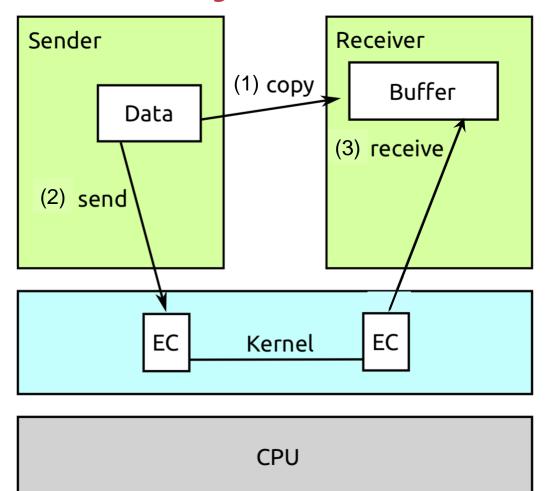
Register IPC



Kernel Memory IPC



Shared Memory IPC



Comparison

Register IPC

- + Very fast
- Amount of data limited to CPU registers

Kernel Memory IPC

- + Larger amount of data (still limited)
- More memory copies

Shared Memory IPC

- + Amount of data not limited
- + Less memory copies
- Security issues (e.g., TOUTOC, manipulation on the fly)

Sync and Async IPC

Synchronous

- Sender is blocked until receiver is ready
- Data and control transfer directly from sender to receiver

Asynchronous

- Data is transferred to temporary location
- Sender continues execution
- If receiver arrives, the data is transferred to him

Comparison

- Sync is typically simpler and faster (no buffering)
- Sync is less prone to DoS attacks (buffer memory)
- Async is typically more functional
- Async allows to do other work while waiting

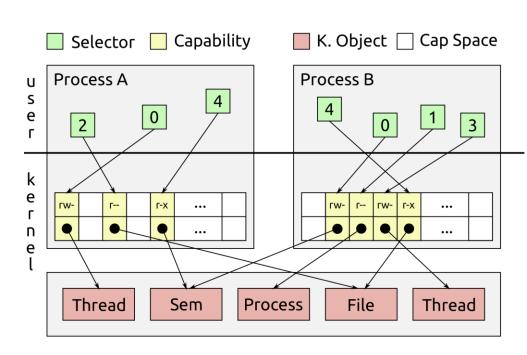
Usages of Sync and Async IPC

Synchronous IPC

- Exchange data
- Exchange capabilities
 - Copy/Move/Revoke

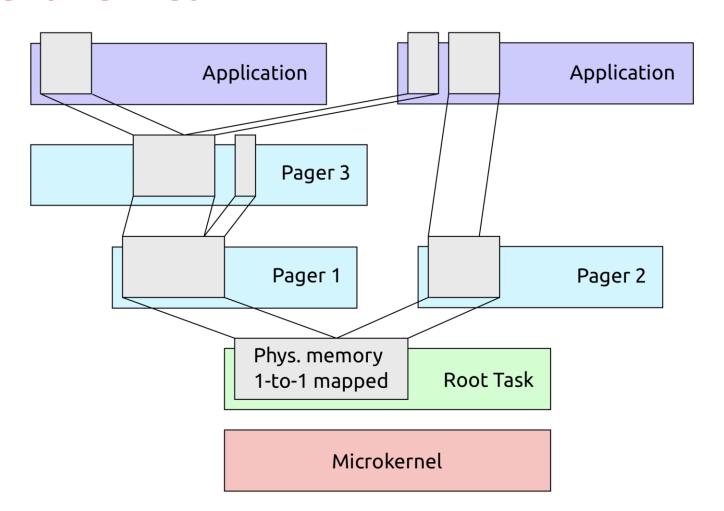
Asynchronous IPC

- Signaling
- Deliver interrupts to user space

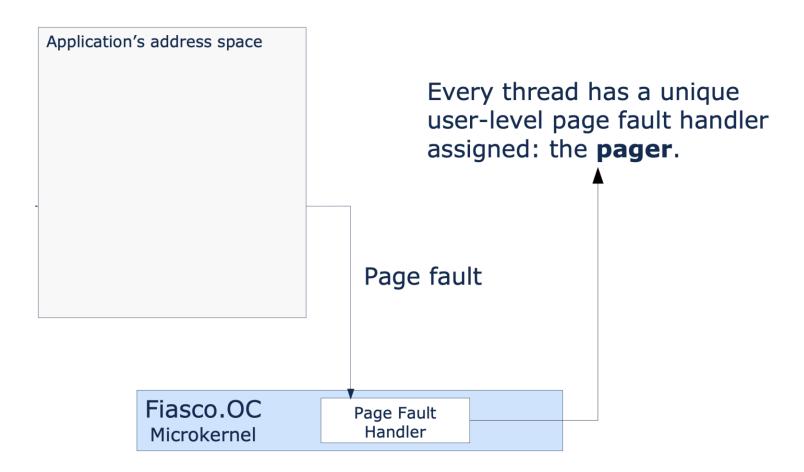


MEMORY MANAGEMENT


Hierarchical MM in L4




Page Fault Handling



Pager



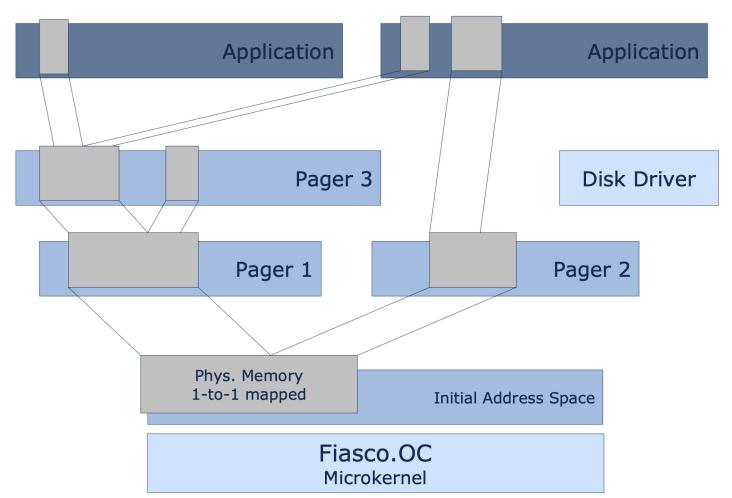
Memory Mapping

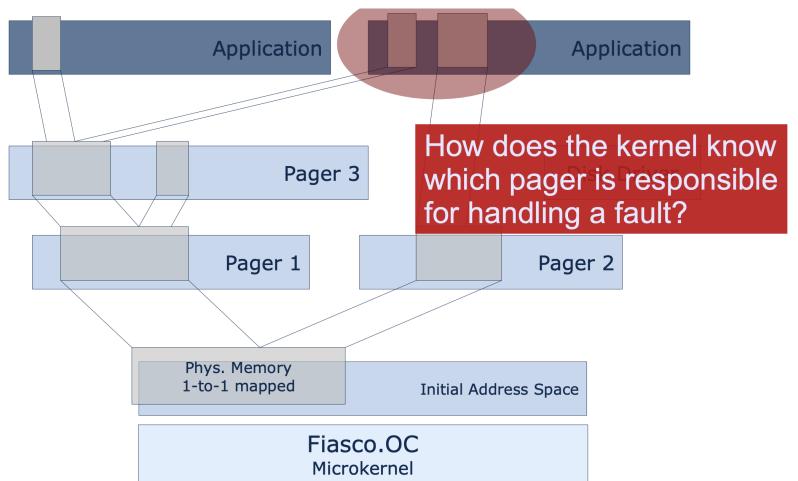


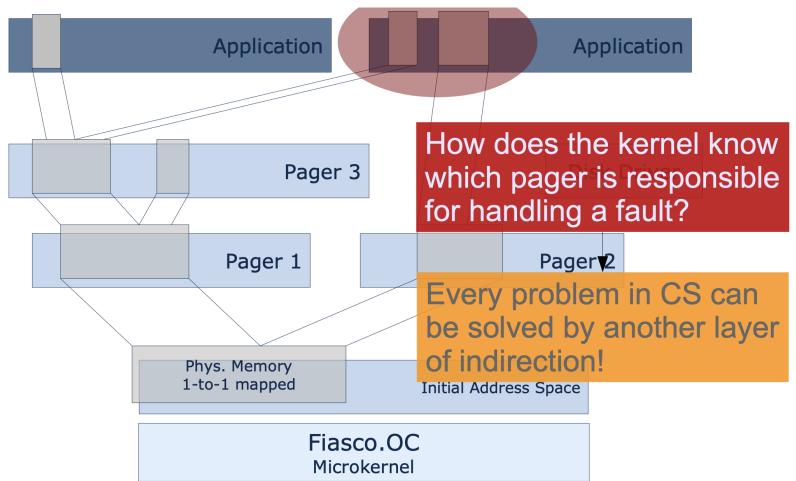
User-Level Pagers


- Chicken vs. egg: who resolves page faults for the Root Task?
 - Fiasco.OC initially maps all memory to it, so that it never raises page faults

- Complex vs. simple management policies
 - Solution: Pagers can be stacked into pager hierarchies



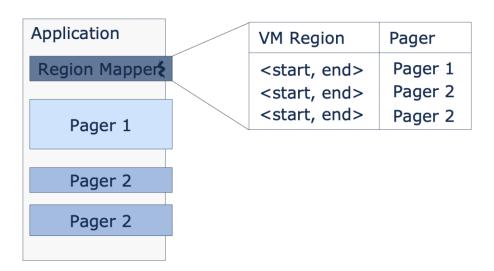


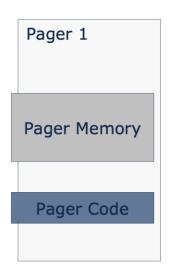




Region Mapper

- Every task has a single pager thread, the region mapper (RM)
 - RM is the first thread within a task in Fiasco.OC
 - RM maintains a region map





Fiasco.OC

Microkernel

Page-Fault
Handler

Region Mapper and Page Faults

<u>Step-1:</u> All threads of a task have the RM assigned as their pager

Step-2: Fiasco.OC redirects all page faults to the RM

Step-3: RM then uses synchronous IPC calls to obtain real memory mappings from the external pager responsible for a region

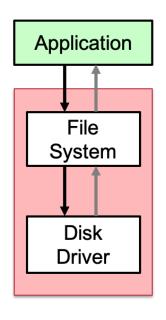
Disadvantages of Pager

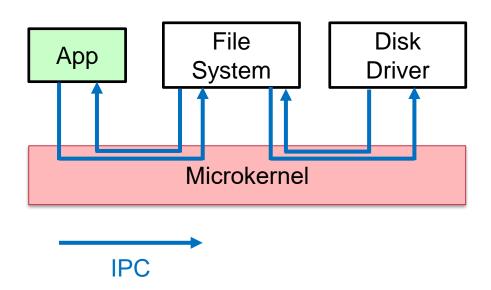
- Poor performance
- Complex developing efforts
 - Is that really necessary?

- Hybrid solution
 - E.g., Google Zircon

SOME RESEARCHES

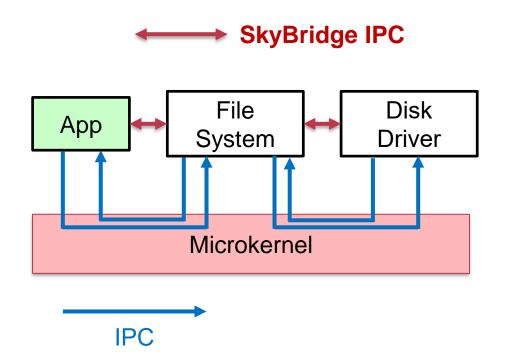
Tradeoff: Performance and Isolation





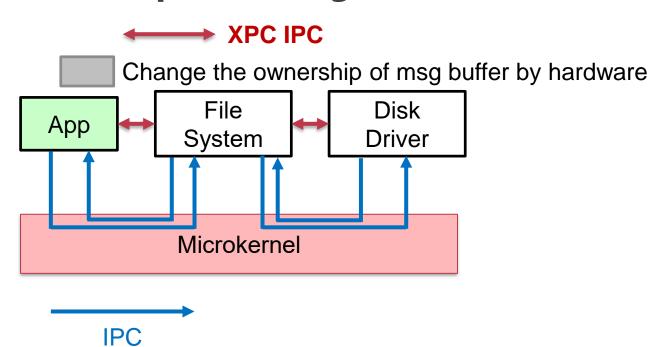
SkyBridge: Make IPC Faster

EuroSys19: No kernel involvements during IPCs




XPC: Make IPC Faster

ISCA19: No kernel involvements during IPCs;
 Zero-copies during IPCs



UnderBridge: Make IPC Faster

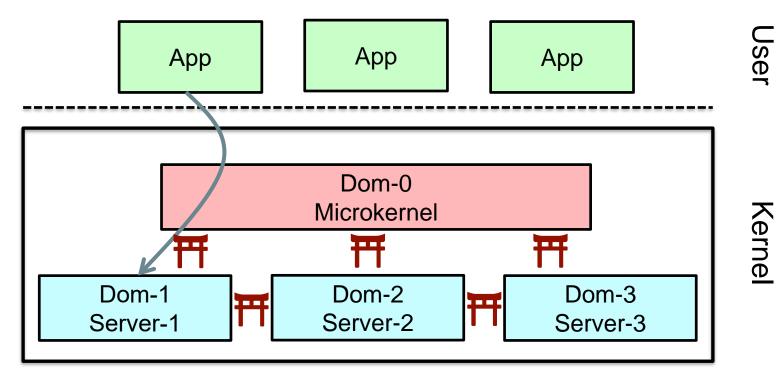
 ATC20: No kernel involvements during IPCs by vertically sinking system servers



Microkernel

UnderBridge: Make IPC Faster

Build execution domains in the kernel space



CuriOS: Improving Reliability through Operating System Structure

Francis M. David & Ellick M. Chan &

Jeffrey C. Carlyle & Roy H. Campbell

UIUC

Is Microkernel Reliable enough?

Benefits: Better Isolation? Fault Recovery?

Problem

 Blindly restarting a service results in the state loss and affects all clients that were using the service

Goal

 Use lightweight distribution, isolation and persistence of OS service state to mitigate the problem of state loss during a restart

Fault Tolerance in Microkernel

- Microkernel designs componentize the OS into servers managed by a minimal kernel
 - Inter-component error propagation is significantly reduced
 - Better fault isolation
- Recovery from a microkernel server failure is typically attempted by restarting it
 - Minix3: a printer driver √
 - Minix3: a file system ×

Simply Restarting is not Enough

 Obstacle: many OS services maintain states related to clients

Take FS as an example

 Reads and writes to existing open files cannot be completed because the restarted server cannot recognize the file handles (i.e., client state lost)

How to Solve the Problem

Design-Choice-1:

- Clients take care of service restarts and state loss
- Increased code complexity ×

Design-Choice-2:

 Provide some form of persistence to the clientrelated state information √

Simply Persisting is not Enough

Intra-component error propagation

 An error that occurs in an OS server can potentially corrupt any part of its state before being detected

A significant limitation of traditional microkernel systems

Server State Management

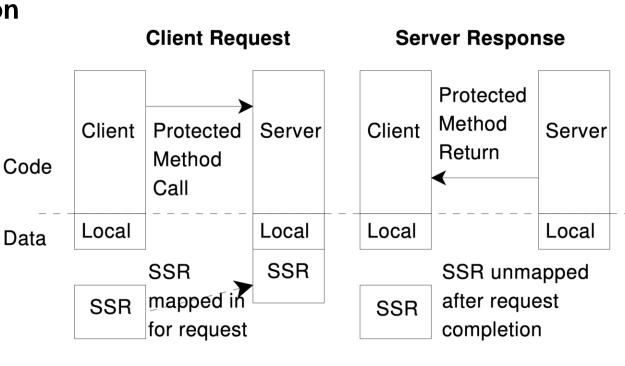
SSR: Server State Region

Stores an OS server's clientrelated information

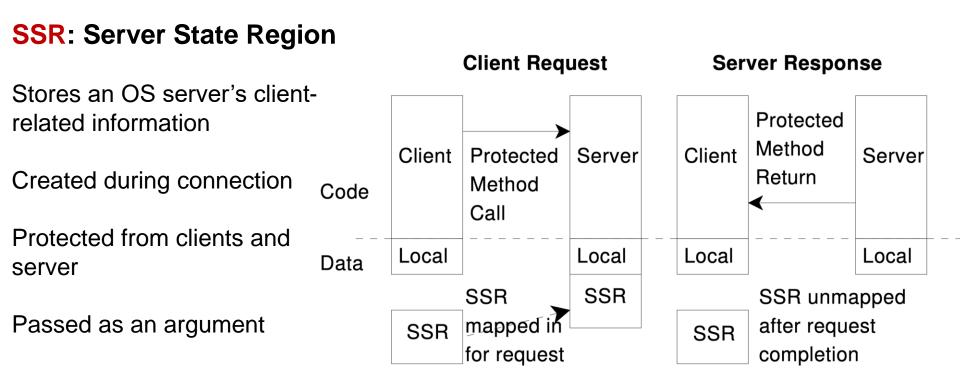
Created during connection

Protected from clients and server

Passed as an argument



Server State Management



Persistence and isolation of client-related state

SSR Manager

- Register a new server
- Bind a client to a server
 - SSR creation
- Undo a client-server binding
 - SSR deletion
- Enumerate all the SSRs associated with a server

OS Service Construction

Three types of servers:

- Stateless (e.g., some drivers)
- No require collective information about all of its clients
 - E.g., a server that provides pseudo random numbers based on a per-client seed
- Require knowledge about all of its clients in order to service a request
 - E.g., a scheduler
 - redundantly cache this information locally to process

Server Recovery

Restarting only is enough for the first two types

- Type-3 servers need to register a recovery routine
 - Query the SSR Manager to obtain all associated SSRs
 - Re-create the internal state of the restarted server
 - Heuristics consistency check is required since some
 SSR may be corrupted (and most clients can survive)

CuriOS

- Make efforts on recovering stateful OS servers
- Fault tolerance
 - Isolation is achieved by the microkernel design and is better since the isolation among SSRs
 - Recovery is achieved by persistence and isolation of client-related state
- There are still many open questions

Conclusion

- Small code size is necessary but not enough
- Capability and IPC are two core abstractions
- Open questions
 - Performance
 - Security
 - Compatibility
 - Reliablity



