

Transaction

Xingda Wei

Nov. 10, 2020



Transaction: key pillar for online trading system





INTRODUCTION

Transaction(TX) is a user specified “program”

Transfer \$100 from A to B

```
A_bal = READ(A)
If (A_bal > 100) {
    B_bal = READ(B)
    B_bal += 100
    A_bal -= 100
    WRITE(A, A_bal)
    WRITE(B, B_bal)
}
```

Report sum of money

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

Transaction executions are concurrent

Transfer \$100 from A to B

```
A_bal = READ(A)
If (A_bal > 100) {
  B_bal = READ(B)
  B_bal += 100
  A_bal -= 100
  WRITE(A, A_bal)
  WRITE(B, B_bal)
}
```

Report sum of money


```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

Executed from different threads,
even different machines (distributed).

Storage
system

Why concurrency ? To achieve high performance

- Moore' s law is dead, single core cannot become faster
- Scaling database to many machines
 - For bigger capacity & better fault tolerance



Wildly execute transactions causes inconsistency

Transfer \$100 from A to B

```
A_bal = READ(A)
If (A_bal >= 100) {
  B_bal = READ(B)
  B_bal += 100
  A_bal -= 100
  WRITE(A, A_bal)
  WRITE(B, B_bal)
}
```

Report sum of money

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

A = 100
B = 100

Transfer B_bal += 100

Report

Timeline

Wildly execute transactions causes inconsistency

Transfer \$100 from A to B

```
A_bal = READ(A)
If (A_bal >= 100) {
  B_bal = READ(B)
  B_bal += 100
  A_bal -= 100
  WRITE(A, A_bal)
  WRITE(B, B_bal)
}
```

Report sum of money

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

A = 100
B = 100

Transfer B_bal += 100

Report B_bal = 200

Timeline

Wildly execute transactions causes inconsistency

Transfer \$100 from A to B

```
A_bal = READ(A)
If (A_bal >= 100) {
  B_bal = READ(B)
  B_bal += 100
  A_bal -= 100
  WRITE(A, A_bal)
  WRITE(B, B_bal)
}
```

Report sum of money

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

A = 100
B = 100

Transfer B_bal += 100

Report

B_bal = 200 A_bal = 100

Timeline

Wildly execute transactions causes inconsistency

Transfer \$100 from A to B

```
A_bal = READ(A)
If (A_bal >= 100) {
  B_bal = READ(B)
  B_bal += 100
  A_bal -= 100
  WRITE(A, A_bal)
  WRITE(B, B_bal)
}
```

Report sum of money

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

A = 100
B = 100

Transfer

B_bal += 100

A_bal -= 100

Report

B_bal = 200 A_bal = 100

Timeline



Solutions?

1. **Leave it to application programmers**
 - Application programmers are responsible for locking data items
2. **System performs automatic concurrency control**
 - Concurrent transactions execute as if serially

Solutions?

1. Leave it to application programmers

- Application programmers are responsible for locking data items

2. System performs automatic concurrency control

- Concurrent transactions execute as if seriall
- Even when the environment is concurrent

System guarantees transaction is ACID

- **A (Atomicity)**
 - All-or-nothing w.r.t. failures
- **C (Consistency)**
 - Transactions maintain any internal storage state invariants
- **I (Isolation)**
 - Concurrently executing transactions do not interfere
- **D (Durability)**
 - Effect of transactions survive failures

What is ACID in details ?

T1: Transfer \$100 from A to B

```
A_bal = READ(A)
If (A_bal > 100) {
  B_bal = READ(B)
  B_bal += 100
  A_bal -= 100
  WRITE(A, A_bal)
  WRITE(B, B_bal)
}
```

T2: Report sum of money

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

- **Atomicity: T1 completes or nothing.**
 - E.g., If $B += 100 \Rightarrow A -= 100$

What is ACID in details ?

T1: Transfer \$100 from A to B

```
A_bal = READ(A)
If (A_bal > 100) {
    B_bal = READ(B)
    B_bal += 100
    A_bal -= 100
    WRITE(A, A_bal)
    WRITE(B, B_bal)
}
```

T2: Report sum of money

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

- **Consistency: guarantees application semantics**
 - E.g., (A_bal + B_bal) is not changed

What is ACID in details ?

T1: Transfer \$100 from A to B

```
A_bal = READ(A)
If (A_bal > 100) {
  B_bal = READ(B)
  B_bal += 100
  A_bal -= 100
  WRITE(A, A_bal)
  WRITE(B, B_bal)
}
```

T2: Report sum of money

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

- **Isolation: T1 & T2 isolated from each other**
 - E.g., T2 does not see T1's intermediate result

What is ACID in details ?

T1: Transfer \$100 from A to B

```
A_bal = READ(A)
If (A_bal > 100) {
  B_bal = READ(B)
  B_bal += 100
  A_bal -= 100
  WRITE(A, A_bal)
  WRITE(B, B_bal)
}
```

T2: Report sum of money

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

- **Durability: storage survives compute failures**
 - E.g., changes of T1 persists after client/server failures

System guarantees transaction is ACID

- **A (Atomicity)**
 - All-or-nothing w.r.t. failures
- **C (Consistency)**
 - Transactions maintain any internal storage state invariants
- **I (Isolation)**
 - Concurrently executing transactions do not interfere
- **D (Durability)**
 - Effect of transactions survive failures

Ideal isolation semantic: serializability

- **Definition:** execution of a set of transactions is equivalent to some serial order
 - Two executions are *equivalent* if they have the same effect on database and produce same output.

Why serializability?

- Suppose each TX transfers system from a consistent state to another consistent state

CS \rightarrow TX \rightarrow CS

Then the final state of executing **all** TXs is consistent

CS \rightarrow TX₁ \rightarrow TX₂ ... TX_n \rightarrow CS

Examples

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

Serializable? R(A),R(B),R(A),R(B),C W(A),W(B),C

Yes, equivalent serial schedule: R(A),R(B),C,R(A),R(B),W(A),W(B),C

Serializable? R(A),R(B), W(A), R(A),R(B),C, W(B),C

Realize serializability

- **Using the standard technique in concurrent programming**
 - Locking-based approach
- **Strawman solution 1:**
 - Grab global lock before transaction starts
 - Release global lock after transaction commits
- **Strawman solution 2:**
 - Grab lock on item X before reading/writing X
 - Release lock on X after reading/writing X

Strawman 2

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

Possible with strawman 2? (short-duration locks)

R(A),R(B), W(A), R(A),R(B),C, W(B),C

Strawman 2

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

Possible with strawman 2? (short-duration locks)

R(A), R(B), W(A), R(A), R(B), C, W(B), C

Read an uncommitted value

Strawman 2

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

Possible with strawman 2? (short-duration locks)

R(A),R(B), W(A), R(A),R(B),C, W(B),C

Locks on writes should be held
till end of transaction

Read an uncommitted value

More Strawman

- **Strawman 3**
 - Grab lock on item X before read/writing X
 - Release locks at the end of transaction

Strawman 2


```
lock(A)
A_bal = READ(A)
lock(B)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
unlock(A)
unlock(B)
```

```
lock(A)
A_bal = READ(A)
lock(B)
B_bal = READ(B)
Print(A_bal+B_bal)
unlock(A)
unlock(B)
```

Possible? R(A),R(B), W(A), **R(A),R(B),C** W(B),C

Strawman 3 is also called 2PL

- **2 phase locking (2PL) guarantees serializability**
 - A growing phase in which the transaction is acquiring locks
 - A shrinking phase in which locks are released



2PL: deadlocks?

```
lock(A)
A_bal = READ(A)
lock(B)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
unlock(A)
unlock(B)
```

```
lock(B)
B_bal = READ(B)
lock(A)
A_bal = READ(A)
Print(A_bal+B_bal)
unlock(A)
unlock(B)
```

2PL: deadlocks?

```
lock(A)
A_bal = READ(A)
lock(B)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
unlock(A)
unlock(B)
```

```
lock(B)
B_bal = READ(B)
lock(A)
A_bal = READ(A)
Print(A_bal+B_bal)
unlock(A)
unlock(B)
```



L(A)




Timeline

2PL: deadlocks?

```
lock(A)
A_bal = READ(A)
lock(B)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
unlock(A)
unlock(B)
```

```
lock(B)
B_bal = READ(B)
lock(A)
A_bal = READ(A)
Print(A_bal+B_bal)
unlock(A)
unlock(B)
```




Timeline

2PL: deadlocks?

```
lock(A)
A_bal = READ(A)
lock(B)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
unlock(A)
unlock(B)
```

```
lock(B)
B_bal = READ(B)
lock(A)
A_bal = READ(A)
Print(A_bal+B_bal)
unlock(A)
unlock(B)
```




Timeline

2PL: deadlocks?

```
lock(A)
A_bal = READ(A)
lock(B)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
unlock(A)
unlock(B)
```

```
lock(B)
B_bal = READ(B)
lock(A)
A_bal = READ(A)
Print(A_bal+B_bal)
unlock(A)
unlock(B)
```




Timeline

2PL: deadlocks?

```
lock(A)
A_bal = READ(A)
lock(B)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
unlock(A)
unlock(B)
```

```
lock(B)
B_bal = READ(B)
lock(A)
A_bal = READ(A)
Print(A_bal+B_bal)
unlock(A)
unlock(B)
```



Disadvantages of locking-based approach

- **Need to detect deadlocks**
 - Distributed implementation needs distributed deadlock detection (bad)
- **Read-only transactions can block update transactions**
 - Big performance hit if there are long running read-only transactions

Optimistic concurrency control

Reads do not acquire locks

- **Use validation-based scheme**

No deadlock

- **Use speculative execution to ensure a deterministic locking order**

Optimistic concurrency control

- **OCC execute transaction speculatively**
 - Execution phase (execute TX with (near no) concurrency control)
 - Validation phase (verify the result of execution phase)
 - Commit phase
- **Only lock the write-set in the validation phase**

OCC: an example

Original transaction

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

OCC: an example

Original transaction

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

Executed using OCC

```
Cache[A_bal] = READ(A)
Cache[B_bal] = READ(B)
Cache[B_bal] += 100
Cache[A_bal] -= 100
```

Execution phase

OCC: an example

Original transaction

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

Executed using OCC

```
Cache[A_bal] = READ(A)
Cache[B_bal] = READ(B)
Cache[B_bal] += 100
Cache[A_bal] -= 100
```

Execution phase

```
Lock(A)
Lock(B)
```

Abort if A_bal or B_bal has changed

Validation phase

OCC: an example

Original transaction

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

Executed using OCC

```
Cache[A_bal] = READ(A)
Cache[B_bal] = READ(B)
Cache[B_bal] += 100
Cache[A_bal] -= 100
```

Execution phase

```
Lock(A)
Lock(B)
Abort if A_bal or B_bal has changed
```

Validation phase

```
A_bal = Cache[A_bal]
B_bal = Cache[B_bal]
// unlock A and B
```

Commit phase

OCC: an example

- **Why no deadlock?**

- after execution phase,

We explore all the

Involved data

- **Thus, we can assign**

a deterministic lock order

Executed using OCC

```
Cache[A_bal] = READ(A)
Cache[B_bal] = READ(B)
Cache[B_bal] += 100
Cache[A_bal] -= 100
```

Execution phase

```
Lock(A)
Lock(B)
Abort if A_bal or B_bal has changed
```

Validation phase

```
A_bal = Cache[A_bal]
B_bal = Cache[B_bal]
// unlock A and B
```

Commit phase

OCC: an example

Original transaction

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

Executed using OCC

```
Cache[A_bal] = READ(A)
Cache[B_bal] = READ(B)
```

Execution phase

Abort if A_bal or B_bal changed

Validation phase

```
Rrint(cache[A_bal] + cache[B_bal])
```

Commit phase


OCC: how to handle aborts?

- **Only retry the execution**
 - Excessive retries cause low performances
- **Cause live lock if there are many contentions !**

Recap

- **We have talked about transaction' s ACID property**
- **We have presented two classic protocols to ensure (ACI) transaction**
- **Most protocols nowadays are those two protocol' s variations !**

However, Serializability is Costly under Contention



Recap: why 2PL & OCC fails under contention ?

- **2PL**
 - Contended transactions block each other
 - Read-only transactions block writers
 - Dead lock (hard in a distributed setting)
- **OCC**
 - Contended transaction frequently retries

Observations

- **Read-only transaction dominates (real) workloads**
 - e.g., 99.8% requests to Facebook TAO are reads[1]
- **Read-only transaction are long-running[2]**
 - Blocking or retry are both unacceptable

[1] TAO [ATC'13]

[2] Challenges to adopting stronger consistency at scale. [HOTOS'15]

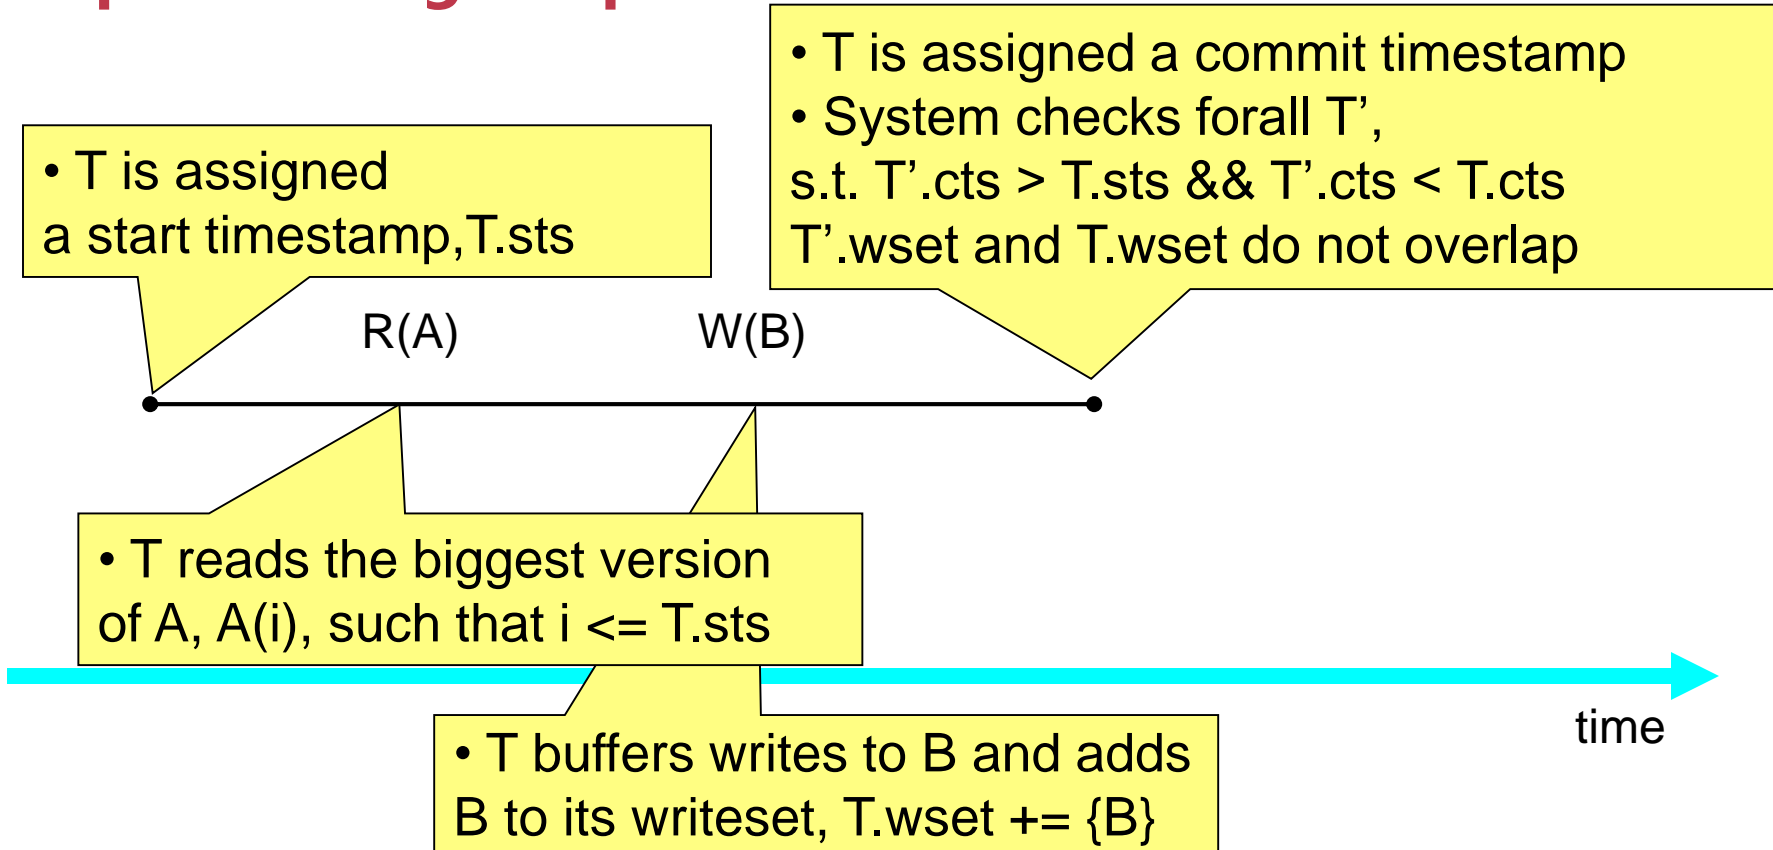
Multi-version protocols

- **Each data item is associated with multiple versions**
- **Multi-version transactions:**
 - Reads choose the appropriate version from a consistent view (no locking and retry!)

Snapshot Isolation (not serializable)

- **A popular multi-version concurrency control scheme**
- **A transaction:**
 - reads a “snapshot” of database image
 - Can commit only if there are no write-write conflict

Implementing Snapshot Isolation



Snapshot Isolation: a concrete example

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

A0 = 100
B0 = 100

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```




Timeline

Snapshot Isolation: a concrete example

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

A0 = 100
B0 = 100

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```




Snapshot Isolation: a concrete example

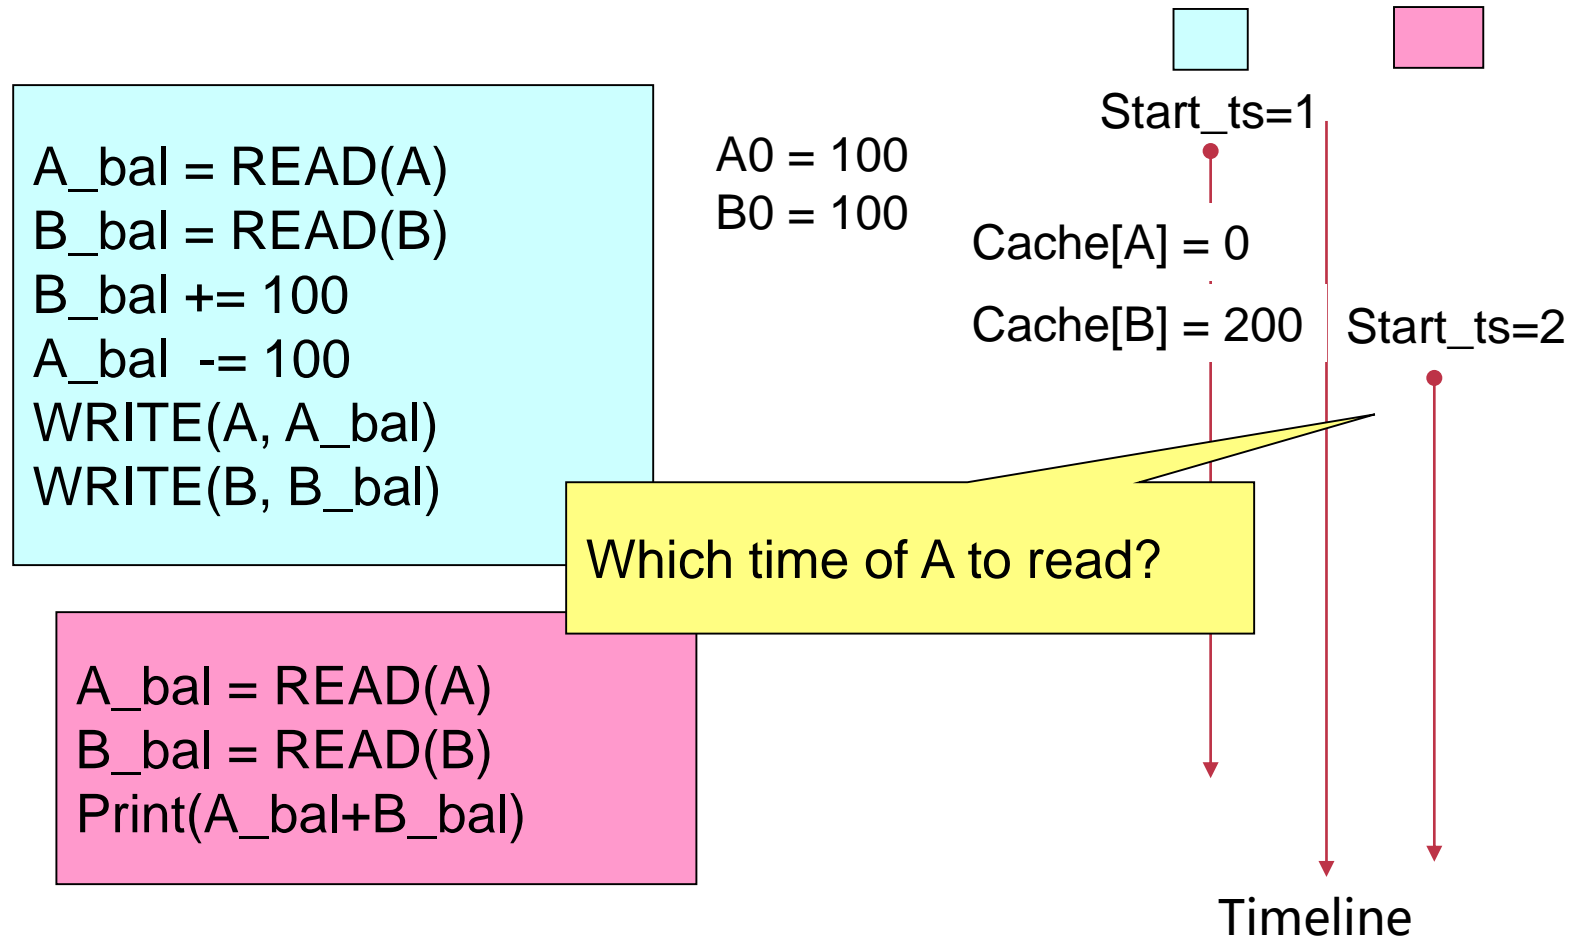
```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

A0 = 100
B0 = 100



Snapshot Isolation: a concrete example

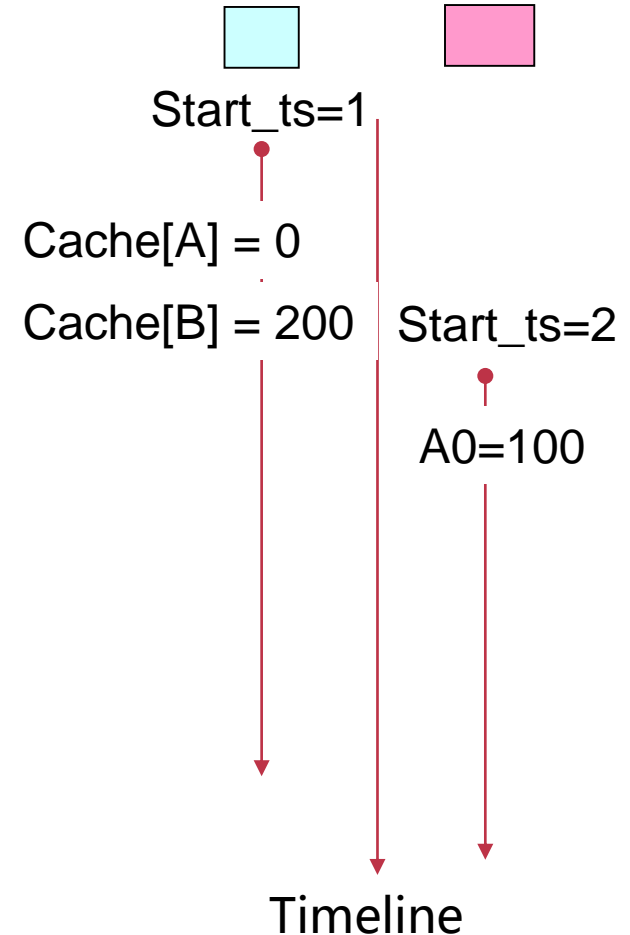


Snapshot Isolation: a concrete example

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

A0 = 100
B0 = 100

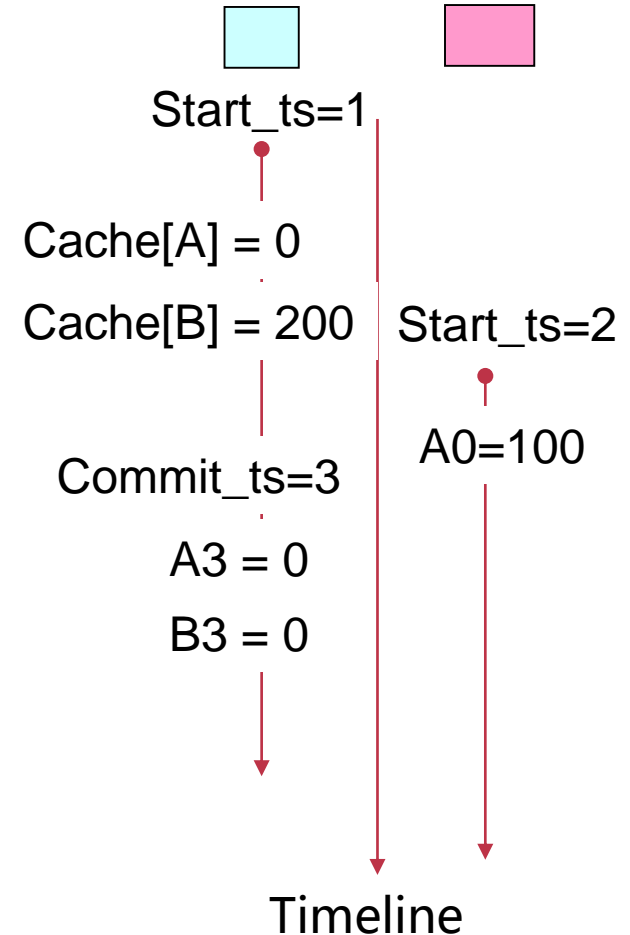


Snapshot Isolation: a concrete example

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

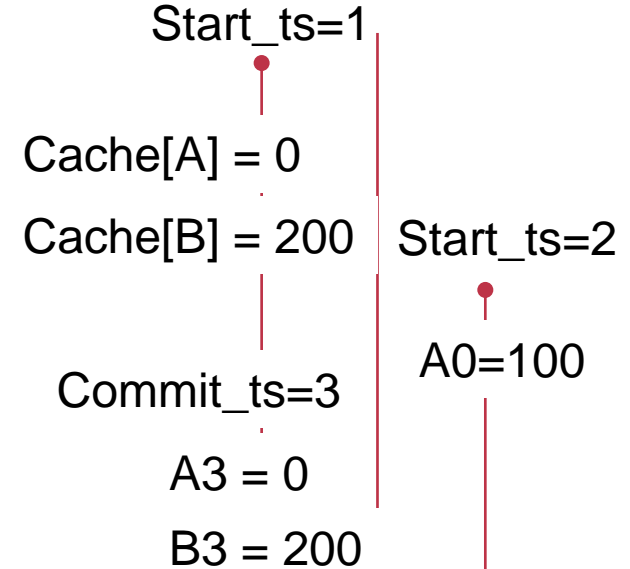
A0 = 100
B0 = 100



Snapshot Isolation: a concrete example

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

A0 = 100
B0 = 100



```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

Which time of B to read?

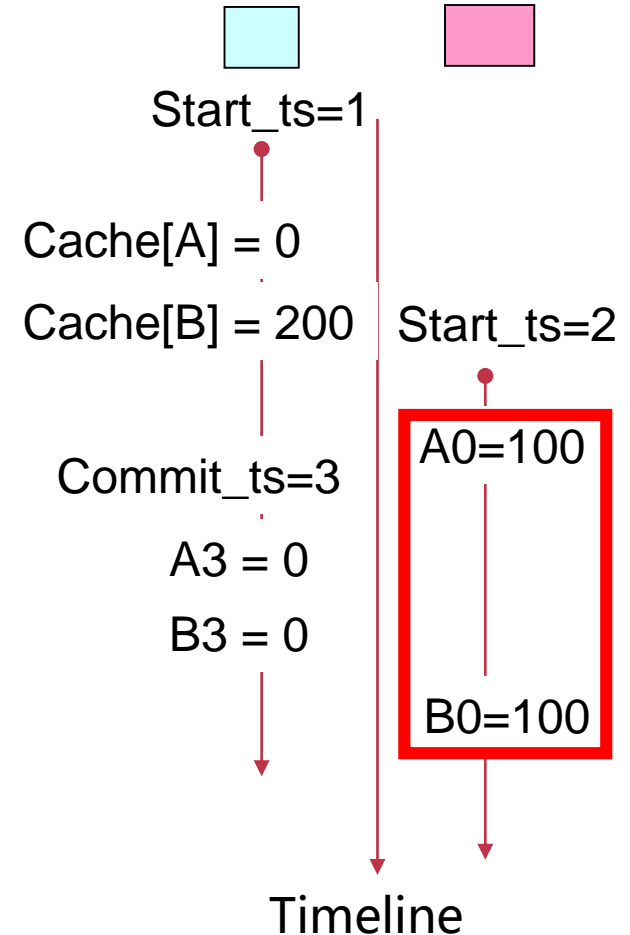
Timeline

Snapshot Isolation: a concrete example

```
A_bal = READ(A)
B_bal = READ(B)
B_bal += 100
A_bal -= 100
WRITE(A, A_bal)
WRITE(B, B_bal)
```

```
A_bal = READ(A)
B_bal = READ(B)
Print(A_bal+B_bal)
```

A0 = 100
B0 = 100



Snapshot isolation < serializability

- The write-skew problem
- Please refer to “A critique of ansi sql isolation levels [SIGMOD’ 95” for more details



FROM ALGORITHMS TO SYSTEMS


Towards modern transaction systems

- **2PL, OCC and SI are decades old**
 - Yet nowadays they are still the golden standard of concurrency control !
- **How can we continuously improve transaction' s performance?**
 - Through better system designs

Transactions are slow in (traditional) databases

Only **4% of wall-clock time** spent on useful data processing, while the rest is occupied with **buffer pools, locking, latching, recovery**.¹

-- Michael Stonebraker




•50% of time spent on ensuring ACI

ACI = algorithms + implementations

- **Algorithms (protocols):**
 - 2PL, OCC, etc
- **Implementations**
 - Locking methods (read-write locks), deadlock detection mechanism, etc
 - How to place data (e.g., in-memory, disk, NVM)
 - Different system design choices matters

Case study: Hstore[VLDB' 08]

- **Put all database in memory**
 - Eliminate buffer pool overhead
- **But other parts need redesign**
 - How to do fault tolerance?




•Totally eliminates buffer pool overhead

A better algorithm does yield better performance

- **Multi-version incurs better performance**
 - E.g., No abort read-only TX, etc
- **But incurs additional scalability bottleneck**
 - A global counter is used to assign timestamps
 - Global timestamp allocation is slow, even on a single modern multi-core machine !

Scalability bottleneck in multi-core server

```
txn_commit()  
{  
    // prepare commit  
    // [...]  
    commit_tid =  
atomic_fetch_and_add(&global_tid);  
    // quickly serialize transactions a la Hekaton  
}
```



Recap: ACI

- **A (Atomicity)**
 - All-or-nothing w.r.t. failures
- **C (Consistency)**
 - Transactions maintain any internal storage state invariants
- **I (Isolation)**
 - Concurrently executing transactions do not interfere
- **D (Durability)**
 - Effect of transactions survive failures



FAULT TOLERANCE

Durability vs. Availability

- **Durability**
 - Storage states survives even if there is failures
 - E.g., TX' modifications flushed to non-volatile storage
- **Availability (for distributed systems)**
 - System survives even if there is failure)

Logging for durability

- **Logging rules**
 - Write log record to disk before modifying persistent state
- **Recovery**
 - After system reboots, drain the logs and recover system states to a consistent state

Challenges of logging for in-memory databases

- **How to recovery logs to a consistent state ?**
 - If T2 reads T1' s modifications, T2' log must be recovered after T1' s log
- **How to hide disk latency ?**
 - Typically several orders of magnitude longer than TX' s lifetime!
- **Please refer to Silo[SOSP' 13]**

Logging + Replication for Availability

- **How to maintain system' s function when there is machine failures ?**
 - Replication: data are replicated to multiple machines; If one machine fails, redirect the workloads to its replica
- **Challenge: How to maintain ACID when there is replications ?**
 - Use logging to sync replica' s states by a consensus protocol, e.g., Paxos

Case study: Spanner[OSDI' 10]

- A distributed (geo-replicated) databases
- Use two-phase locking for concurrency control
- Uses **paxos** to sync replication' s state
 - So that each replica group function as a single machine with no failure

Case study: FaRM[SOSP' 15]

- **Drawback of paxos:**
 - At least two network-roundtrips for a single request (e.g., lock, read/write)
- **FaRM uses primary backup replication with vertical paxos**
 - OCC for concurrency control
 - Only the commit phase need to ship log to replicas

Summary

- **Protocols to ensure transaction' s ACI property**
- **How system implementation affect transaction' s performance**

Thanks!

