

Data Privacy

Protecting Your Data from Anyone

Yubin Xia

Data Privacy

- Data can be used everywhere
 - Risk management
 - Medicine
 - Recommended system
 - ...
- Data can be stolen easily
 - Everyone who uses your data can steal it

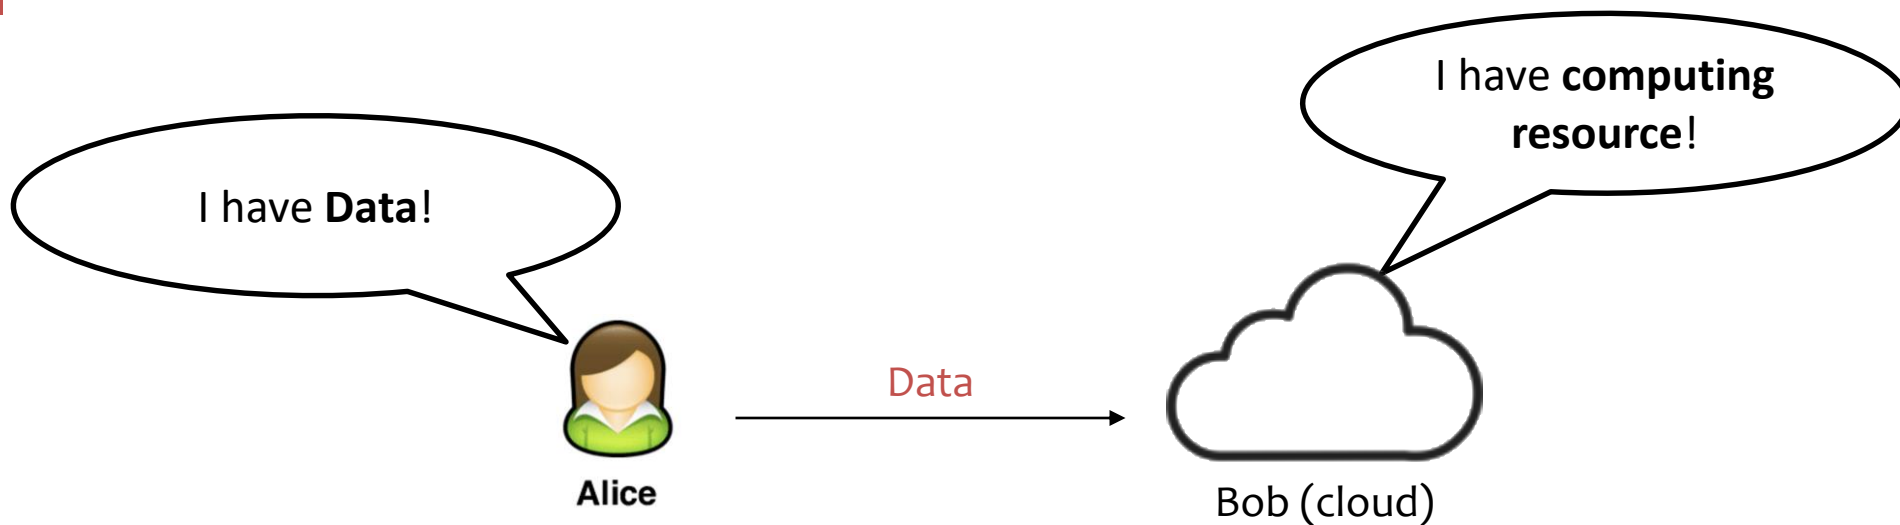
Data Privacy

- What's the target of data privacy system?
 - Allow data **to be used**, and
 - **Protect data from being stolen**
- What will be introduced?
 - Basic data privacy method
 - ZKP, OT, HE, sMPC, TEE, DP
 - Systems which try to enforce data privacy

Trusted Execution Environment

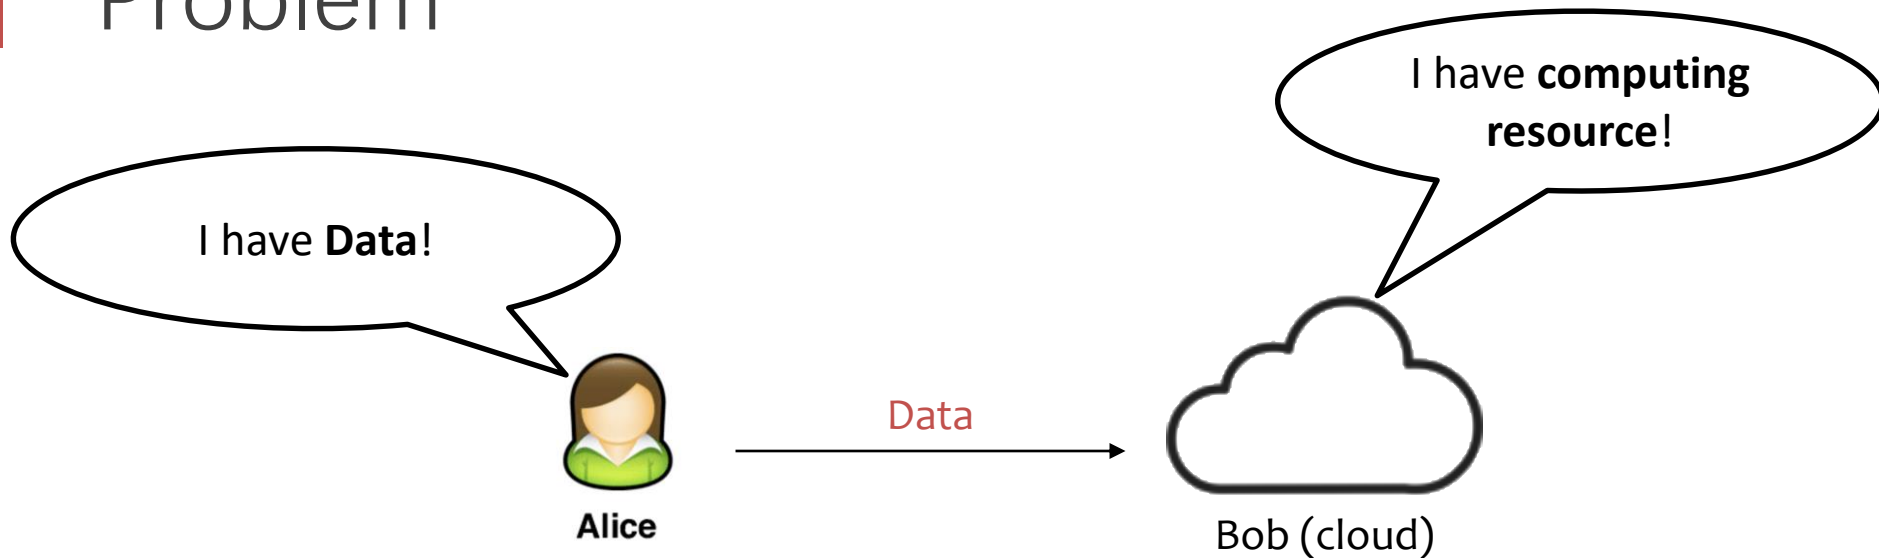
TEE

Problem



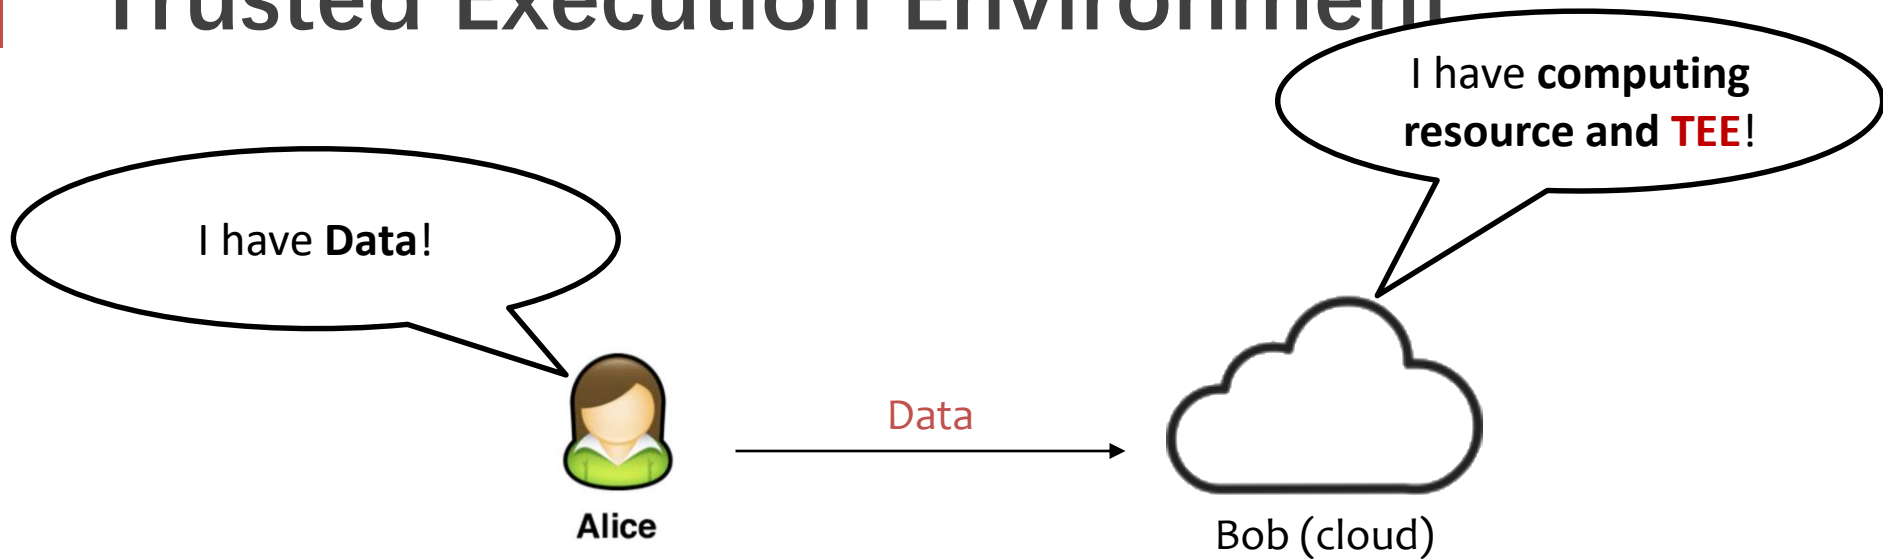
- Alice wants to ask Bob (e.g., a cloud) to perform calculation on her data
- Naïve method: Sending Data to Bob

Problem



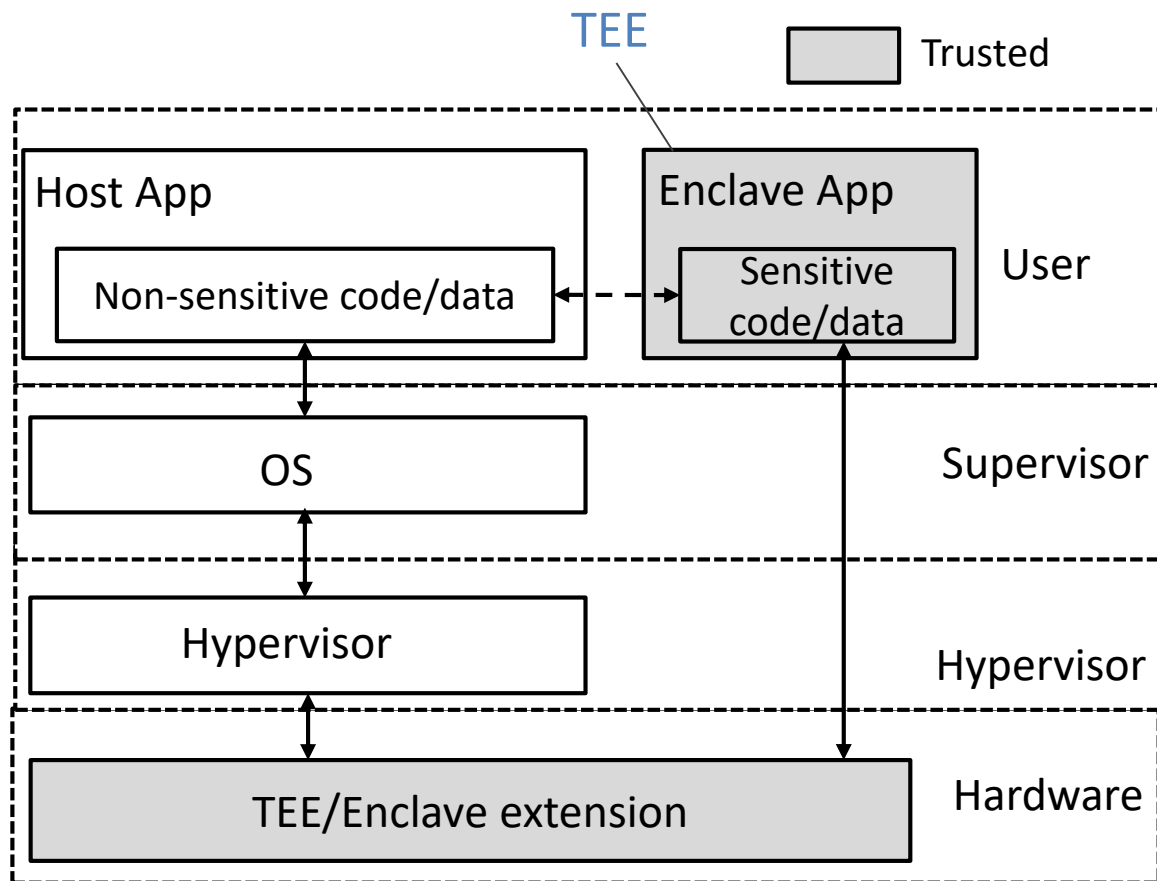
- Alice wants to ask Bob (e.g., a cloud) to perform calculation on her data
- ~~Naïve method: Sending Data to Bob~~

Trusted Execution Environment



- Alice wants to ask Bob (e.g., a cloud) to perform calculation on her data
- ~~Naïve method: Sending Data to Bob~~
- **Bob cloud construct a TEE**

What is TEE



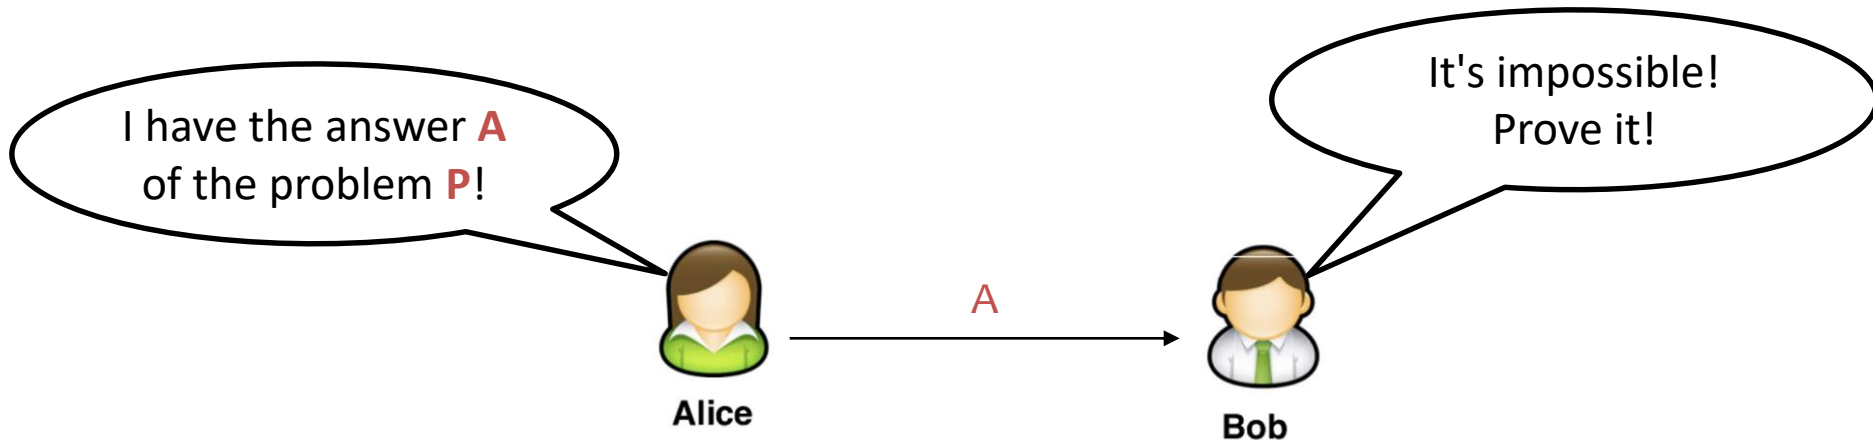
Different TEEs

- Software TEE
 - VM-based TEE
 - Same privilege protection
- ARM TrustZone
- Intel SGX
- AMD SME/SEV
- SANCTUM

Zero-Knowledge Proof

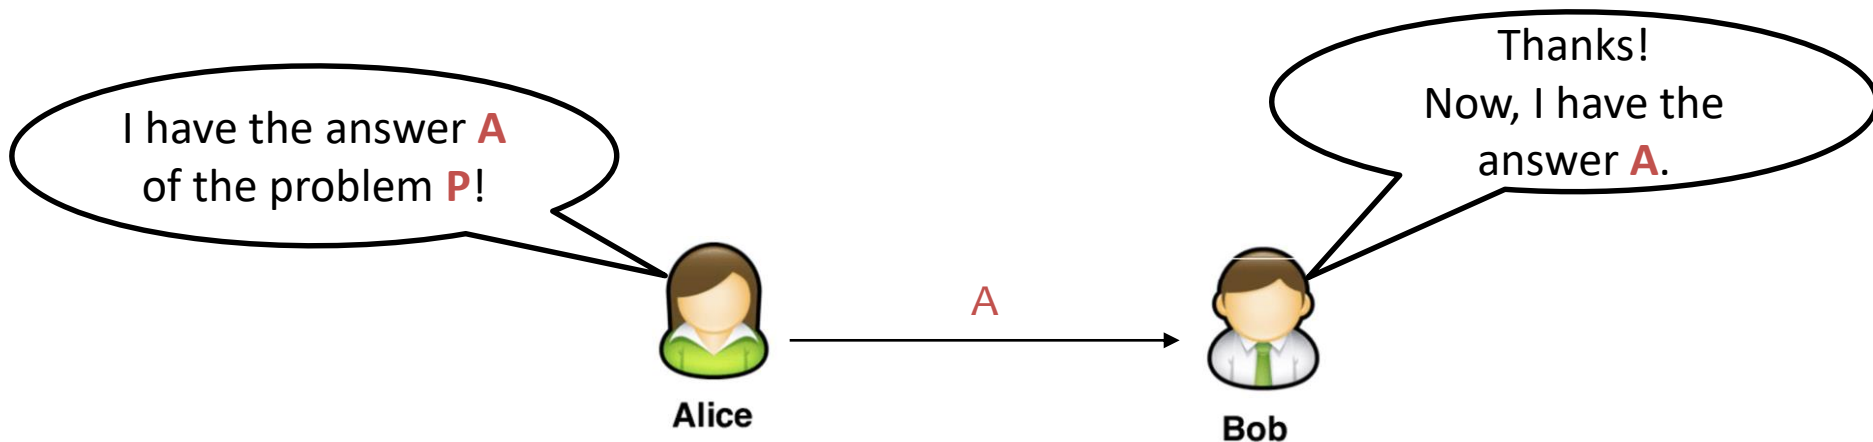
ZKP

Problem



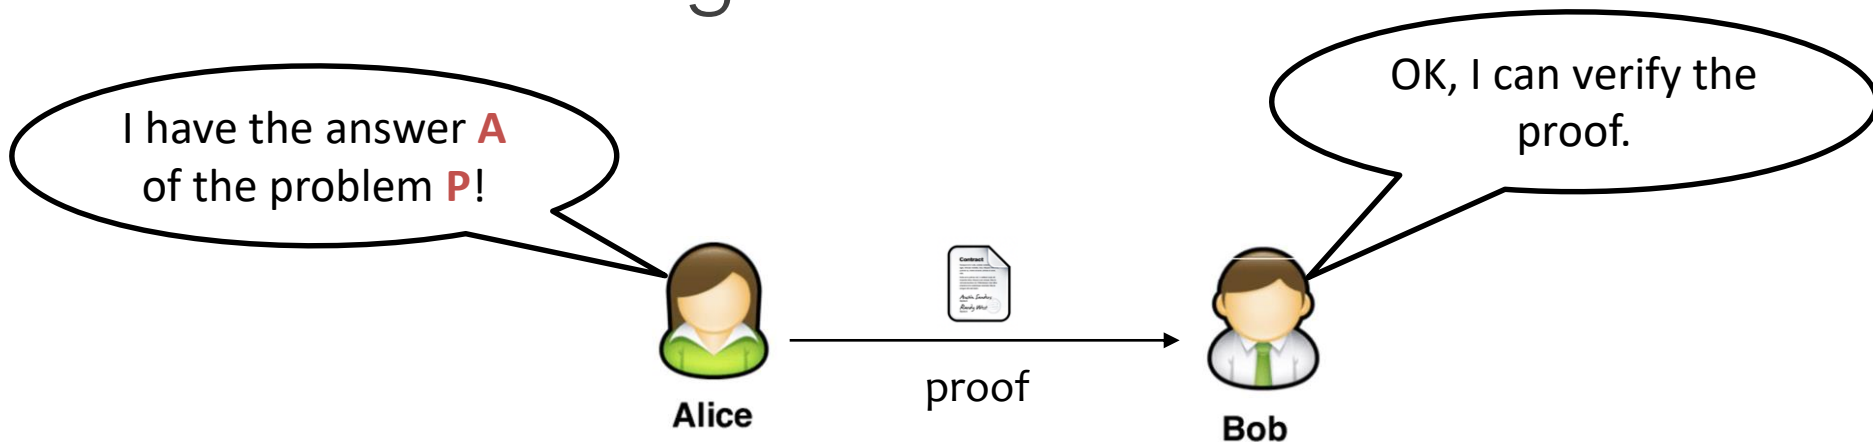
- Alice tries to **prove** to Bob that she **has the answer of a difficult problem** (e.g., a NP problem)
- Naïve method: Sending A to Bob

Problem



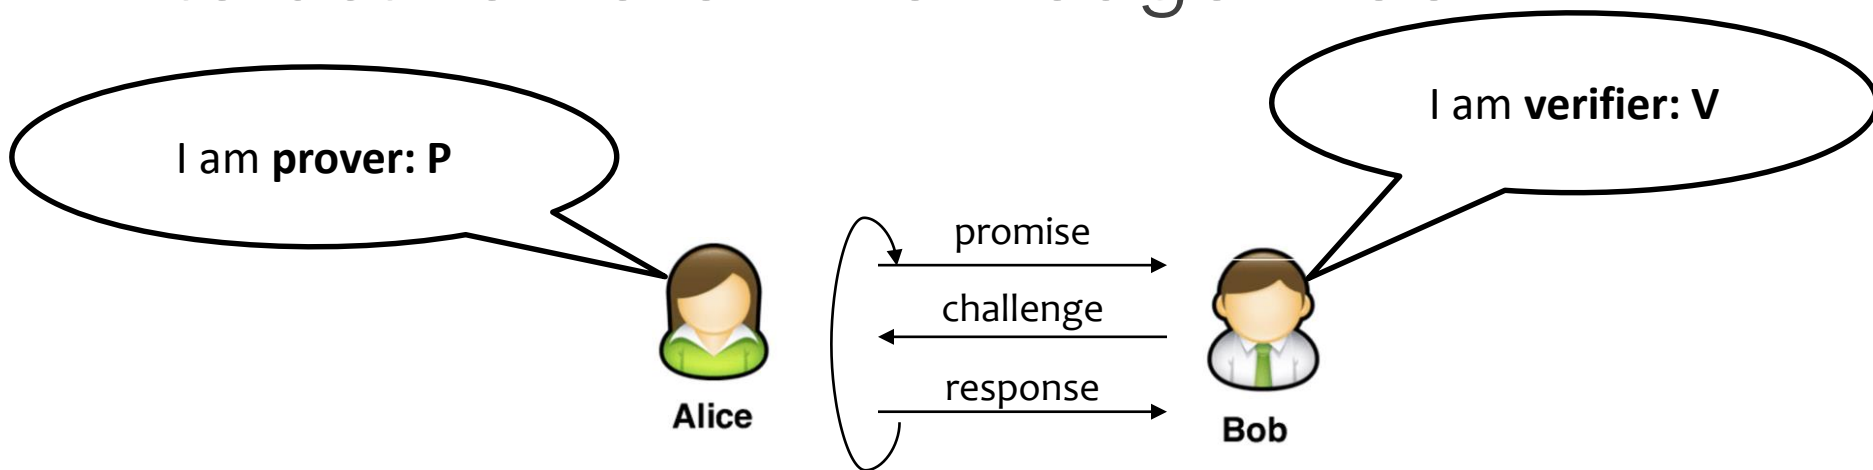
- Alice tries to **prove** to Bob that she **has the answer of a difficult problem** (e.g., a NP problem)
- Naïve method: Sending A to Bob
 - Problem: Bob will get the answer A

Zero-Knowledge Proof



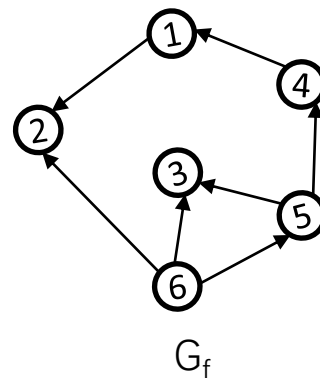
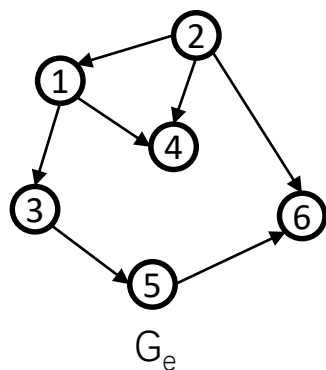
- Alice tries to **prove** to Bob that she **has the answer of a difficult problem** (e.g., a NP problem).
- Zero-Knowledge Proof
 - **Completeness**: Alice **can** construct the proof if she has A
 - **Soundness**: Alice **cannot** construct the proof if she doesn't have A
 - **Zero-knowledge**: Bob **knows nothing about A**

Interactive Zero-Knowledge Proof



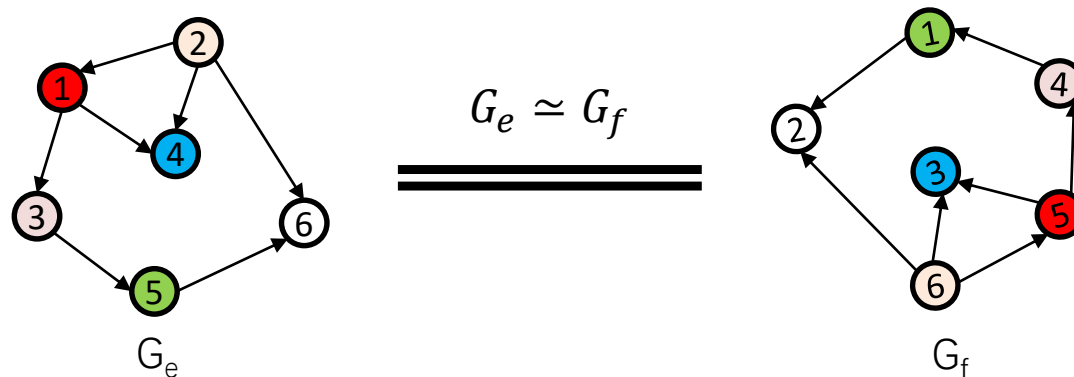
- P has answer x of a problem L , and tries to prove it with > 1 iterations:
 - Step-1: P transfers L to L' , and promises that L' is transferred from L and she has the answer x'
 - Step-2: V challenges P
 - Step-3: P shows the proof of the answer x' , **which will not leak x**
 - **V trusts that P has x when P always meets the challenge**

Graph Isomorphism



- If $\mathbf{G}_1 = (\mathbf{V}_1, \mathbf{E}_1)$ and $\mathbf{G}_2 = (\mathbf{V}_2, \mathbf{E}_2)$ are isomorphic, there exist a bijection function $\boldsymbol{\phi}$, that for any $(\mathbf{u}, \mathbf{v}) \in \mathbf{E}_1$, exist $\boldsymbol{\phi}(\mathbf{u}, \mathbf{v}) \in \mathbf{E}_2$

Graph Isomorphism



- If $\mathbf{G}_1 = (\mathbf{V}_1, \mathbf{E}_1)$ and $\mathbf{G}_2 = (\mathbf{V}_2, \mathbf{E}_2)$ are isomorphic, there exist a bijection function ϕ , that for any $(u, v) \in \mathbf{E}_1$, exist $\phi(u, v) \in \mathbf{E}_2$
- \mathbf{G}_e and \mathbf{G}_f are isomorphic
 - $\phi = \{1_1 \rightarrow 2_5, 1_2 \rightarrow 2_6, 1_3 \rightarrow 2_4, 1_4 \rightarrow 2_3, 1_5 \rightarrow 2_1, 1_6 \rightarrow 2_2\}$

Graph Isomorphism w/ Interactive ZKP

G_e and G_f are
Isomorphism,
 $\emptyset = \dots$



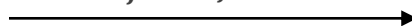
Alice



Bob

- Select a random bijection function π
- Calculate $G'_f \simeq G_f$ with π

$$G'_f \simeq G_f \simeq G_e$$



Ask Alice to prove $G_\sigma \simeq G'_f$



If $\sigma = f$, send π
If $\sigma = e$, send $\pi \cdot \emptyset$



- Select a random $\sigma \in \{e, f\}$
- Verify that $G_\sigma \simeq G'_f$

Graph Isomorphism w/ Interactive ZKP

G_e and G_f are
Isomorphism,
 $\emptyset = \dots$



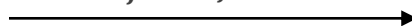
Alice



Bob

- Select a random bijection function π
- Calculate $G'_f \simeq G_f$ with π

$$G'_f \simeq G_f \simeq G_e$$



Ask Alice to prove $G_\sigma \simeq G'_f$



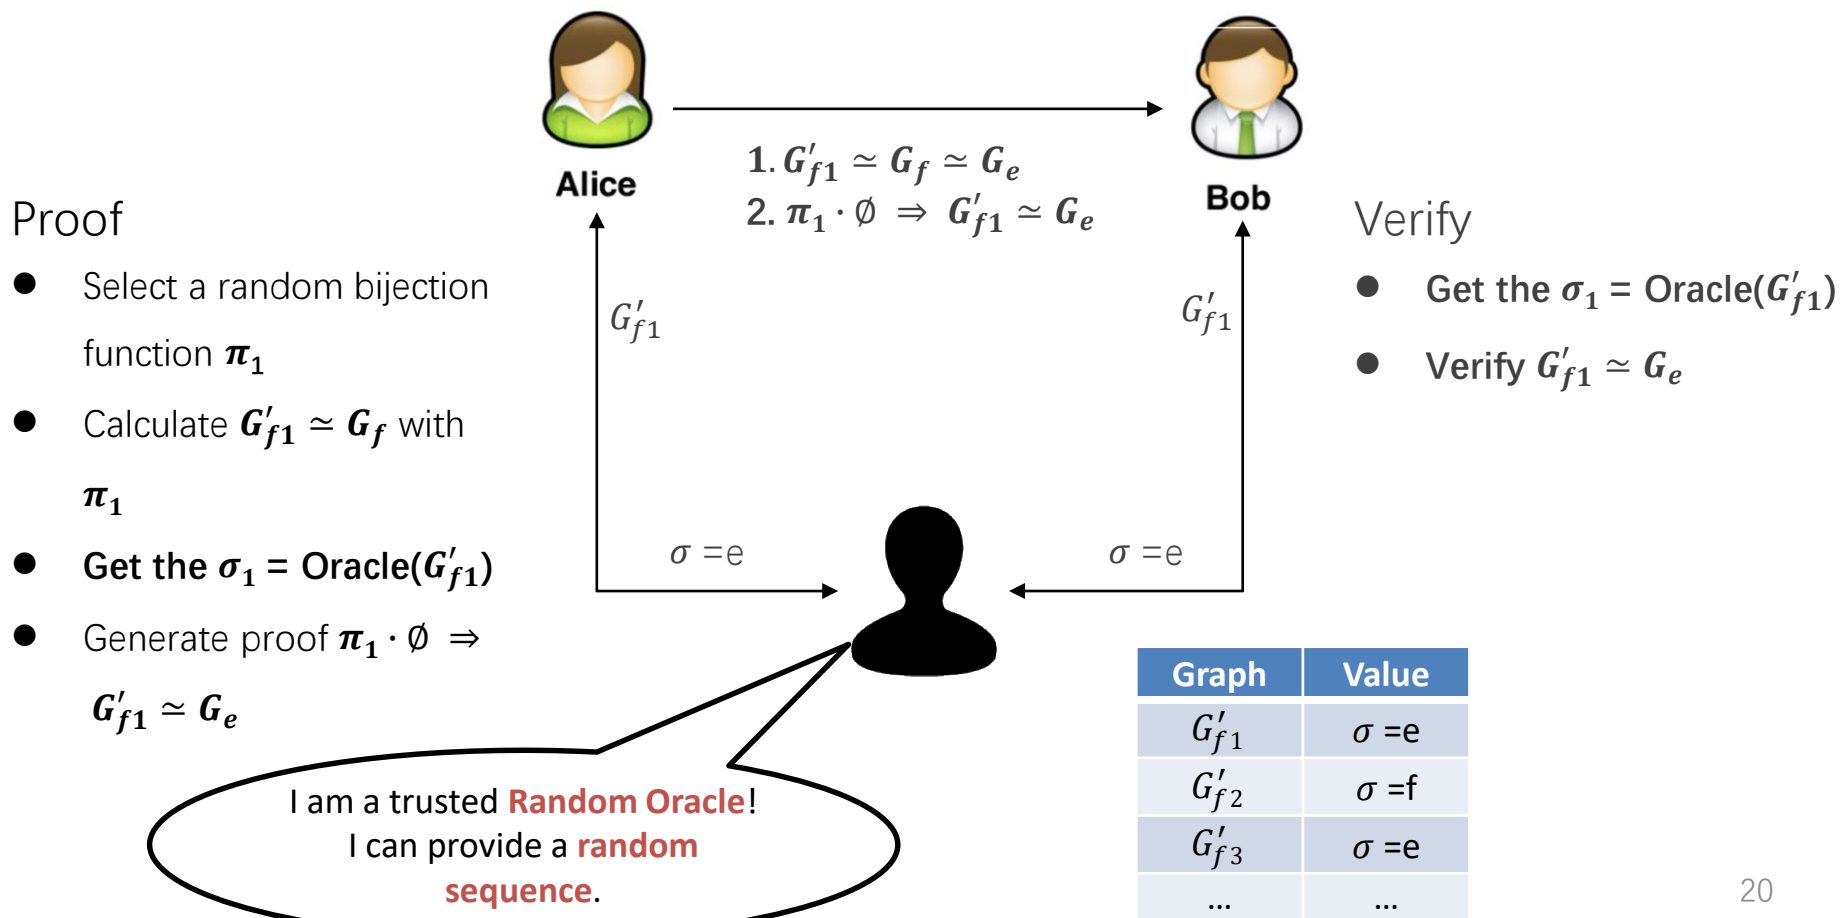
- Select a random $\sigma \in \{e, f\}$

Require too much
interactions!



- Verify that $G_\sigma \simeq G'_f$

Non-Interactive ZKP



Oblivious Transfer

OT

Problem

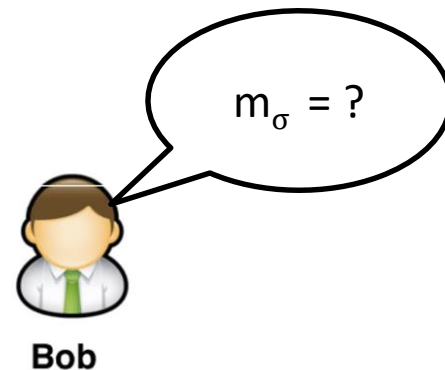
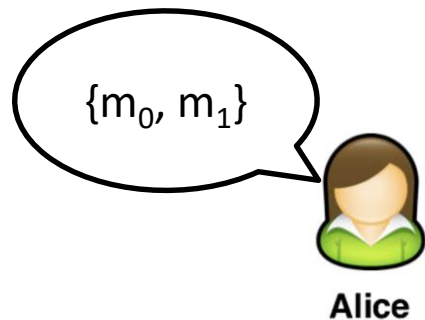


- Alice has $\{m_0, m_1\}$ and Bob wants to get m_σ
 - Alice may know the m_σ
 - Bob may get both m_0 and m_1

Oblivious Transfer

- Scenario: message transfer
 - A sender has a message list $\{m_0, m_1, \dots, m_n\}$
 - A receiver wants to get **k** target messages from sender
- Properties: oblivious and secure
 - Oblivious: sender **cannot know which messages are received**
 - Secure: receiver can **only get the target messages**

1-out-of-2 OT



- Generate (K_{pub}, K_{prv})
- Select random numbers r_0, r_1

K_{pub}, r_0, r_1

→

V

←

- $k_0 = Dec(K_{prv}, V \oplus r_0)$
- $k_1 = Dec(K_{prv}, V \oplus r_1)$
- $C_0 = Enc(k_0, m_0)$
- $C_1 = Enc(k_1, m_1)$

C_0, C_1

→

- Generate a key k
- $V = Enc(K_{pub}, k) \oplus r_\sigma$
 $\sigma \in \{0, 1\}$
- $m_\sigma = Dec(k, C_\sigma)$

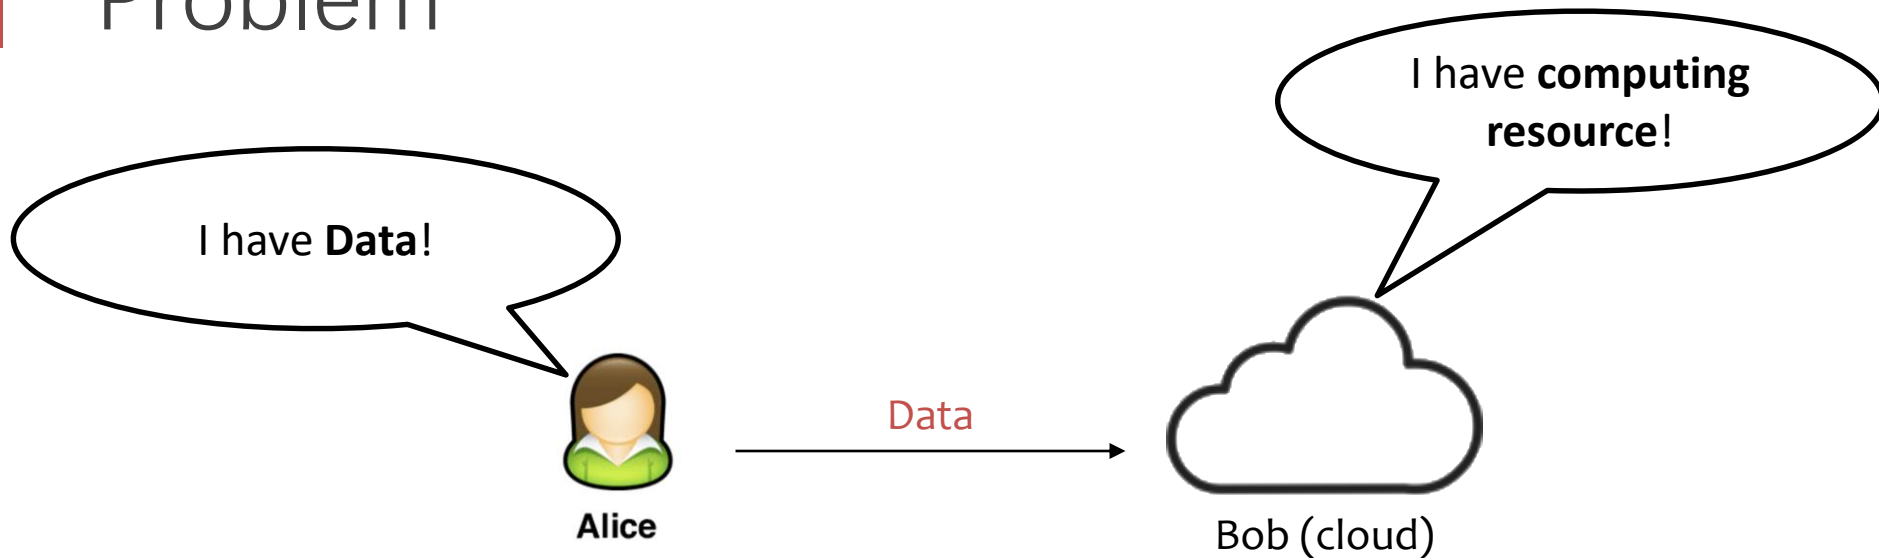
More OT Protocols

- Different numbers of selected messages
 - 1-out-of-2 OT:
 - 1-out-of- n OT
 - k -out-of- n OT
- Implementation method
 - Non-adaptive OT
 - Adaptive OT
 - Publicly Verifiable OT
 - ...

Homomorphic Encryption

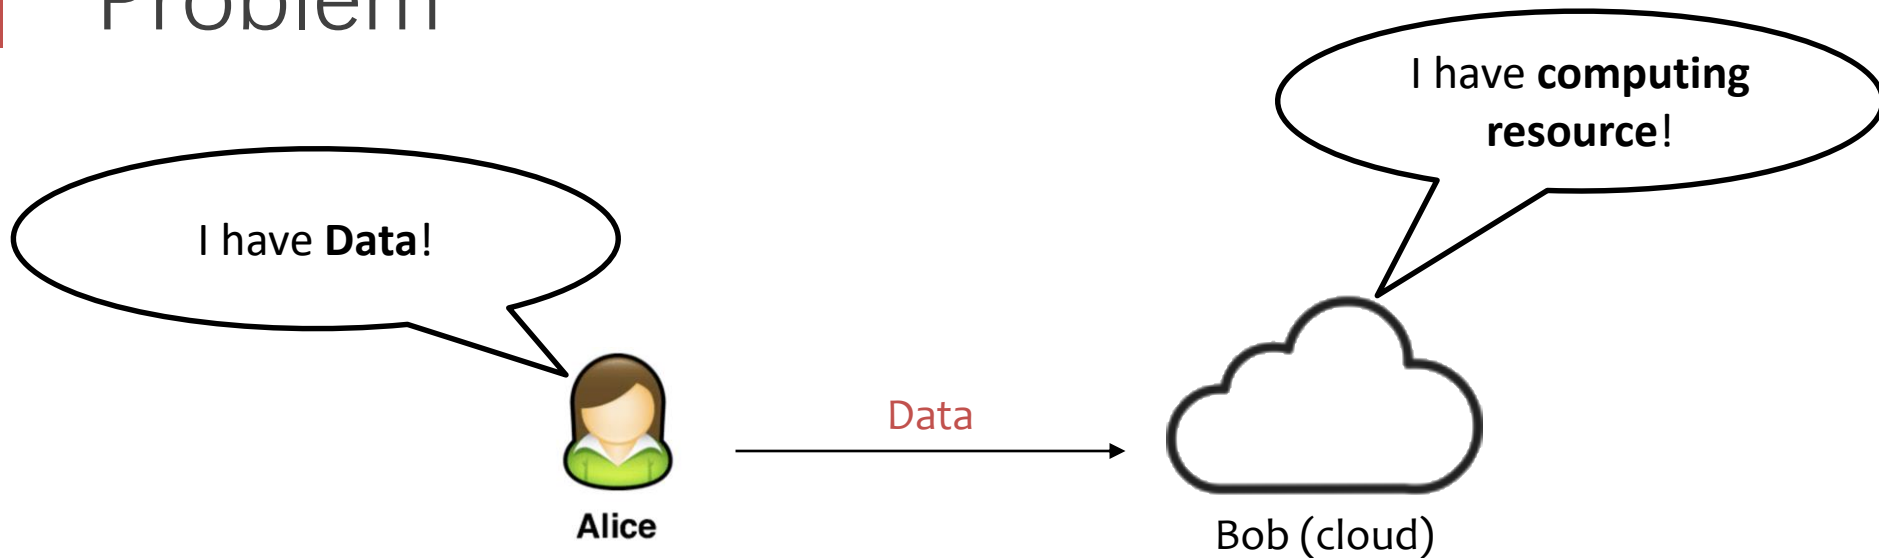
HE

Problem



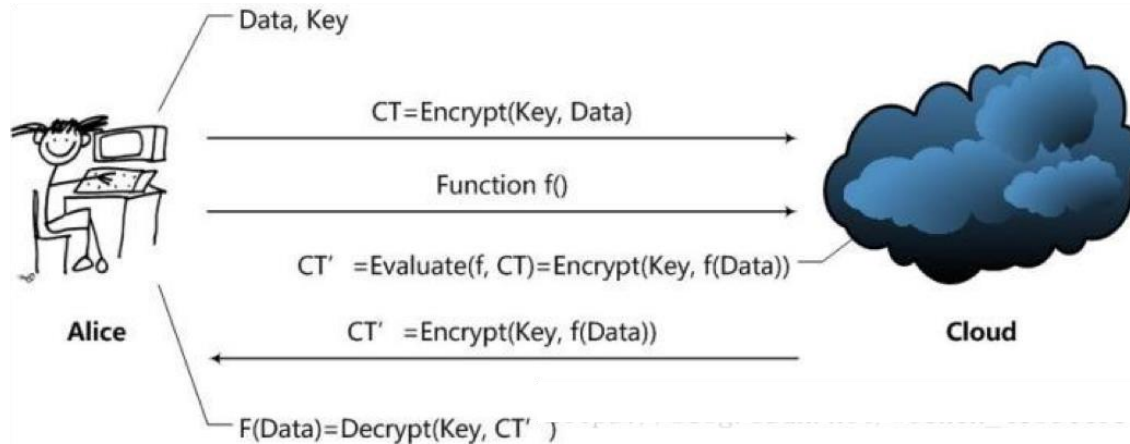
- Alice wants to ask Bob (e.g., a cloud) to perform calculation on her data
- Naïve method: Sending data to Bob

Problem



- Alice wants to ask Bob (e.g., a cloud) to perform calculation on her data
- Naïve method: Sending Data to Bob
 - Bob will get the data

Homomorphic Encryption



- Alice sends $CT = \text{encrypt}(\text{Key}, \text{Data})$ and function f to the Cloud
- Cloud calculates $CT' = \text{Evaluate}(f, CT) = \text{Encrypt}(\text{Key}, f(\text{Data}))$
- Alice gets $f(\text{Data}) = \text{Decrypt}(\text{Key}, CT')$

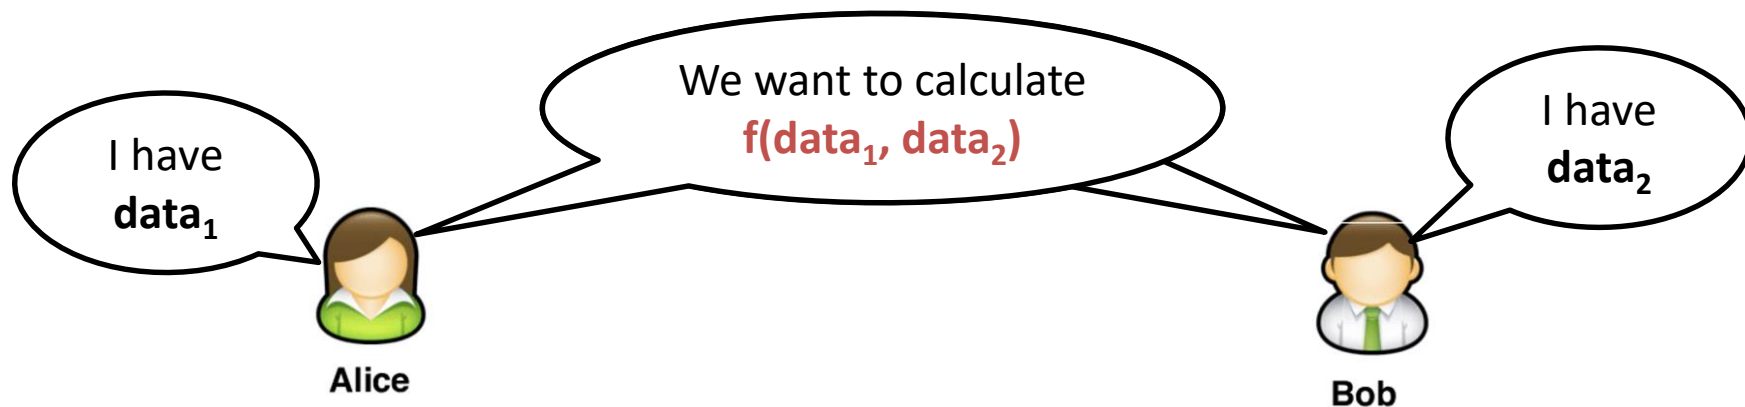
SWHE and FHE

- HE: Homomorphic Encryption
 - $Enc(f(m_1, m_2)) = Eval_f(Enc(m_1), Enc(m_2))$
- SWHE: Somewhat Homomorphic Encryption
 - Support **limited** kinds of operation
 - $f(m_1, m_2) = m_1 \cdot m_2$ (e.g., RSA)
 - $f(m_1, m_2) = m_1 + m_2$
- FHE: Full Homomorphic Encryption
 - Support **all** kinds of operations
 - Addition and multiplication

Secure Multi-Party Computing

SMPC

Problem



- Multiple parties (at least 2) work together to calculate a function
 - Enforce the **data privacy** for each party

Yao's Protocol

- Two-party computing
- Semi-honest adversary
 - Each party must **follow the protocol**
- Generic protocol
 - Can securely compute **any functionality**
- GC(Garbled Circuits)+OT(Oblivious Transfer)

Millionaire Problem

- Money: Wang has i , Lee has j , i and j are between 1 to 10
- Lee (money: j)
 - Chose a random big integer x
 - $K = \text{enc}\{\text{Key}_{\text{pub}} \text{ of Wang}, x\}$
 - Send $c=K-j$ to Wang
- Wang (money: i)
 - Decrypt with Key_{pri} of Wang ten number: $c+1, c+2 \dots c+10$, get $y_1, y_2 \dots y_{10}$
 - Chose a prime number p
 - Calculate $d_1 = y_1 \bmod p$
 - For $n=i$ to 10, d_n++ , other no change
 - Send d_1 to d_{10} to Lee
- Lee
 - Check d_j , if $d_j == x \bmod p$, then $i \geq j$; else $i < j$

Garbled Circuits

- Represent functions as **Boolean circuits**
 - Basic gates: AND, OR, NOT
 - Adding numbers
 - Comparing numbers
 - Multiplying numbers
 - Computing AES
- Represent input and output as wires

Garbled Circuits

- An **encrypted circuits** together with a pair of keys (k_0, k_1) for every wire so that for any gate, given one key for every input wire:
 - It is possible to compute the key of the corresponding gate output
 - It is **impossible to learn anything else**

Yao's Protocol on GC

- **Input:** \mathbf{x} and \mathbf{y} of length n from P_1 and P_2
- P_1 generates a garbled circuit $\mathbf{G(C)}$
 - K_L^0, K_L^1 are the keys on wire w_L
 - Let w_1, \dots, w_n be the input wires of P_1 and w_{n+1}, \dots, w_{2n} be the input wires of P_2
- P_1 sends to P_2 $\mathbf{G(C)}$ and strings $K_1^{x_1}, \dots, K_n^{x_n}$
- P_1 and P_2 run n OTs in parallel
 - P_1 inputs (K_{n+1}^0, K_{n+1}^1)
 - P_2 inputs y_i
- Given all keys, P_2 computes $\mathbf{G(C)}$ and obtains $\mathbf{C(x,y)}$

Different sMPC Protocols

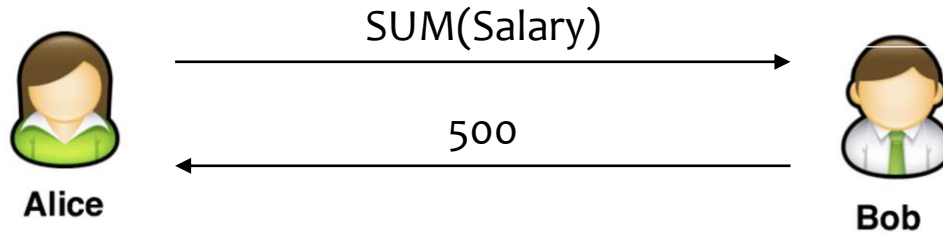
- Two-party
 - Yao's protocol
 - TinyOT protocol
 - Obliv-C
- Multi-party
 - BMR protocol
 - GMW protocol
 - SPDZ protocol



Differential privacy

DP

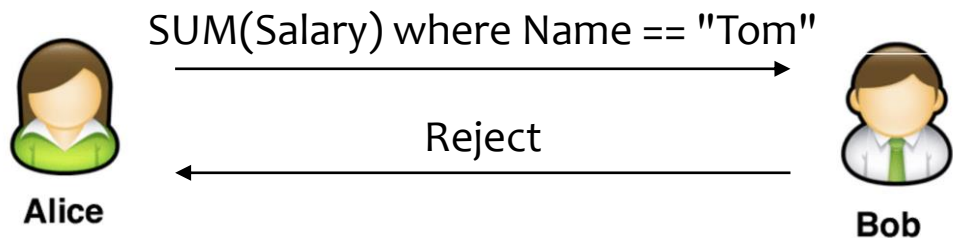
Problem



Name	Salary
Alice	100
Bob	80
Brown	200
Tom	120

- Alice can perform queries on Bob's database, but cannot access a single database entry

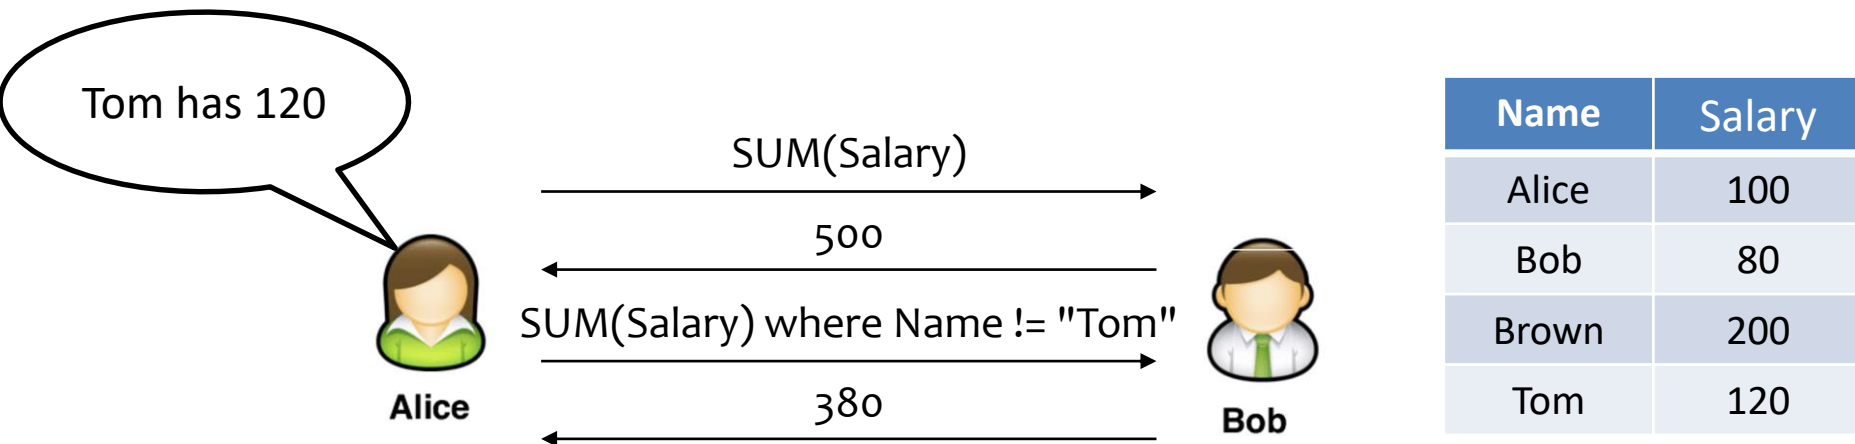
Problem



Name	Salary
Alice	100
Bob	80
Brown	200
Tom	120

- Alice can perform queries on Bob's database, but cannot access a single database entry
 - Naïve method: reject Alice to access single entry

Problem

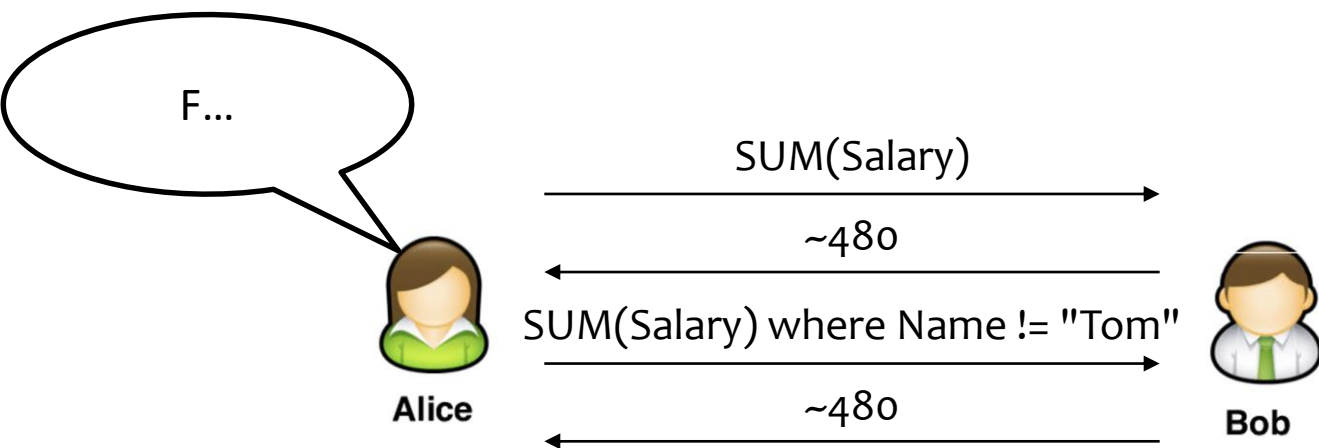


- Alice can perform queries on Bob's database, but cannot access a single database entry
 - Naïve method: reject Alice to access single entry

Differential privacy

- Allow user to perform a **random function M** on data set $\mathbf{D} = \{a_1, \dots a_n\}$, but get nothing about any individual entry of \mathbf{D}
- \mathbf{M} is ϵ -DP if:
 - For all datasets \mathbf{D} and \mathbf{D}' that **differ on a single element**, $\Pr(\mathbf{M}(\mathbf{D}) = x) \leq e^\epsilon * \Pr(\mathbf{M}(\mathbf{D}') = x)$
- Security properties
 - Robustness to post-processing
 - **User can perform any operation on the result of M, and get nothing about the individual entry of D**
 - Composability
 - Group privacy

When DP is Enabled



Name	Salary
Alice	100
Bob	80
Brown	200
Tom	120

How to Implement a DP Algorithm

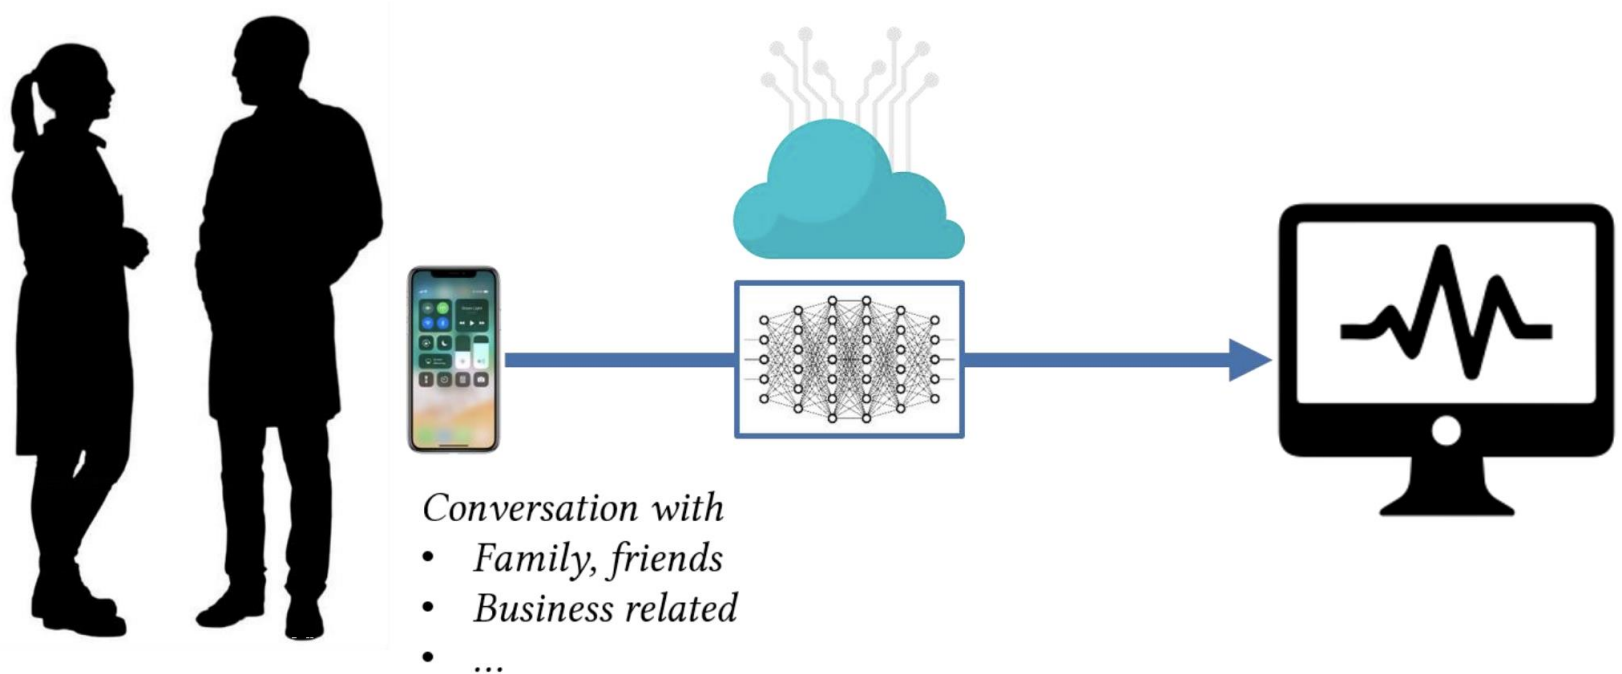
- Adding noisy to the function that we want to compute
 - Translate the function f to a random algorithm M
- Existing Mechanism
 - Laplace mechanism
 - Gaussian mechanism
 - ...

Occlumency: Privacy-preserving Remote Deep-learning Inference
Using SGX (MobiCom'2019)



OCCLUMENCY

Concerns in Cloud-driven Deep Learning



Sensitive data is exposed to leak / tampering

Cloud Offloading is Inevitable, but Risky

- Practical method to support mobile DL
 - Easily supports a large model with high accuracy
 - Consumes less resources
 - Addresses device heterogeneity
- Privacy concerns!
 - User data can be disclosed
 - Image, video, audio, activities, health/medical data
- More critical for mobile/IoT services
 - Life-immersive: data from users' daily life

Cloud Offloading is Inevitable, but Risky

Cloud

- High computational power
- Sacrifices [privacy](#)

On-device

- High privacy protection
- Sacrifices [accuracy/speed](#)

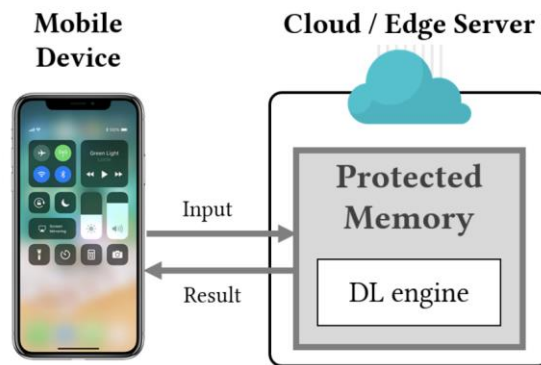
This paper aims to build a **secure cloud**-based solution

to strike the balance between [privacy](#), [speed](#) and [accuracy](#)

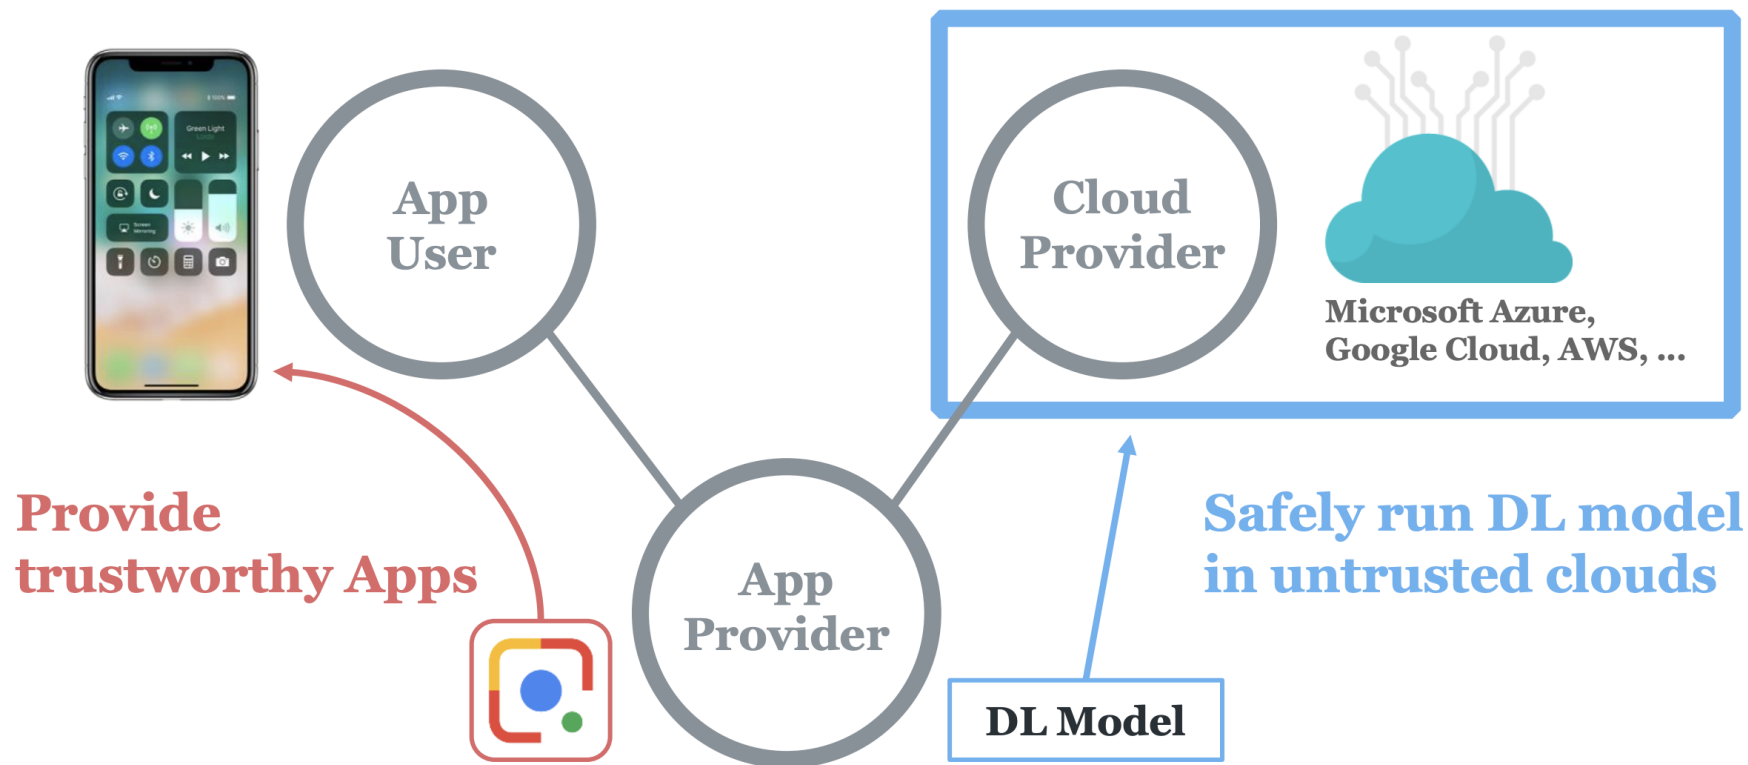
Occlumency

A cloud-driven DL inference system preserving user privacy

- Key approach: [SGX enclaves](#)
 - Commodity TEE with the highest protection level
 - Prevent memory access even from OS / hypervisor
- Protects:
 - User data disclosure
 - Inference result manipulation

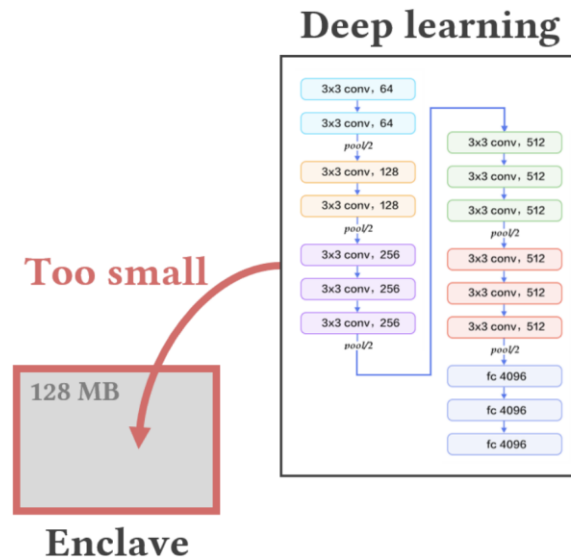


Stakeholders of Occlumency



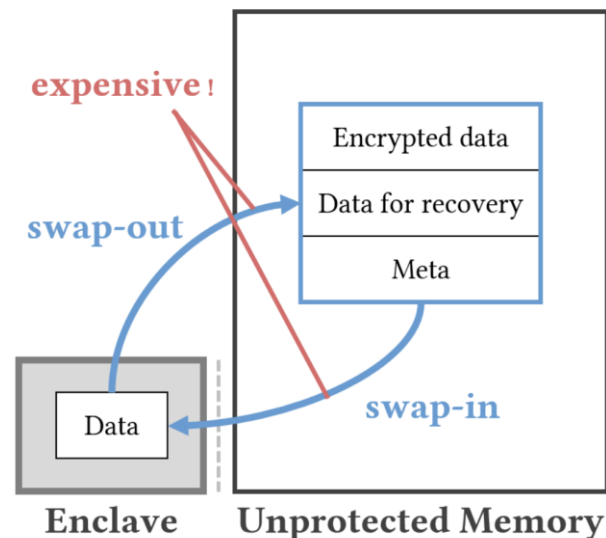
Challenge: Limited Memory Size

- SGX's physical memory is very small (128 MB)
 - DNN requires 100MB ~ 1GB of memory
 - Windows -- fails
 - Linux -- makes frequent page swapping



Challenge: Limited Memory Size

- SGX's physical memory is very small (128 MB)
 - DNN requires 100MB ~ 1GB of memory DNN
⇒ Frequent page swapping
- SGX's paging is expensive
 - Swaps-out memory into untrusted memory
 - Involves encryption, redundancy checking, ...
- SGX slows down the inference speed
 - Takes 7x longer latency to infer VGG, YOLO



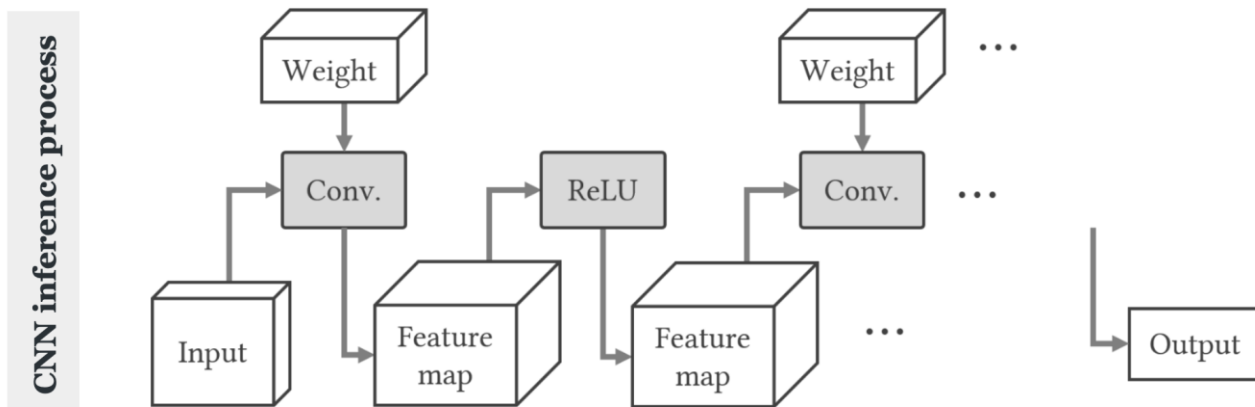
Enabling DL to Run within Small Memory

Observation -- three dominant memory usages of DNN:

- 1) Model weights (parameters)
- 2) Intermediate feature maps
- 3) Conv. layer computation

Memory Usage of VGG-16

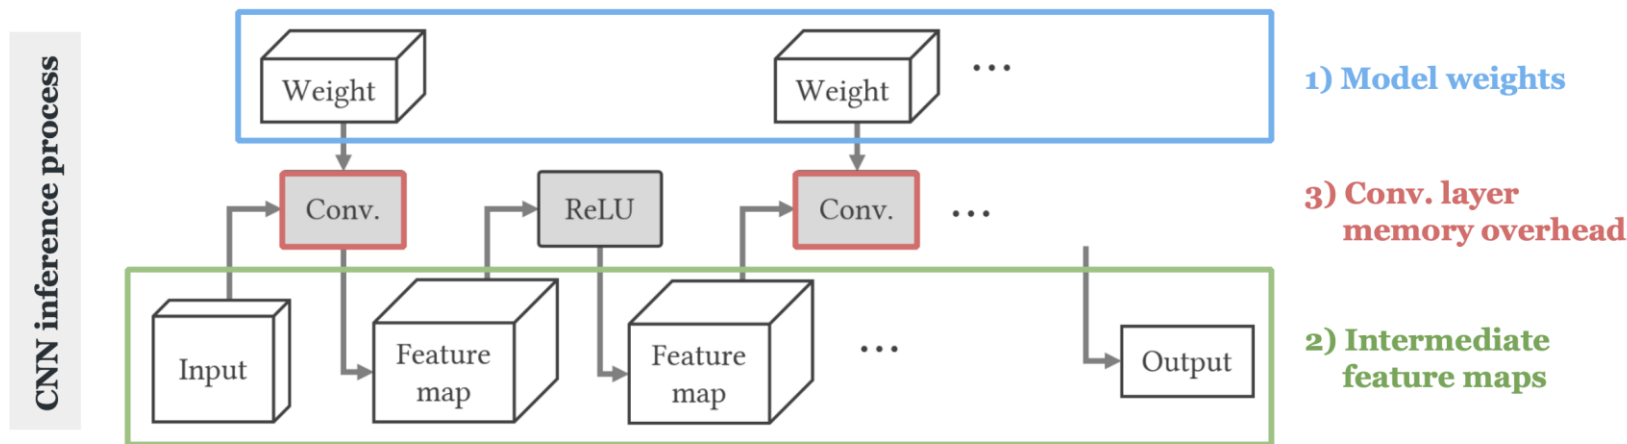
1) Model weights	540 MB
2) Feature maps	61 MB
3) Conv. layer	327 MB
Others	6 MB



Enabling DL to Run within Small Memory

Observation -- three dominant memory usages of DNN:

- 1) Model weights (parameters)
- 2) Intermediate feature maps
- 3) Conv. layer computation

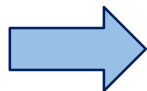


Approach of Occlumency

1) Model weights (parameters)

2) Intermediate feature maps

3) Conv. layer computation



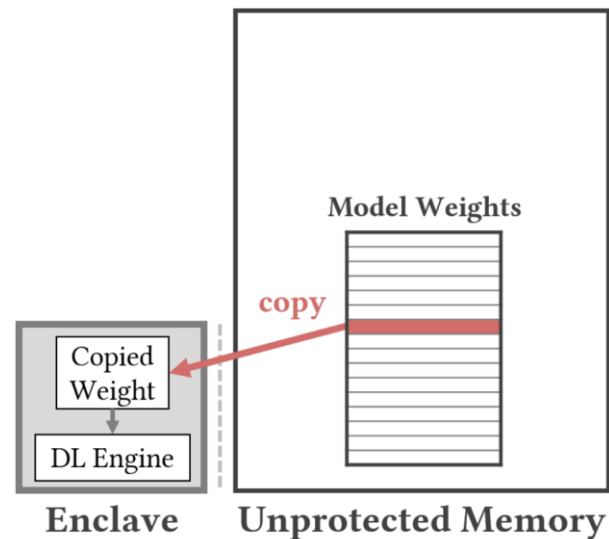
1) On-demand weight loading

2) Memory-efficient FM allocation

3) Partitioned convolution

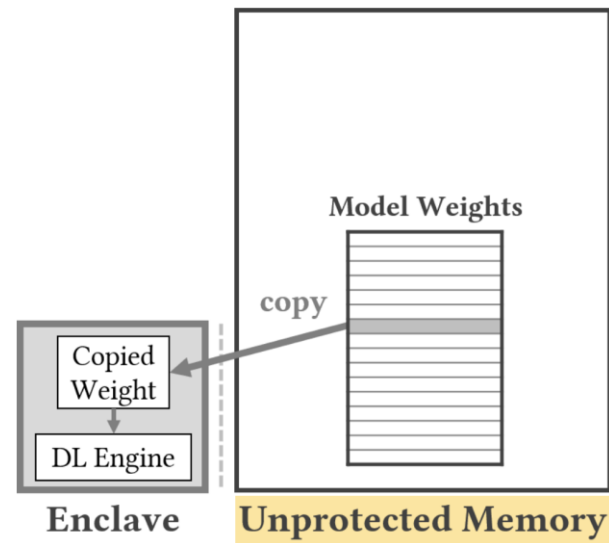
1) On-demand Weights Loading

- Saves memory used to load weights (~500MB for VGG-16)
- Idea: Not protecting model weights
 - Our goal is to protect the user privacy
 - Model weights are irrelevant to user data
- Keeps weight in **unprotected memory** & copies into enclaves on demand



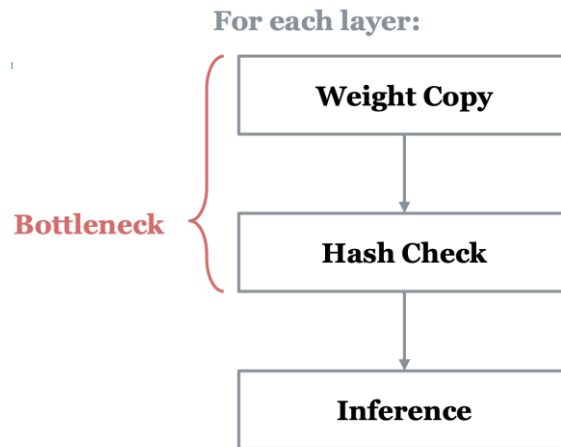
1-1) New Problem 1: Weight Corruption Threat

- Weight may be corrupted
 - Weights are no longer protected in enclaves
 - Weight manipulation attacks will lead to wrong inference results
- Solution: Model integrity checking
 - [Hash checking](#)-based weight modification detection
 - Compares hash value of each layer



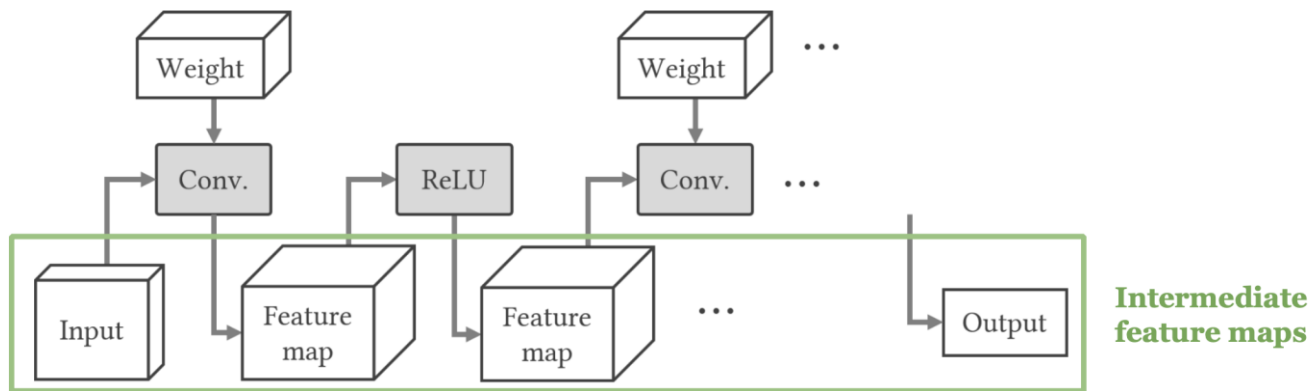
1-2) New Problem 2: Computation Bottleneck

- Additional computation bottlenecks:
 - Weight copying (on-demand weight loading)
 - Hash checking (model integrity checking)
- Parallel **pipeline**
 - Weight copy / Hash check / Inference
 - Reduces ~17.5% of latency



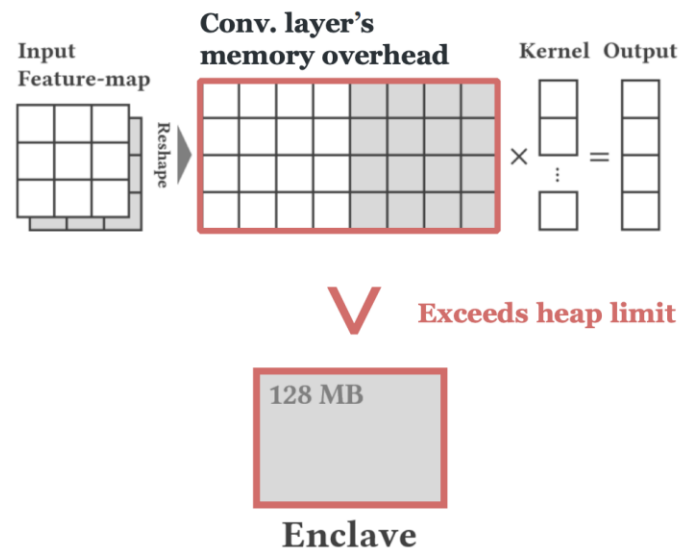
2) Memory-efficient Feature Map Allocation

- Reduces the required memory to load **intermediate feature-maps (FM)**
- Idea: Releasing unnecessary FMs
 - **Profiles** when each FM can be released in advance
 - **Immediately deallocates** FMs already used



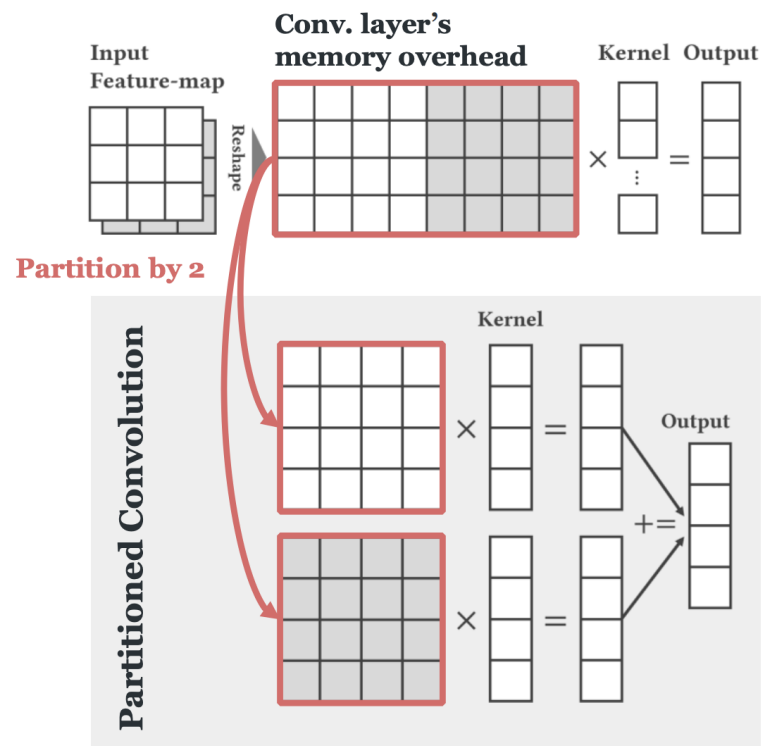
3) Partitioned Convolution

- Reduces memory overhead of Conv. that **exceeds** the SGX heap limit



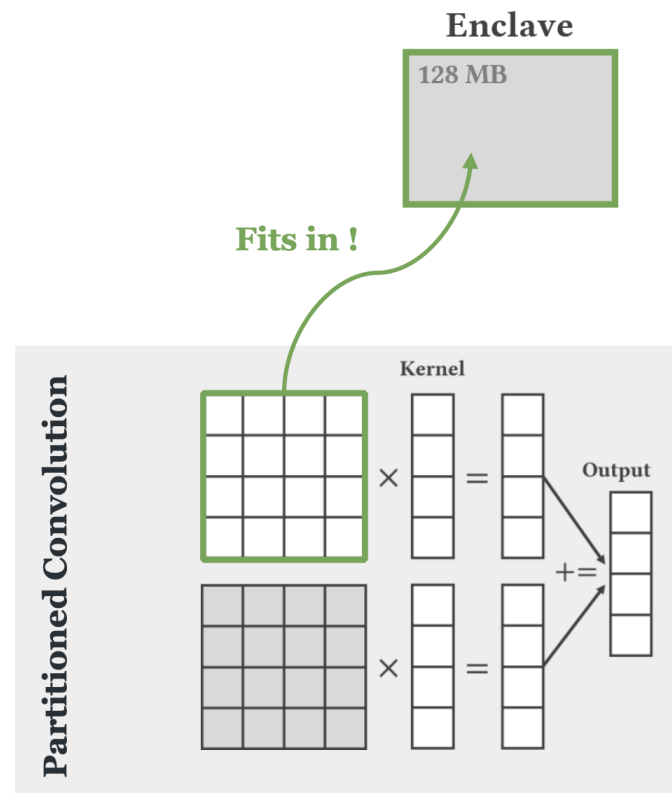
3) Partitioned Convolution

- Reduces memory overhead of Conv. that **exceeds** the SGX heap limit
- Idea: Breaking down big operations into smaller jobs
 - e.g., Partition by 2 \rightarrow requires **2x less** memory



3) Partitioned Convolution

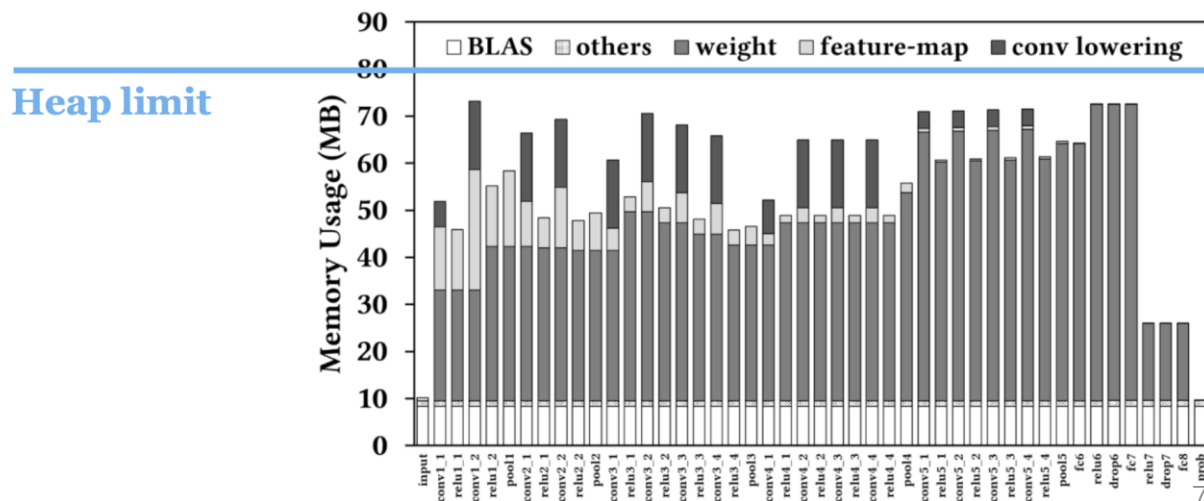
- Reduces memory overhead of Conv. that **exceeds** the SGX heap limit
- Idea: Breaking down big operations into smaller jobs
 - e.g., Partition by 2 \rightarrow requires **2x less** memory
- Adaptively partitions by 2, 4, 8, ...
 - Runs Conv. layer within limited memory size



Memory Usage Evaluation

1) Can Occlumency run DNN within SGX's memory limit?

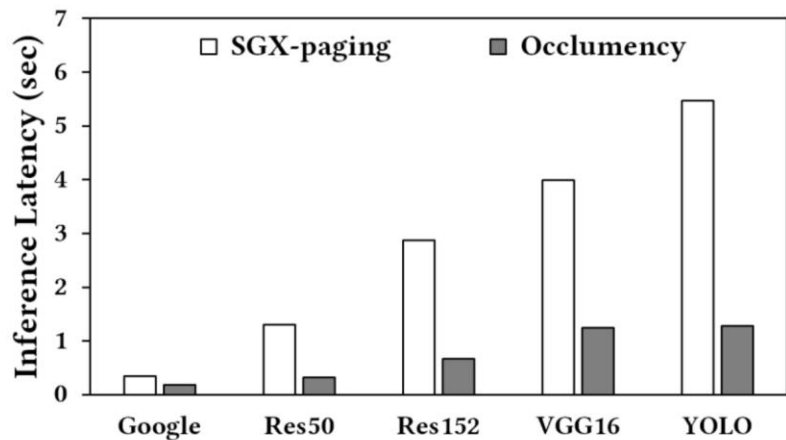
- Occlumency successfully runs DNN models within [SGX's heap limit](#)
- Ex) Reduced memory for VGG-19:
 - Original: [980MB](#)
 - Occlumency: [74MB](#)



Inference Latency Evaluation

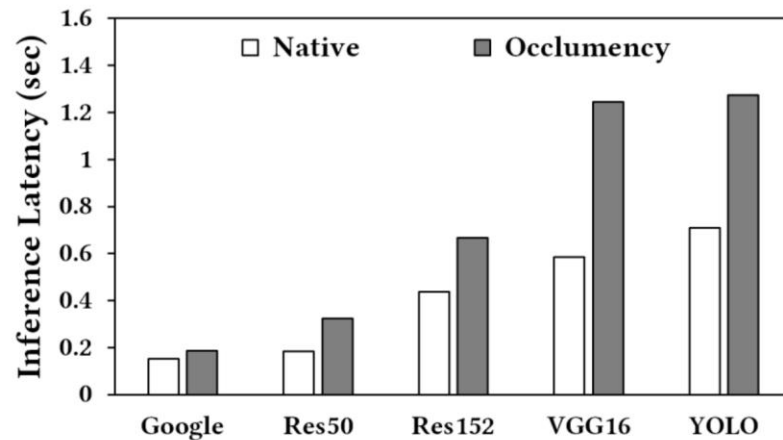
2) How much the Occlumency can enhance the inference speed?

- 3.0 ~ 4.3x faster than **SGX-paging**



3) How much is the overhead of Occlumency?

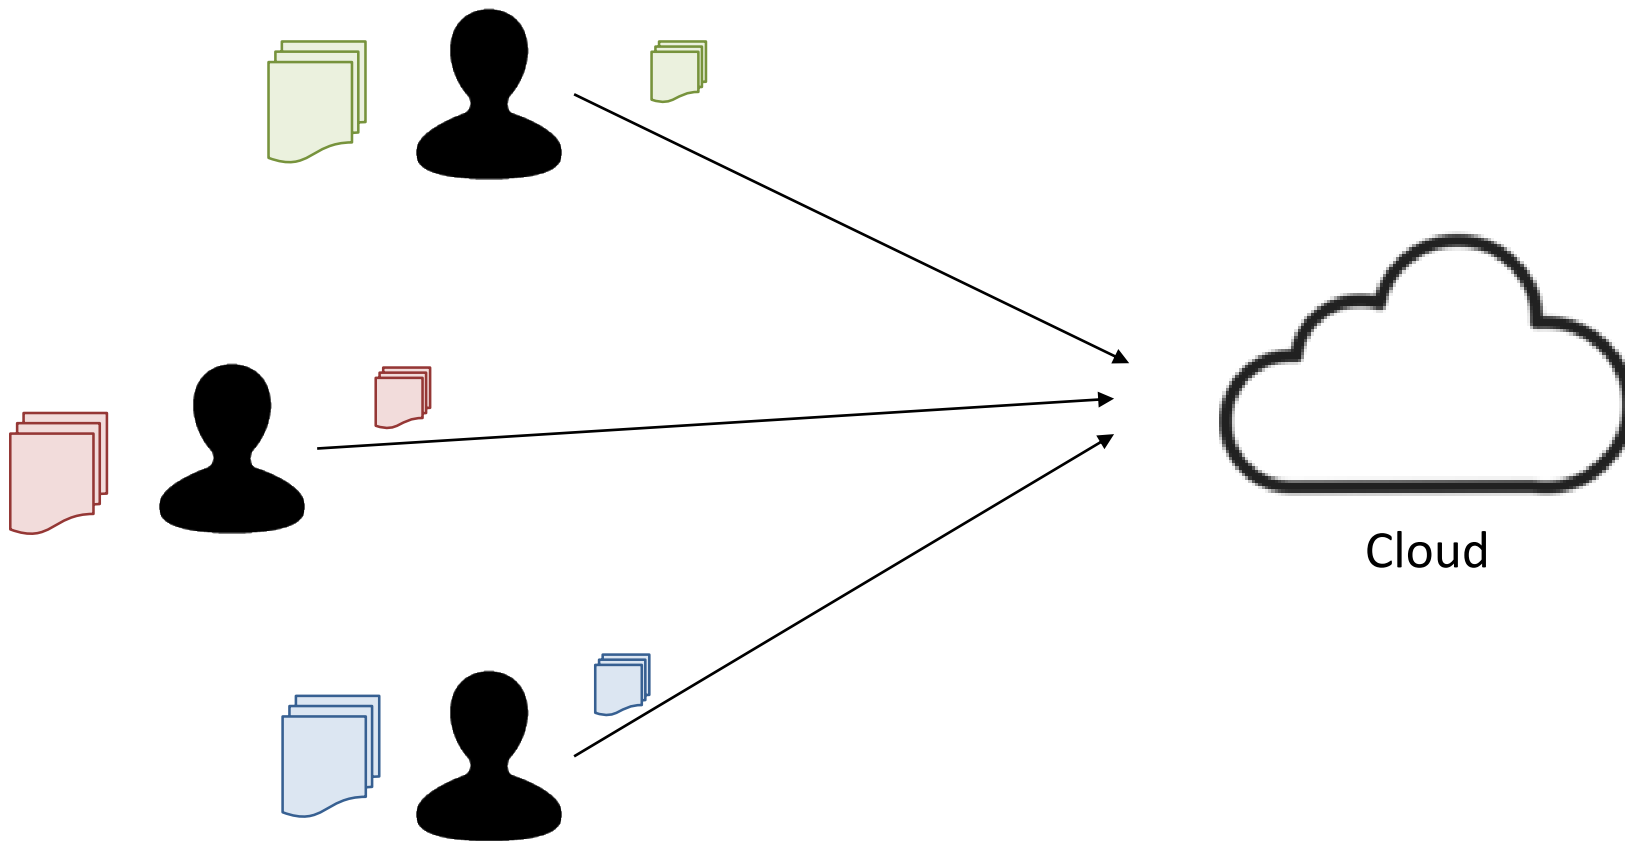
- 72% overhead compared to **Native**



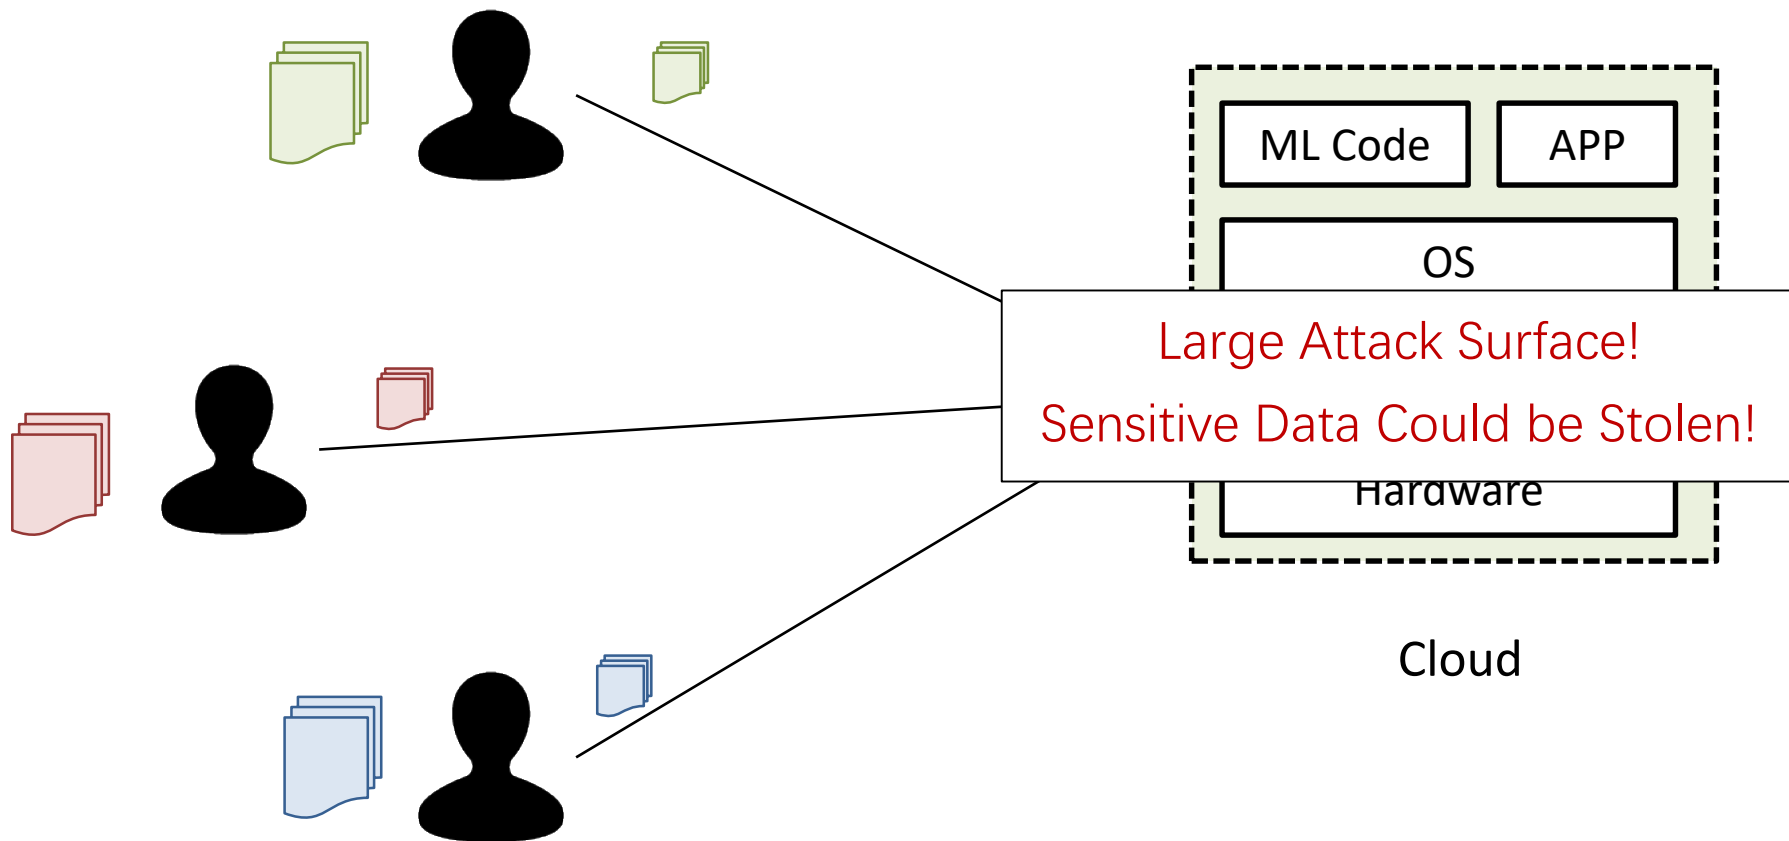
Oblivious Multi-Party Machine Learning on Trusted Processors
(USENIX Security'2016)

OBLIVIOUS MULTI-PARTY MACHINE

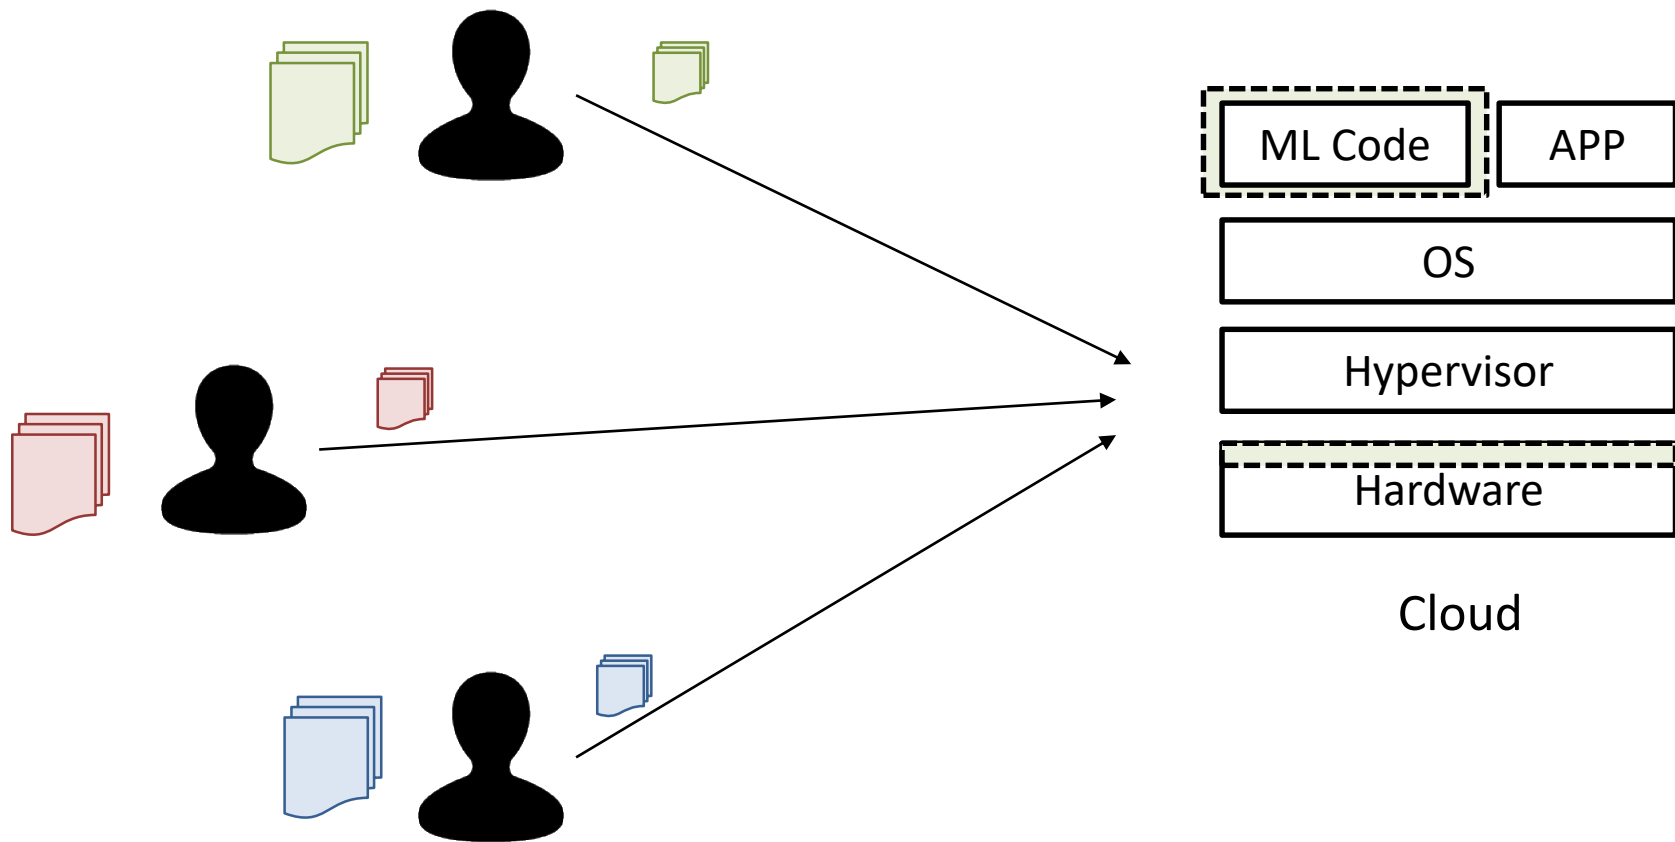
Machine Learning on Cloud



Machine Learning on Cloud



Machine Learning on Cloud with TEE

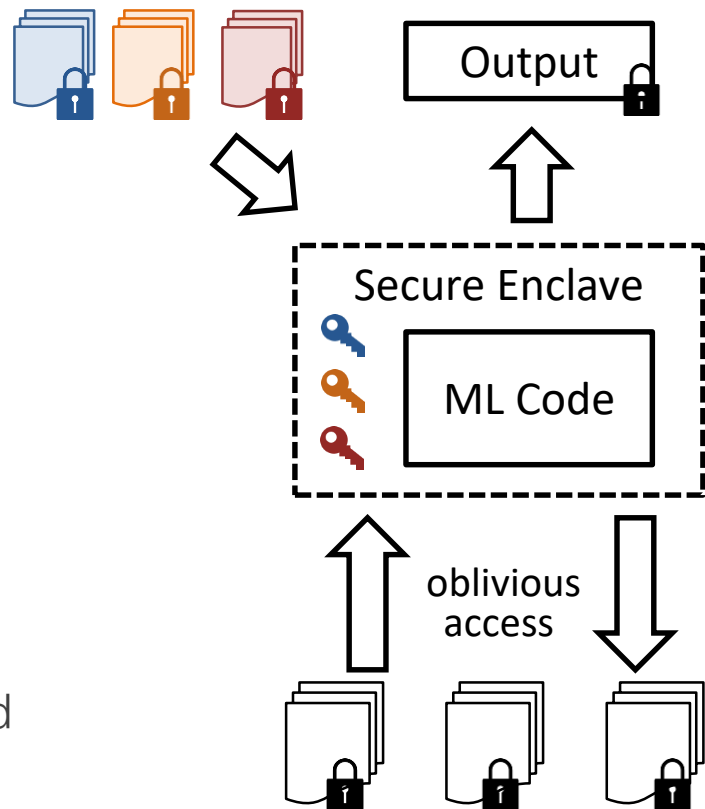


Threat Model

- Any party may be malicious
- The cloud can be malicious
 - **Memory & Network observer**
 - Hardware attackers (on mem bus)
 - Perform **side-channel attacks**
- Assumptions:
 - Code does not leak secrets
 - Do not consider leakage through time or power channels

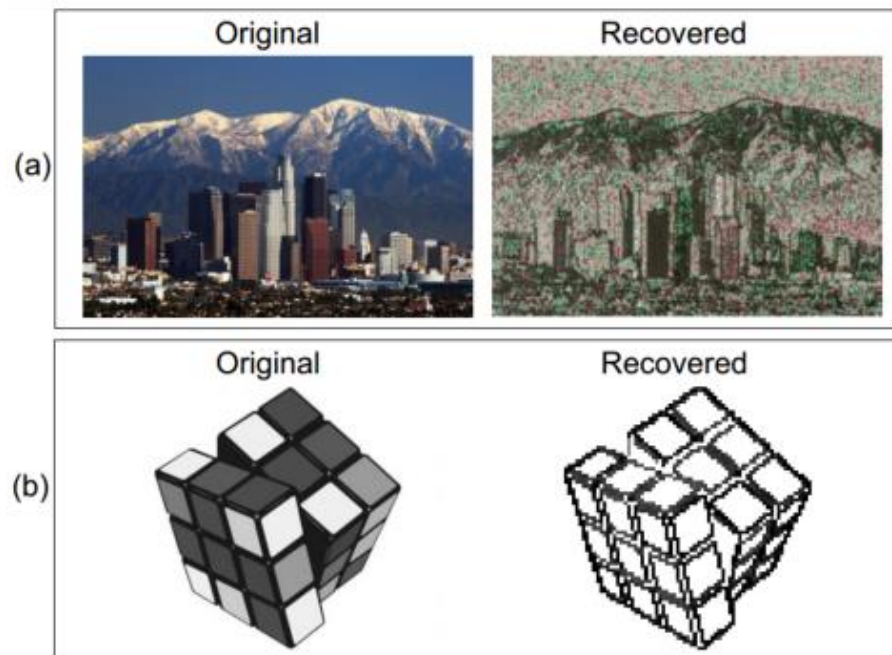
System Design

- Data: encryption
 - Input and output are encrypted
 - Data outside of enclave is encrypted
- Code: secure enclave
 - Trusted processors (SGX)
- Data accesses:
 - Side channel protection
 - Memory, disk, and network are accessed obliviously



SGX's Vulnerabilities

- Side channel attack
 - Malicious OS still controls page fault handler
 - If know the photo processing algorithm, can get the image by monitoring page fault
 - Not 100% accurate, but still good enough



Side-channel Protection

- Memory side-channel
- Security guarantee:
 - **Data oblivious**
 - Given two inputs and a memory trace, one cannot distinguish which one was executed
- Memory accesses **only depends on public information**
 - E.g., number of instances, number of labels
- Assumption: register-to-register manipulation is data oblivious

Library of Oblivious Primitives LibO

- In assembly:
 - ogreater, omove, oless, oequal
 - oget
 - get the ith array element (hide i)

ogreater()

mov	rcx, x
mov	rdx, y
cmp	rcx, rdx
setg	al
retn	

omove()

mov	rcx, cond
mov	rdx, x
mov	rax, y
test	rcx, rcx
cmovz	rax, rdx
retn	

Oblivious Operation

Non-oblivious

```
int max(int x, int y) {  
    if (x > y) return x;  
    else return y;  
}
```

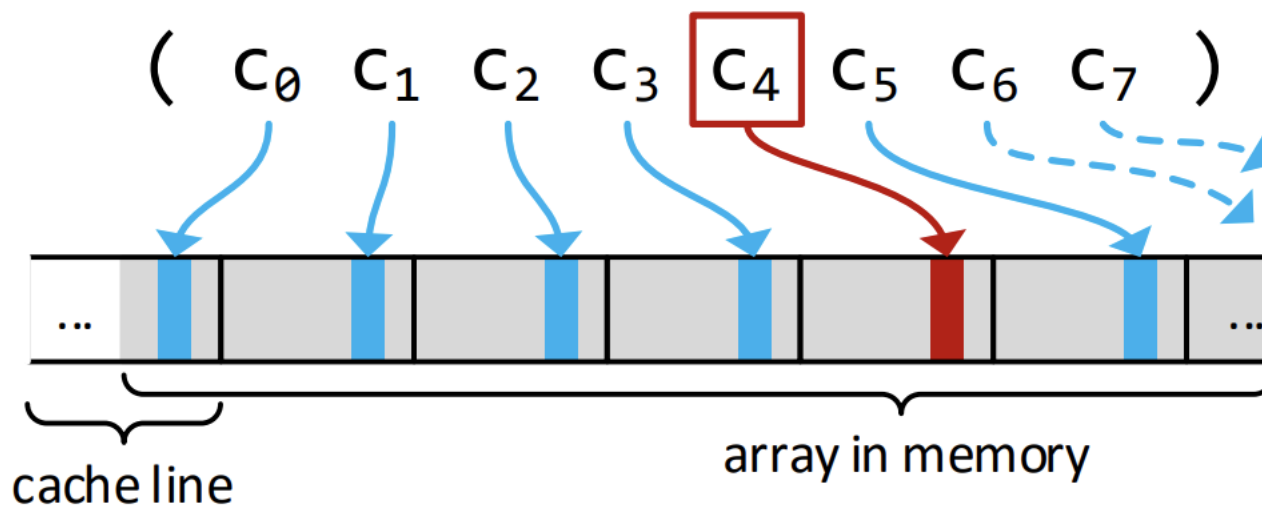
Oblivious

```
int max(int x, int y) {  
    bool getX = ogreater(x, y);  
    return omove(getX, x, y);  
}
```

Optimized Oblivious Array Access

- Naïve method
 - Scan the array
 - Just actually load/store a single element
- Observation
 - Attacker can only trace memory access at cache line granularity
- Optimization
 - Scan arrays at cache line granularity
 - Leveraging AVX2 vector instructions

Optimized Oblivious Array Access



Other Oblivious ML Algorithms

- Decision trees
- Support Vector Machines
- Neural Network
- Matrix Factorization
- K-Means clustering

Evaluation

Algorithm	SGX+enc.	SGX+enc.+obl.	Dataset	Parameters	Input size	# Instances
K-Means	1.91	2.99	MNIST	$k=10, d=784$	128MB	70K
CNN	1.01	1.03				
SVM	1.07	1.08	SUSY	$k=2, d=18$	307MB	2.25M
Matrix fact.	1.07	115.00	MovieLens	$n=943, m=1,682$	2MB	100K
Decision trees	1.22	31.10	Nursery	$k=5, d=27$	358KB	6.4K

Baseline is processing the data in plaintext without SGX protection.

Different Data Privacy Systems

- Data privacy + ML
 - Sage (SOSP'19)
 - Oblivious multi-party ML (USENIX Security'16)
 - Chiron, ...
- Data privacy + Database
 - CryptDB (SOSP'11)
 - EnclaveDB (S&P'18), ...
- Data privacy + Data analysis
 - Opaque (NDSI'17)
- More ...