

第 3 讲: Virtual Machine Monitor

第三节: Hardware-assisted Virtualization

陈渝

清华大学计算机系

yuchen@tsinghua.edu.cn


2020 年 2 月 29 日




Virtualization holes

Conventional Intel® 64, or IA-32, is not virtualizable architecture

- Sensitive register instructions: read or change sensitive registers and/or memory locations such as a clock register or interrupt registers
 - SGDT, SIDT, SLDT
 - SMSW
 - PUSHF, POPF
- Protection system instructions: reference the storage protection system, memory or address relocation system
 - LAR, LSL, VERR, VERW
 - POP, PUSH
 - CALL, JMP, INT n, RET
 - STR, MOVE




Source from proceeding of 2000 USENIX ATC



Controlling the CPU Resource

With an Intel® 64 CPU, a VMM must be able to retain control over


- Access to privileged state (CRn, DRn, MSRs)
- Exceptions (#PF, #MC, etc.)
- Interrupts and interrupt masking
- Address translation (via page tables)
- CPU access to I/O (via I/O ports or MMIO)



CPU Control via “Ring Deprivileging”

With an Intel® 64 CPU, a VMM must be able to retain control over


- Guest OS kernel runs in a less privileged ring than usual
- VMM runs in the most privileged ring 0
- prevent guest OS from accessing privileged instructions / state
- prevent guest OS from modifying VMM code and data



Problems with Ring Deprivileging

Ring deprivileging encounters various issues in current Intel® 64 architecture:


- Ring compression/aliasing
- Non-faulting reads of privileged state
- Excessive faulting
- Interrupt-virtualization issues



Addressing “Virtualization Holes”

- Paravirtualization
- Binary translation (or patching)
- Hardware-assisted virtualization
 - Closing virtualization holes in hardware
 - Simplify VMM software
 - Optimizing for performance

Solution – hardware-assisted VT




Intel® Virtualization Technology

A hardware-assisted virtualization technology, named as Intel® Virtualization Technology, or Intel® VT

- For Intel® 64, VT-x: CPU MEM
- For directed I/O, VT-d: DMA/Interrupt remapping
- For connectivity, VT-c: offload work from CPU to IO device

Virtual Machines (VMs)

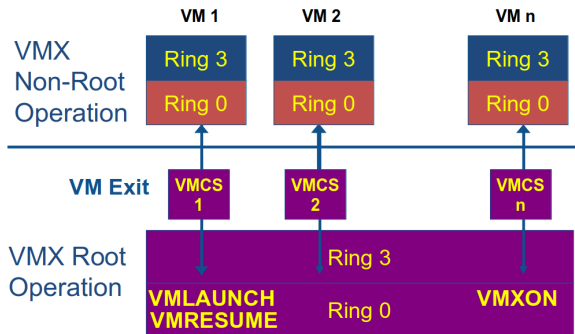


VM entry

- Transition from VMM to guest
- Enters VMX non-root operation
- VMLAUNCH used for initial entry
- VMRESUME used subsequently

VM exit

- Guest-to-VMM transition
- Enters VMX root operation
- Caused by external events, exceptions, some instructions



Virtual Machine Control Structure (VMCS)

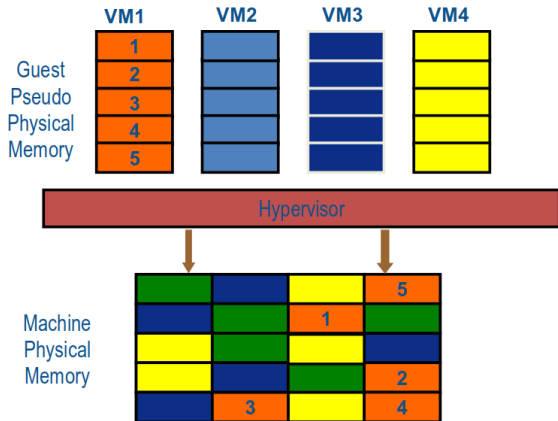
Controls VMX non-root operation/transitions

- Each virtual CPU should have its own VMCS
- Only one VMCS active at a time on a physical CPU

Each VMCS stored in a memory region

- Physical address set by new **VMPTRLD** instruction
- Need not reside in linear-address space
- **VMPTRLD** also used to switch VMCS

Physical page frame redirection



Memory virtualization challenges

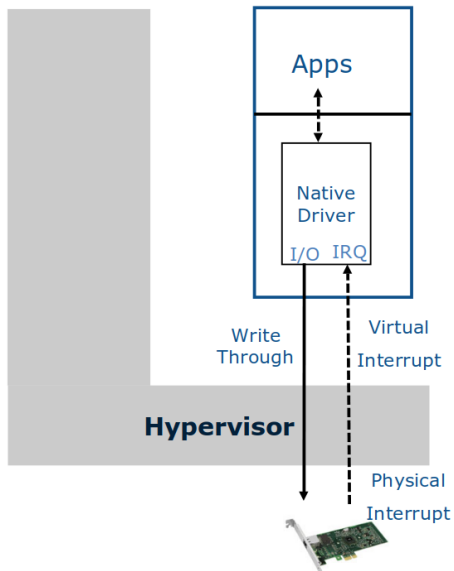
- OS expect to see physical memory starting from 0
- OS expect to see contiguous memory in address space
- BIOS/Legacy OS are designed to boot from address low 1M
- DMA, TLB, etc.

Extended Page Table (EPT)



Memory virtualization challenges

- Guest can have full control over its page tables and events:
 - CR3, INVLPG, page fault
- VMM controls Extended Page Tables:
 - BIOS/Legacy OS are designed to boot from address low 1M
 - DMA, TLB, etc.

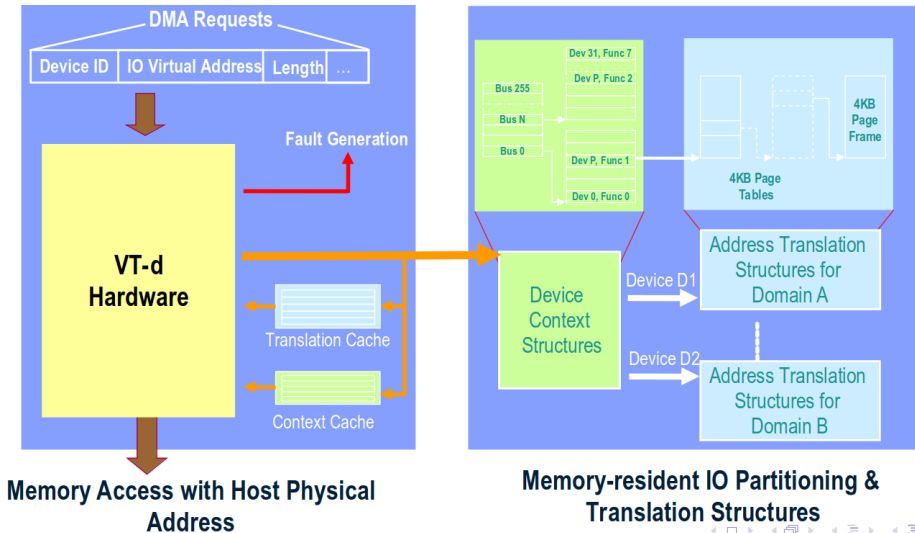



Direct assignment

- Guest runs native driver
- I/O is written through
- Physical interrupt is captured by hypervisor (pIRQ)
- Virtual interrupt is signaled for guest (vIRQ)
- Remapping from pIRQ <-> vIRQ in hypervisor

Maximizing performance, but sacrificing sharing

Intel® VT for directed I/O (VT-d)





Single Root I/O Virtualization (SR-IOV)

- Close to native performance
- Limited CPU overhead
- Flexible sharing
- Security control through PF

Maximizing performance

RISC-V H-Extension: Compare ARM64

How is RISC-V H-Extension compared to ARM64 virtualization?

RISC-V H-Extension v0.4 draft	ARM64 (ARMv8.x) Virtualization
No separate privilege mode for hypervisors. Extends S-mode with hypervisor capabilities (HS-mode) and Guest/VM run in virtualized S-mode/U-mode (VS-mode or VU-mode).	Separate EL2 exception-level for hypervisors with it's own <xyz>_EL2 MSRs. The Guest/VM will run in EL1/EL0 exception levels.
Well suited for both Type-1 (baremetal) and Type-2 (hosted) hypervisors. The S<xyz> CSRs access from VS-mode map to special VS<xyz> CSRs which are only accessible to HS-mode and M-mode.	Special ARMv8.1-VHE Virtualization Host Extension for better performance of Type-2 (hosted) hypervisor. Allows Host kernel (meant for EL1) to run in EL2 by mapping <xyz>_EL1 MSRs to <abc>_EL2 MSRs in Host mode.
Virtual interrupts for Guest/VM injected using VSIP CSR. The hypervisor does not require any special save/restore but it will emulate entire PLIC in software.	Virtual interrupts for Guest/VM injected using LR registers of GICv2/GICv3 with virtualization extension. The hypervisor will save/restore LR registers and emulate all GIC registers in software except GIC CPU registers.

RISC-V H-Extension: Compare ARM64 (Contd.)

How is RISC-V H-Extension compared to ARM64 virtualization?

RISC-V H-Extension v0.4 draft	ARM64 (ARMv8.x) Virtualization
Virtual timer events for Guest/VM using SBI calls emulated by hypervisor. The SBI calls trap to hypervisor so save/restore of virtual timer state not required.	Virtual timer events for Guest/VM using ARM generic timers with virtualization support. The hypervisor will save/restore virtual timer state and manage virtual timer interrupts.
Virtual inter-processor interrupts for Guest/VM using SBI calls emulated by hypervisor. The hypervisor does not require any special save/restore.	Virtual inter-processor interrupts for Guest/VM by emulating ICC_SGI1R_EL1 (virtual GICv3) or GICD_SGIR (virtual GICv2). The save/restore will be handled as part of LR registers save/restore.
Nested virtualization supported using HSTATUS.TVM and HSTATUS.TSR bits. The hypervisor will trap-n-emulate Guest hypervisor capabilities.	Special ARMv8.3-NV for supporting nested virtualization on ARMv8. The hypervisor will trap-n-emulate Guest hypervisor capabilities. The ARMv8.4-NV further enhances nested virtualization support.

Summary

Linux KVM RISC-V

