

# 第 4 讲: Optimization of Virtual Machine Monitor

## 第一节: Introduction

陈渝


清华大学计算机系

*yuchen@tsinghua.edu.cn*

2020 年 3 月 8 日




# Introduction



Before optimization of VMM ...

- Where is the bottlenecks of VMM?
- How to find these bottlenecks?


# Introduction – Where is the bottlenecks of VMM?




bottlenecks of VMM

- VCPU
- VMEM
- VIO
- etc...

# Introduction – bottlenecks of VMM



## Containers



## Comparison between VMM and Container

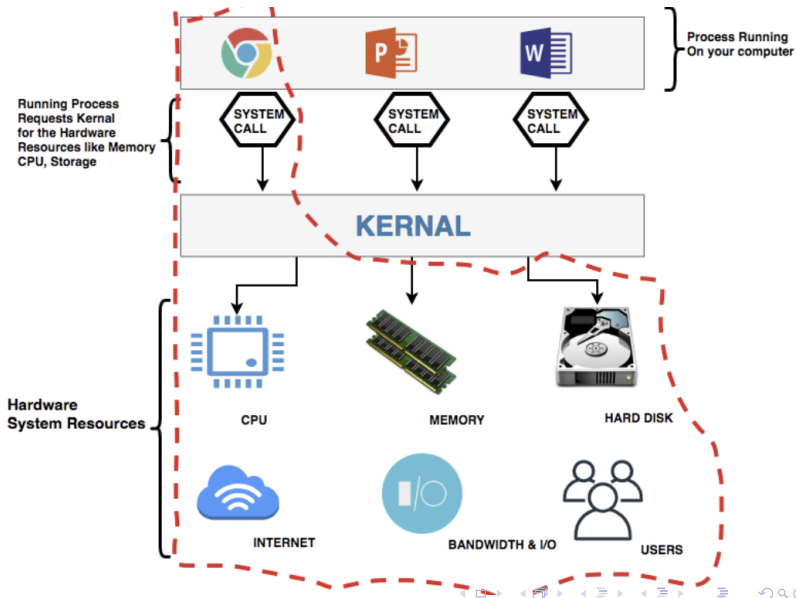

- Target: Docker v.s. KVM/XEN
- Benchmark: Environment and Workload
- Testing & Evaluation

from: An Updated Performance Comparison of Virtual Machines and Linux Containers, TR of IBM, 2014

from: My VM is Lighter (and Safer) than your Container and Linux Containers, SOSPP'17, 2017



## Containers

container container container




## Containers

container container container




*Isolating resources per process (or group of processes)*



*Limit amount of resources used per process*

## A Container Skeleton



## What is inside a Container

- Program Binaries/configuration
- Runtime libraries
- Dependency Products/tools
- A Piece of Kernel
- System Resources:  
CPU/MEM/IO/Net/Storage

We are isolating the program and delicately providing its own system resources and runtime libraries.

# Introduction – bottlenecks of VMM – CPU/MEM

Workload		Native	Docker	KVM-untuned	KVM-tuned
PXZ (MB/s)		76.2 [ $\pm 0.93$ ]	73.5 (-4%) [ $\pm 0.64$ ]	59.2 (-22%) [ $\pm 1.88$ ]	62.2 (-18%) [ $\pm 1.33$ ]
Linpack (GFLOPS)		290.8 [ $\pm 1.13$ ]	290.9 (-0%) [ $\pm 0.98$ ]	241.3 (-17%) [ $\pm 1.18$ ]	284.2 (-2%) [ $\pm 1.45$ ]
RandomAccess (GUPS)		0.0126 [ $\pm 0.00029$ ]	0.0124 (-2%) [ $\pm 0.00044$ ]	0.0125 (-1%) [ $\pm 0.00032$ ]	Tuned run not warranted
Stream (GB/s)	Add	45.8 [ $\pm 0.21$ ]	45.6 (-0%) [ $\pm 0.55$ ]	45.0 (-2%) [ $\pm 0.19$ ]	
	Copy	41.3 [ $\pm 0.06$ ]	41.2 (-0%) [ $\pm 0.08$ ]	40.1 (-3%) [ $\pm 0.21$ ]	
	Scale	41.2 [ $\pm 0.08$ ]	41.2 (-0%) [ $\pm 0.06$ ]	40.0 (-3%) [ $\pm 0.15$ ]	
	Triad	45.6 [ $\pm 0.12$ ]	45.6 (-0%) [ $\pm 0.49$ ]	45.0 (-1%) [ $\pm 0.20$ ]	

## Environment

- two E5-2665 (16 cores with HT), 256 GB of RAM
- Direct 10 Gbps Ethernet link between two Mellanox ConnectX-2 EN NICs
- Ubuntu 13.10, kernel 3.11, Docker 1.0, QEMU 1.5.0
- 32vCPU (Power management was disabled)



# Introduction – bottlenecks of VMM – CPU/MEM

Workload		Native	Docker	KVM-untuned	KVM-tuned
PXZ (MB/s)		76.2 [ $\pm 0.93$ ]	73.5 (-4%) [ $\pm 0.64$ ]	59.2 (-22%) [ $\pm 1.88$ ]	62.2 (-18%) [ $\pm 1.33$ ]
Linpack (GFLOPS)		290.8 [ $\pm 1.13$ ]	290.9 (-0%) [ $\pm 0.98$ ]	241.3 (-17%) [ $\pm 1.18$ ]	284.2 (-2%) [ $\pm 1.45$ ]
RandomAccess (GUPS)		0.0126 [ $\pm 0.00029$ ]	0.0124 (-2%) [ $\pm 0.00044$ ]	0.0125 (-1%) [ $\pm 0.00032$ ]	Tuned run not warranted
Stream (GB/s)	Add	45.8 [ $\pm 0.21$ ]	45.6 (-0%) [ $\pm 0.55$ ]	45.0 (-2%) [ $\pm 0.19$ ]	
	Copy	41.3 [ $\pm 0.06$ ]	41.2 (-0%) [ $\pm 0.08$ ]	40.1 (-3%) [ $\pm 0.21$ ]	
	Scale	41.2 [ $\pm 0.08$ ]	41.2 (-0%) [ $\pm 0.06$ ]	40.0 (-3%) [ $\pm 0.15$ ]	
	Triad	45.6 [ $\pm 0.12$ ]	45.6 (-0%) [ $\pm 0.49$ ]	45.0 (-1%) [ $\pm 0.20$ ]	

## Results

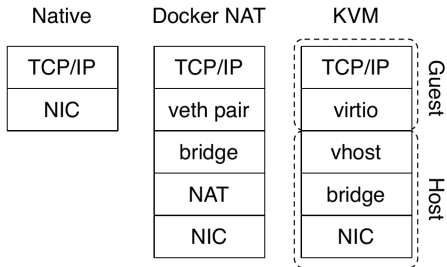
- PXZ: extra TLB pressure of nested paging.
- Linpack: pin vCPUs  $\rightarrow$  pCPU, expose the underlying cache topology
- STREAM: bandwidth of mem, cost of handling TLB misses, uses large pages

TABLE II. STREAM COMPONENTS

Name	Kernel	Bytes per iteration	FLOPS per iteration
COPY	$a[i] = b[i]$	16	0
SCALE	$a[i] = q * b[i]$	16	1
ADD	$a[i] = b[i] + c[i]$	24	1
TRIAD	$a[i] = b[i] + q * c[i]$	24	2

# Introduction – bottlenecks of VMM – Net


## Network configurations




netperf to measure network bandwidth

- All three configurations reach 9.3 Gbps
- NAT noticeably increases overhead
- vhost reduces overhead
- NIC is the bottleneck

## TCP bulk transfer efficiency (CPU cycles/byte)




## Network round-trip latency ( $\mu$ s)




# Introduction – bottlenecks of VMM – Storage

## Storage configurations




- 20 TB SSD, two 8 Gbps Fibre Channel links
- fio with the libaio backend in O DIRECT mode
- Fibre Channel HBA is the bottleneck

## Sequential I/O throughput (MB/s)



## Random I/O throughput (IOPS)




# Introduction – bottlenecks of VMM – Real Apps


## MySQL configurations

Configuration	Network (Figure 1)	Storage
Native	Native	Native
Docker net=host Volume	Native	Native
Docker NAT Volume	NAT	Native
Docker NAT AUFS	NAT	AUFS
KVM	vhost-net	virtio + qcow


## MySQL throughput (transactions/s)




## Redis performance (requests/s)




## Average latency (in ms) of operations on Redis



# Introduction – bottlenecks of VMM – Summary



## Containers



## Summary

- Containers and VMs impose almost no overhead on CPU and memory usage; they only impact I/O and OS interaction.
- This overhead comes in the form of extra cycles for each I/O operation, so small I/Os suffer much more than large ones.
- Several additional topics worthy of investigation: performance isolation when multiple workloads run on the same server, live resizing of containers and VMs, tradeoffs between scale-up and scale-out, and tradeoffs between live migration and restarting.
- <https://github.com/thewmf/kvm-docker-comparison>