

第 5 讲：The Interface of OS

第二节：Overview of POSIX

陈渝

清华大学计算机系

yuchen@tsinghua.edu.cn

2020 年 3 月 15 日



Introduction – History




First Research Edition (1971)


- Complete rewrite (4213 lines kernel)
- Reference architecture
 - 34 system calls
 - 18 common with PDP-7 version
 - 18 survive until today
- Binary code API
- Abstraction of standard I/O
- Devices as files

Half Century of Unix: History, Preservation, and Lessons Learned, Diomidis Spinellis, 2017
Analyzing a Decade of Linux System Calls, Mojtaba Bagherzadeh, etc. ICSE 2018


Introduction – History



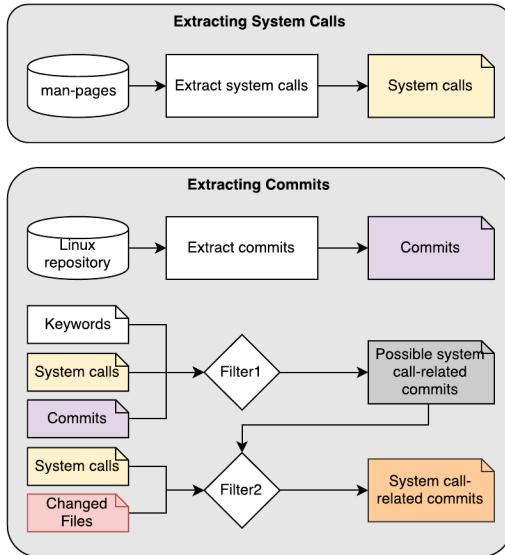
Introduction – History



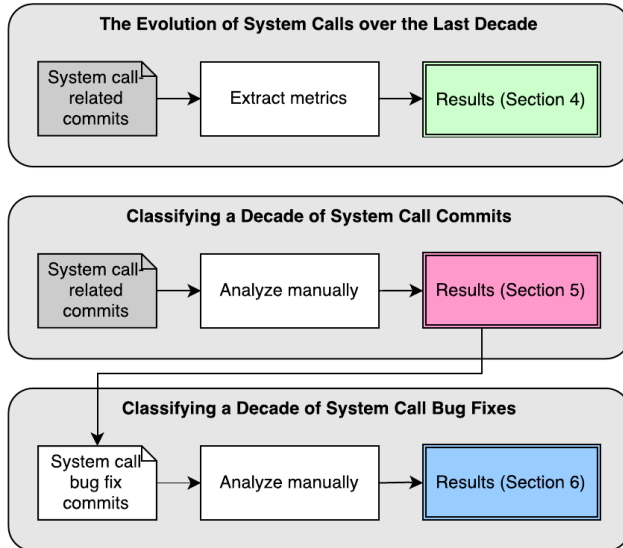
Introduction – The sequence of a system call



Introduction – An overview of syscall data collection



Introduction – An overview of syscall empirical study



Introduction – Syscall Categories

Code	Category	Example	Total
FS	File system & I/O	Reading and writing a file.	147
PM	Process management	Creating, cloning or debugging a process.	71
IPC	IPC ⁴ & network	Sharing memory between processes.	51
MM	Memory management	Mapping pages in memory.	27
SH	Signal handling	Killing a process.	24
TO	Time operations	Setting and querying the time.	23
SI	System info & settings	Retrieving information about the system.	21
SC	Scheduling	Thread prioritization.	14
SEC	Security & capabilities	Performing security checks.	8
MO	Modules	Loading a module.	6
All system calls			393

Introduction – sibling syscalls

Type of sibling	Pattern	# of calls	Example
<i>Parameter extension</i>	<i>*[1..4]</i>	12	<code>dup()</code> , <code>dup2()</code>
<i>Architecture</i>	<i>*[32/64]</i>	32	<code>truncate()</code> , <code>truncate64()</code>
<i>Working directory</i>	<i>*at</i>	14	<code>open()</code> , <code>openat()</code>
<i>Backwards compatibility</i>	<i>*old</i>	6	<code>vm86()</code> , <code>vm86old()</code>
<i>Real time</i>	<i>rt*</i>	8	<code>sigreturn()</code> , <code>rt_sigreturn()</code>
<i>Others</i>	-	30	<code>waitpid()</code> , <code>wait4()</code>
Total number		102	

Introduction – new syscalls

Functionality	# of system calls	Example
<i>Monitoring</i>	8	<code>inotify()</code> , <code>getcpu()</code>
<i>Synchronization</i>	7	<code>eventfd()</code> , <code>signalfd()</code>
<i>Hardware-specific</i>	6	<code>cacheflush()</code> , <code>move_pages()</code>
<i>Message passing</i>	5	<code>process_vm_readv()</code> , <code>tee()</code>
<i>Security</i>	3	<code>bpf()</code> , <code>seccomp()</code>
<i>Other¹</i>	7	<code>setns()</code> , <code>clock_adjtime()</code>
Total number	36	

Introduction – Classifying a Decade of System Call Commits

- **Add/remove:** The commit was made to add or remove one or more system calls.
- **Bug fix:** The commit was made to fix a bug.
- **Improvement:** The commit was made to make an improvement.
- **Restructuring:** The commit was made to conduct code restructuring, such as cleaning up comments or refactoring.

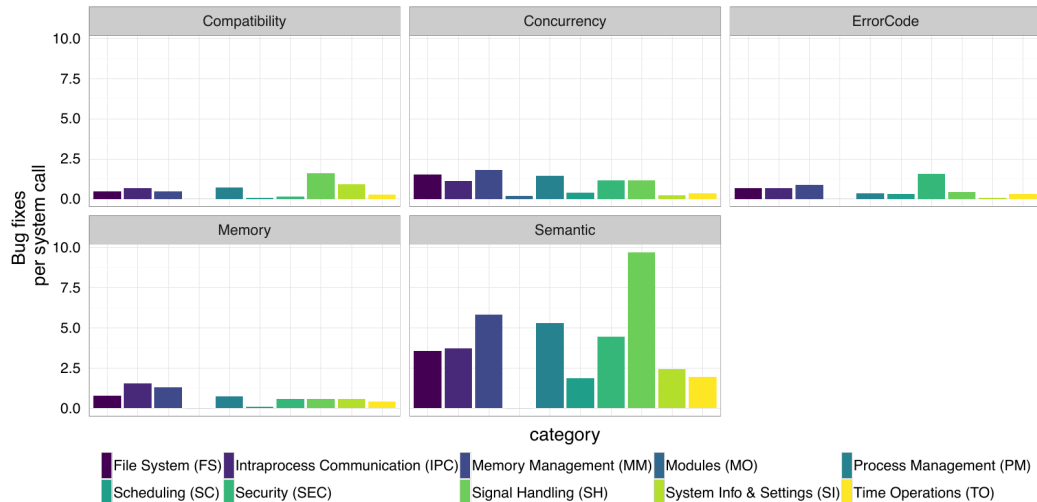
Introduction – The system calls with the most commits

System call	ALL	Restructuring ¹	Bug fix	Improvement
ptrace()	743	46%	35%	21%
signal()	714	53%	33%	18%
ioctl()	438	44%	32%	25%
futex()	257	35%	43%	23%
ipc()	253	51%	23%	30%
mmap()	213	30%	43%	31%
perf_event_open()	199	10%	46%	45%
readdir()	169	46%	41%	14%
splice()	166	30%	40%	25%

Introduction – Classifying Bug Fixes for Sysalls

- **Compatibility:** Compatibility-related bugs are caused by compatibility issues between architectures (e.g., 32-bit versus 64-bit).
 - **Concurrency:** Concurrency-related bugs are caused by issues with atom-icity, execution order, synchronization or locking, and lead to problems such as deadlock or race conditions.
 - **Error code:** Error code-related bugs are caused by returning the wrong error code or handling a returned error code incorrectly.
 - **Memory:** Memory-related bugs are caused by incorrect usage of the mem-ory, thereby introducing an issue such as a memory leak.
 - **Semantic** : Semantic bugs are bugs in the implementation of the system call-specific behaviour, such as the logic of the service provided by the system call.
- Signal handling system calls have the highest number of semantic (9.33) and compatibility-related (1.70) bug fixes per system call.

Introduction – Classifying Bug Fixes for Sysalls



Signal handling system calls have the highest number of semantic (9.33) and compatibility-related (1.70) bug fixes per system call.