

第 14 讲: Concurrency in OS Kernel

第一节: Introduction


陈渝

清华大学计算机系

yuchen@tsinghua.edu.cn

2020 年 5 月 18 日





Is Parallel Programming Hard, and, if so, What Can You Do About It?

Edited by Paul E. McKenney



Reference:

"Is Parallel Programming Hard, And, If So, What Can You Do About It?", Paul McKenney;

"C++ Concurrency in Action", ANTHONY WILLIAMS;

"CS510 - Advanced Topics in Concurrency", Jonathan Walpole

Locking In The Linux Kernel

- Why do we need locking in the kernel?
- Which problems are we trying to solve?
- What implementation choices do we have?
- Is there a one-size-fits-all solution?



Concurrency in Linux

- Linux is a symmetric multiprocessing (SMP) preemptible kernel
- Its has true concurrency
- And various forms of pseudo concurrency



Software-based preemption

- Voluntary preemption (sleep/yield)
- Involuntary preemption (preemptable kernel)
- And various forms of pseudo concurrency
- Solutions: don't do the former, disable preemption to prevent the latter



Hardware preemption

- Interrupt/trap/fault/exception handlers can start executing at any time
- Solutions: disable interrupts



Uniprocessor Example

```
preempt disable;  
    r1 = x;  
    r2 = x;  
preempt enable;  
assert (r1 == r2)
```



Uniprocessor Example

```
preempt disable;  
    r1 = x;  
    yield();  
    r2 = x;  
preempt enable;  
assert (r1 == r2)
```



Uniprocessor Example

```
interrupt disable;  
    r1 = x;  
    r2 = x;  
interrupt enable;  
assert (r1 == r2)
```



Solutions to pseudo-concurrency do not work in the presence of true concurrency

Alternatives include atomic operators, various forms of locking, RCU, and non-blocking synchronization

Locking can be used to provide mutually exclusive access to critical sections

- Locking can not be used everywhere, i.e., interrupt handlers can't block
- Locking primitives must support coexistence with various solutions for pseudo concurrency, i.e., we need hybrid primitives



Multiprocessor Example

```
interrupt disable;  
    r1 = x;  
    r2 = x;  
interrupt enable;  
assert (r1 == r2)
```



Simplest synchronization primitives

- Primitive operations that are indivisible

Two types

- methods that operate on integers
- methods that operate on bits

Implementation

- Assembly language sequences that use the atomic read- modify-write instructions of the underlying CPU architecture



Memory Invariance Example

```
r1 = atomic read x;  
r2 = atomic read x;  
assert (r1 == r2)
```



Atomic Integer Operators

```
atomic_t v;  
atomic_set(&v, 5); /* v = 5 (atomically) */  
atomic_add(3, &v); /* v = v + 3 (atomically) */  
atomic_dec(&v);    /* v = v - 1 (atomically) */  
  
printf("This will print 7: %d\n", atomic_read(&v));
```

”atomic_add(3,&v);” is NOT the same as ”atomic_add(1,&v); atomic_add(2,&v);”



Mutual exclusion for larger (than one operator) critical sections requires additional support
Spin locks are one possibility

- Single holder locks
- When lock is unavailable, the acquiring process keeps trying



Basic Use of Spin Locks

```
spinlock_t mr_lock = SPIN_LOCK_UNLOCKED;  
spin_lock(&mr_lock);  
/* critical section ... */  
spin_unlock(&mr_lock);
```

spin_lock()

- Acquires the spinlock using atomic instructions required for SMP

spin_unlock()

- Releases the spinlock



Interrupting a spin lock holder may cause problems

- Spin lock holder is delayed, so is every thread spin waiting for the spin lock
- Not a big problem if interrupt handlers are short
- Interrupt handler may access the data protected by the spin-lock
 - Should the interrupt handler use the lock?
 - Can it be delayed trying to acquire a spin lock?
 - What if the lock is already held by the thread it interrupted?



Solutions for Spin Locks and Interrupts

- Should the interrupt handler use the lock?
 - Can it be delayed trying to acquire a spin lock?
 - What if the lock is already held by the thread it interrupted?
-

- If data is only accessed in interrupt context and is local to one specific CPU we can use interrupt disabling to synchronize
- If data is accessed from other CPUs we need additional synchronization – spin lock
- Normal code (kernel context) must disable interrupts and acquire spin lock



Spin Locks & Interrupt Disabling

```
spinlock_t mr_lock = SPIN_LOCK_UNLOCKED;  
unsigned long flags;  
spin_lock_irqsave(&mr_lock, flags); /* critical section ... */  
spin_unlock_irqrestore(&mr_lock, flags);
```

spin_lock_irqsave()

- Disables interrupts locally
- Acquires the spinlock using instructions required for SMP

spin_unlock_irqrestore()

- Restores interrupts to the state they were in when the lock was acquired
 - Contention for semaphores causes blocking not spinning
 - Should not be used for short duration critical sections!
 - Can be used to synchronize with user contexts that might block or be preempted



Bottom Halves and Softirqs

Softirqs, tasklets and BHs are deferrable functions

- delayed interrupt handling work that is scheduled
- they can wait for a spin lock without holding up devices
- they can access non-CPU local data

Softirqs –the basic building block

- statically allocated and non-preemptively scheduled
- can not be interrupted by another softirq on the same CPU
- can run concurrently on different CPUs, and synchronize with each other using spin-locks

Bottom Halves

- built on softirqs
- can not run concurrently on different CPUs



Spin Locks & Deferred Functions

`spin_lock_bh()`

- Implements the standard spinlock
- Disables softirqs
- Allows the softirq to use non-preemption only

`spin_unlock_bh()`

- Releases the spinlock
- Enables softirqs



Spin Lock Rules

- Do not try to re-acquire a spinlock you already hold!
- Spinlocks should not be held for a long time!
- Do not sleep while holding a spinlock!



Semaphores are locks that are safe to hold for longer periods of time

- Contention for semaphores causes blocking not spinning
- Should not be used for short duration critical sections!
- Can be used to synchronize with user contexts that might block or be preempted



Semaphore Implementation

Implemented as a wait queue and a usage count

- wait queue: list of processes blocking on the semaphore
- usage count: number of concurrently allowed holders

down()

- Attempts to acquire the semaphore by decrementing the usage count and testing if it is negative
- Blocks if usage count is negative

up()

- releases the semaphore by incrementing the usage count and waking up one or more tasks blocked on it



Semaphore Implementation

Implemented as a wait queue and a usage count

- wait queue: list of processes blocking on the semaphore
- usage count: number of concurrently allowed holders

`down_interruptible()`

- Returns `-EINTR` if signal received while blocked
- Returns 0 on success

`down_trylock()`

- Attempts to acquire the semaphore
- On failure it returns nonzero instead of blocking



No need to synchronize concurrent readers unless a writer is present

Both spin locks and semaphores have reader/writer variants



Reader-Writer Spin Locks

```
rwlock_t mr_rwlock = RW_LOCK_UNLOCKED;  
  
read_lock(&mr_rwlock);  
/* critical section (read only) ... */  
read_unlock(&mr_rwlock);  
  
write_lock(&mr_rwlock);  
/* critical section (read and write) ... */  
write_unlock(&mr_rwlock);
```



Reader-Writer Semaphores

```
struct rw_semaphore mr_rwsem;  
init_rwsem(&mr_rwsem);  
  
down_read(&mr_rwsem); /* critical region (read only) ... */  
up_read(&mr_rwsem);  
  
down_write(&mr_rwsem); /* critical region (read and write) ... */  
up_write(&mr_rwsem);
```



Wow! Why does one system need so many different ways of doing synchronization?

- Actually, there are more ways to do synchronization in Linux, this is just “locking” !



One size does not fit all:

- Need to be aware of different contexts in which code executes (user, kernel, interrupt etc) and the implications this has for whether hardware or software preemption or blocking can occur
- The cost of synchronization is important, particularly its impact on scalability

